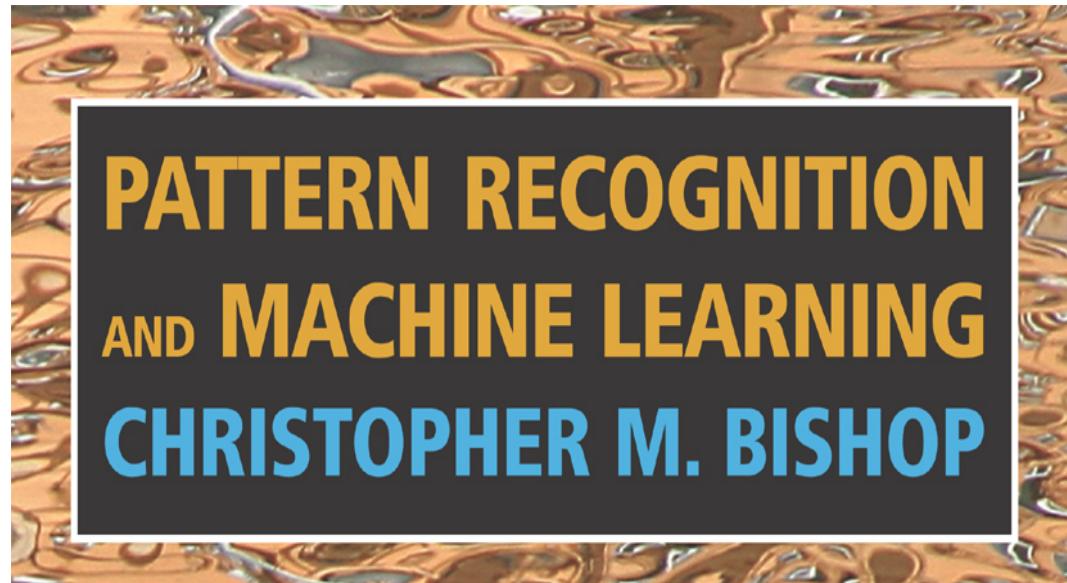


Image Segmentation

Suyash P. Awate



K-means Clustering

- **Given:** Data points in D-dimensional space
 - Data points can be pixel values
 - Values can be 1-D (intensities) or multi-D
- **Goal:** Partition data points into 'k' clusters
 - 'k' is known / fixed
 - Label each point to belong to one (and exactly one) of the groups
 - Hard (discrete) decision made for each data point

K-means Clustering

- Consider k groups S_1, \dots, S_k
- Consider a **representative** point for each group
 - Called “**mean**”
- **Optimization problem**
 - Find : **Optimal Groups** and their **Means**
 - Penalize sum of squared distances of each data point to the mean of the group it is assigned to
 - Minimize means μ_i and labels S_i

$$\sum_{i=1}^k \sum_{x_j \in S_i} \|x_j - \mu_i\|^2$$

K-means Clustering

$$\sum_{i=1}^k \sum_{\mathbf{x}_j \in S_i} \|\mathbf{x}_j - \boldsymbol{\mu}_i\|^2$$

- Optimization algorithm (2 steps)

(0) An initial solution for the set of means is given

- Can equivalently assume an initial solution for the groups

(1) Given means, what are the optimal groups ?

- Which group will you assign to each data point so that the objective function is guaranteed to decrease ?
- Assign the group whose mean is the closest to data point

(2) Given groups, what are the optimal means ?

- For a specific group, what mean will minimize the penalty for that group ?
- Assign mean = average of the data points in that group

K-means Clustering

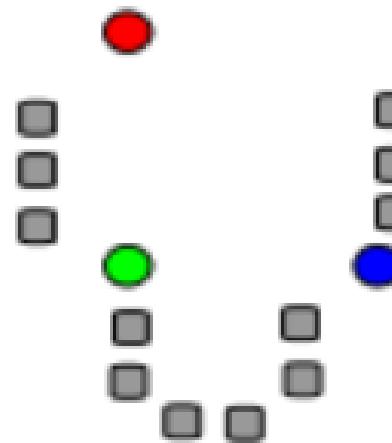
- Optimization algorithm
 - Is it guaranteed to converge / terminate. Why or why not ?
 - Will it converge to a global optimum ? Why or why not ?
 - Is the objective function quadratic ? convex ?

K-means Clustering

- Optimization algorithm
 - Is it guaranteed to converge / terminate. Why or why not ?
 - Yes. Sample size is finite. So # labelings is finite. Each relabeling reduces a non-negative penalty function; else you've stopped.
 - Will it converge to a global optimum ? Why or why not ?
 - No guarantees. See a counter example soon
 - Is the objective function quadratic ? convex ?
 - No. Combinatorial optimization. NP hard.
 - Given groups, objective function for means is quadratic

K-means Clustering

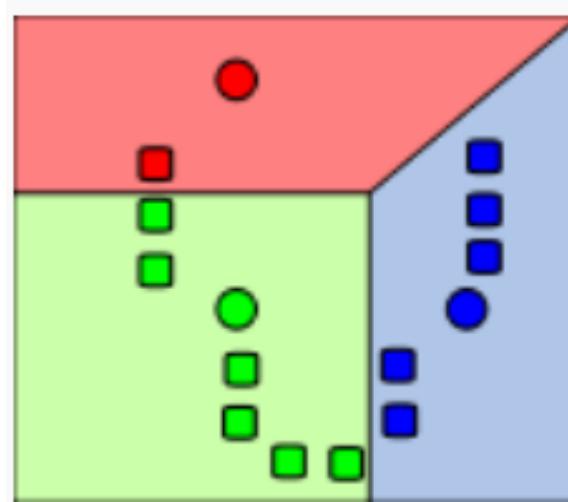
- K-means example
 - Data points in black
 - $K = 3$ (fixed)
 - Initial means in R, G, B
 - 1st K points in the list ?
 - Randomly selected K points in the list ?
 - Randomly generated K points in the space ?
 - Can we do better than any of these ? Yes. Will see soon.



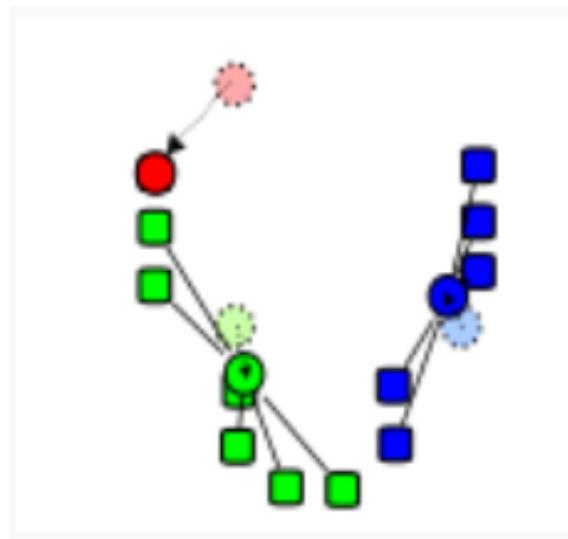
K-means Clustering

- K means example

- Given means,
update groups



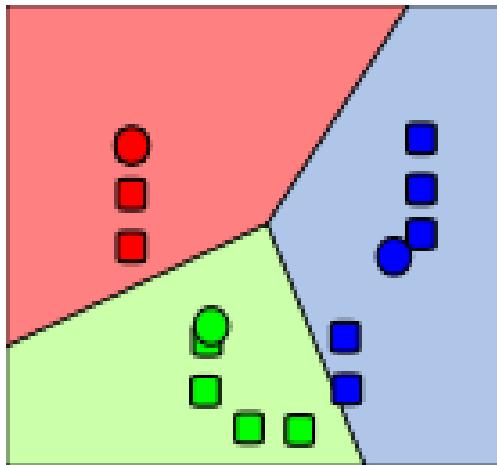
- Given groups,
update means



K-means Clustering

- K means example

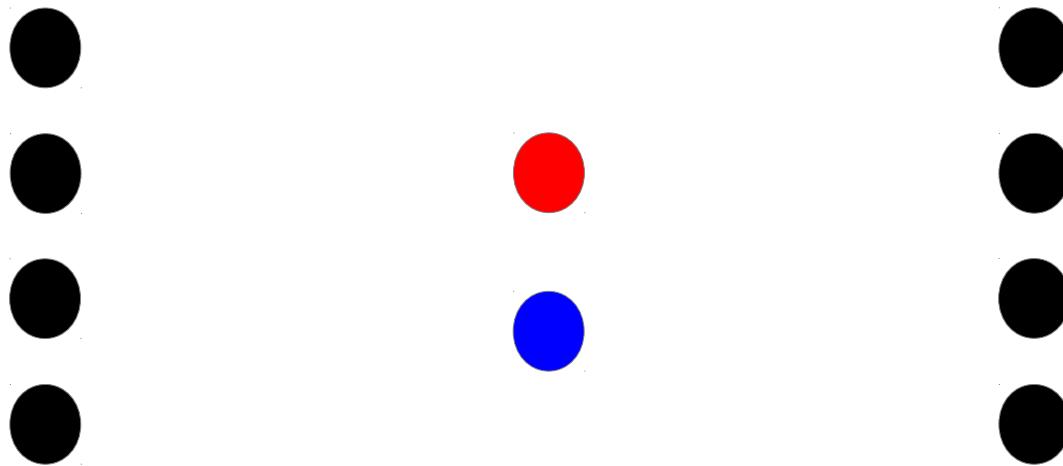
- Given means,
update groups



- Repeat until objective function stops reducing

K-means Clustering

- Convergence to local minimum
 - Example 1

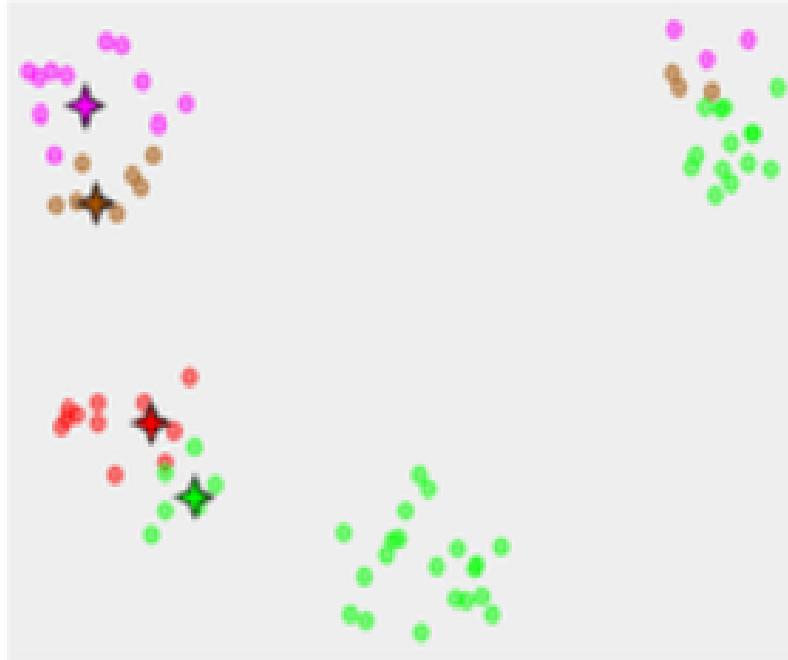


K-means Clustering

- Convergence to local minimum

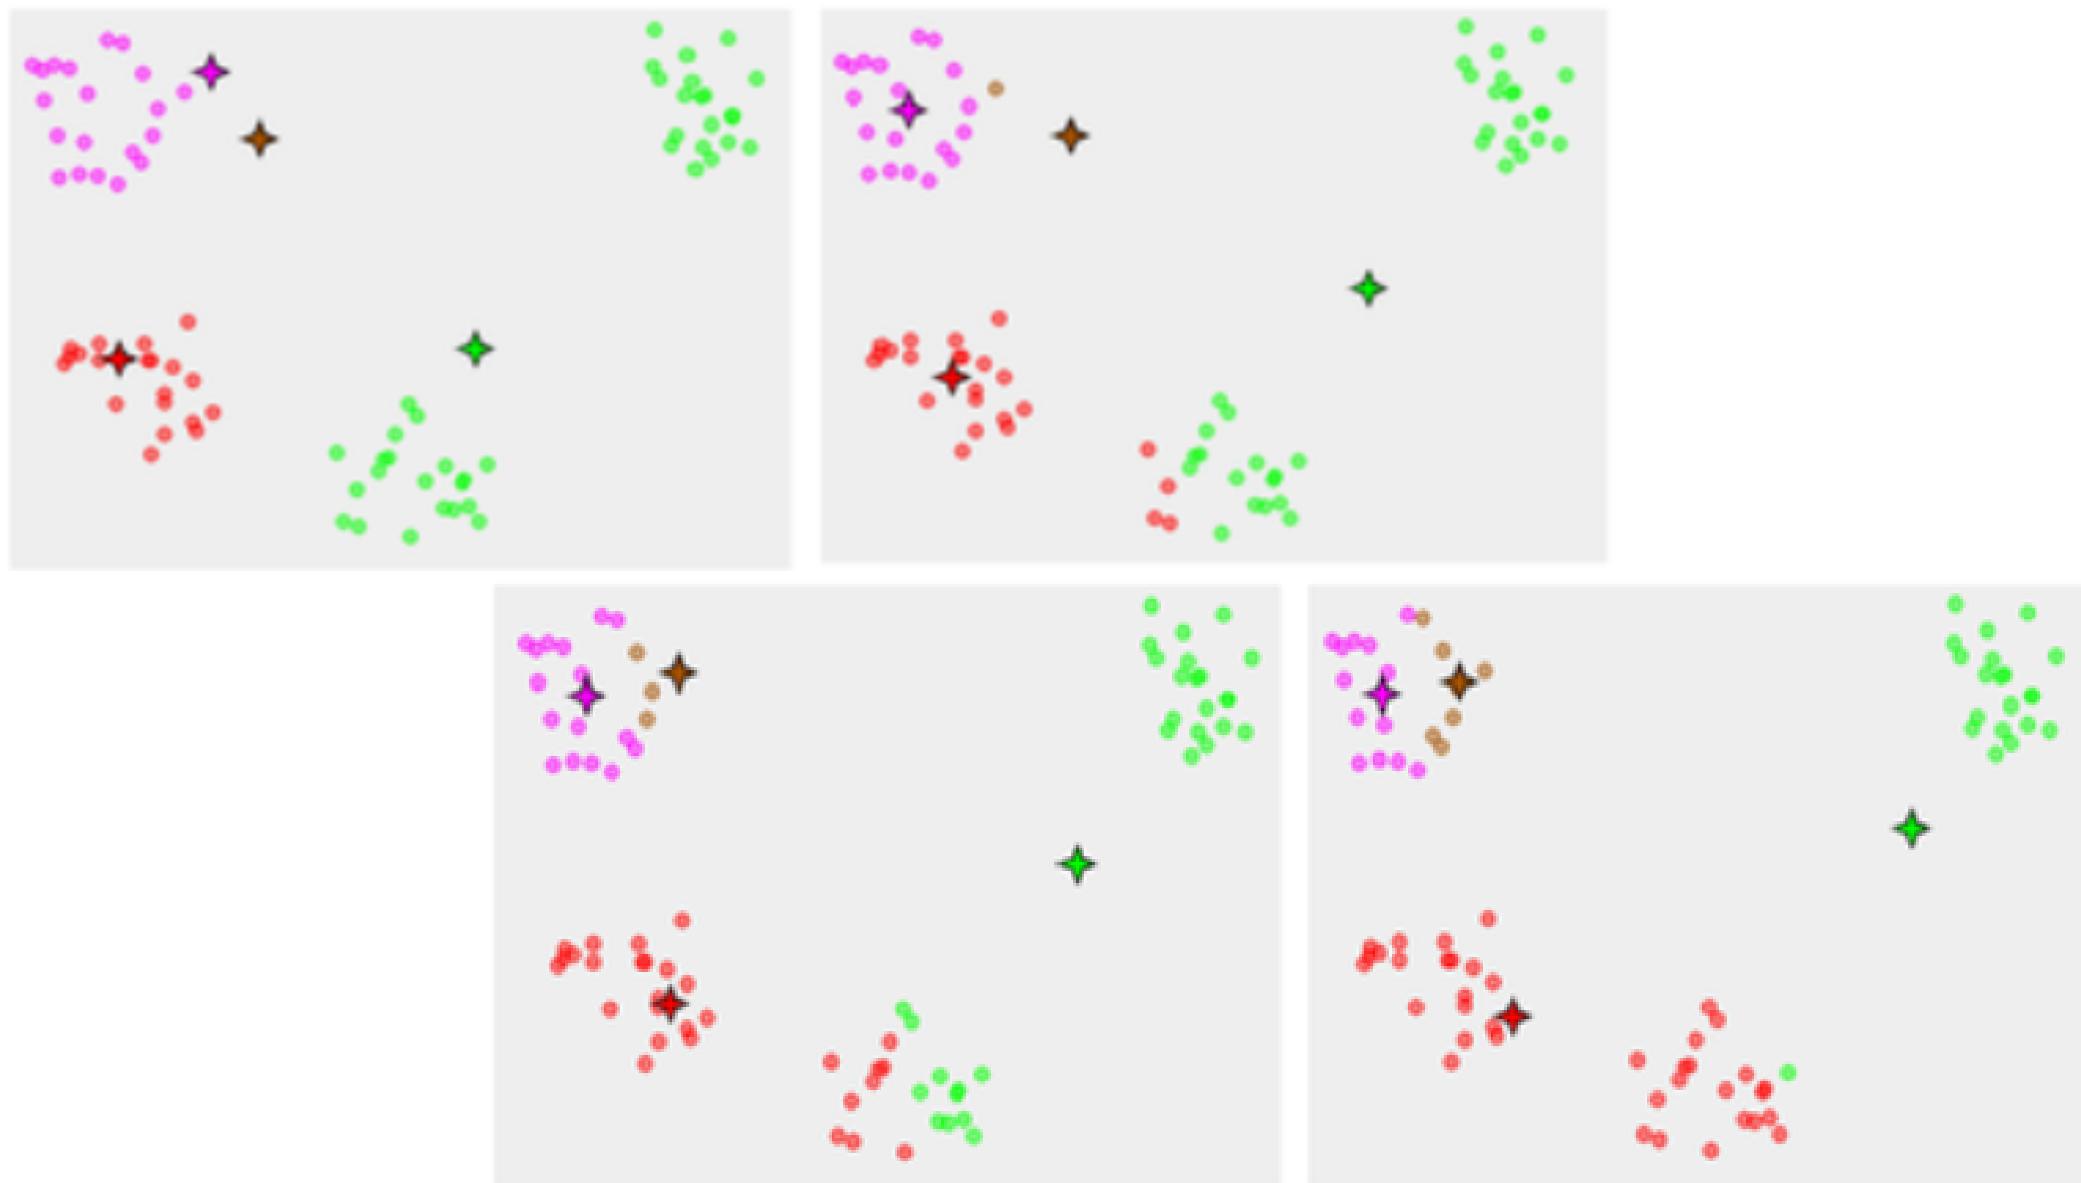
- Example 2

- $k = 4$
 - Initial means



K-means Clustering

- Convergence to local minimum

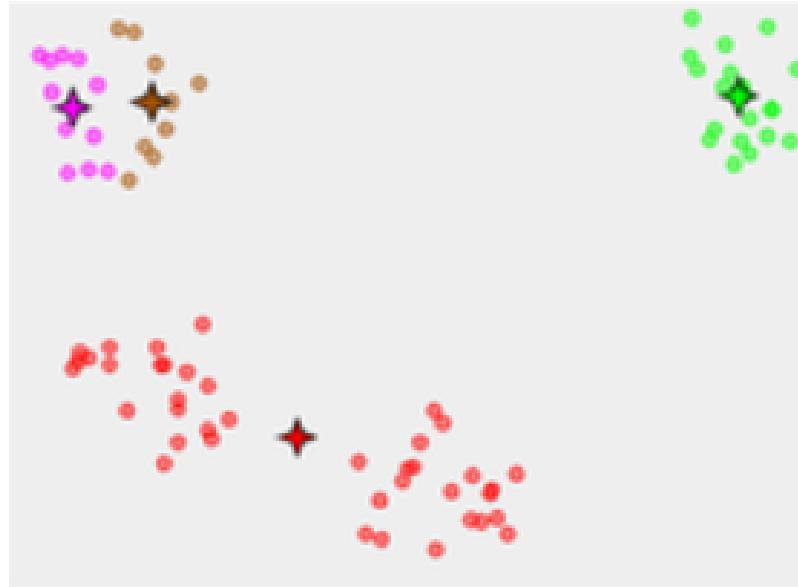


K-means Clustering

- Convergence to local minimum

- Example 2

- Final solution
 - Local minimum



K-means Clustering

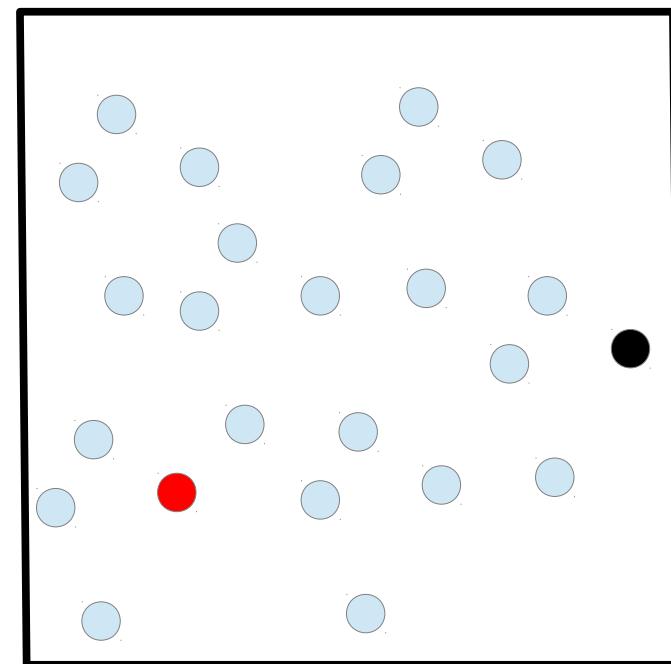
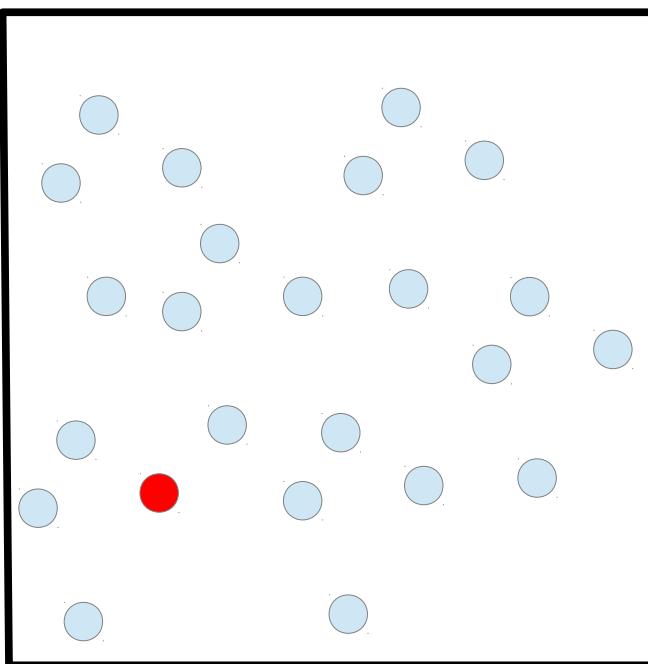
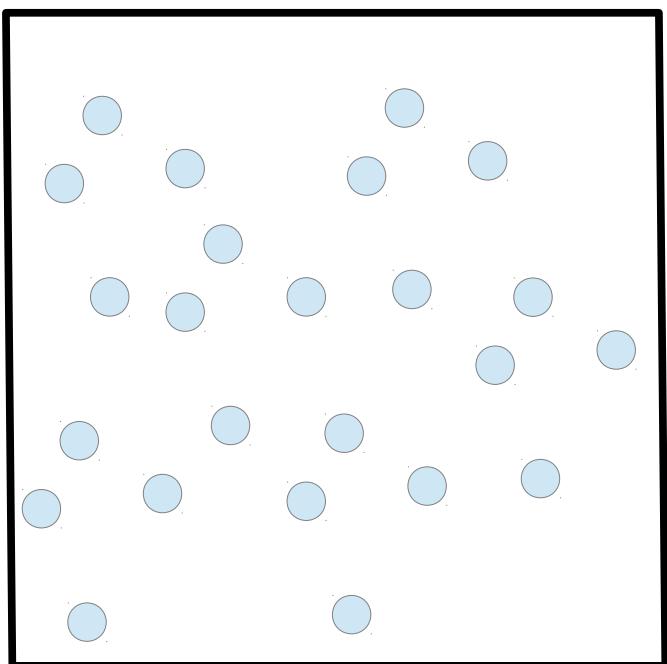
- Hope to get global optimum ?
 - Brute-force approach
 - Check objective-function value for all possible clusterings
 - N points and k clusters $\rightarrow k^N$ possible clusterings
 - Re-run algorithm using different initial mean estimates
 - Introduce randomization in the initialization
 - Better initialization

K-means Clustering

- Better initialization
 - Farthest-point clustering
 - Algorithm
 - Heuristic
 - Given: data points, value of 'k'
 - Select 1st “mean” at random from among the data points
 - Select the other means, among the data points, iteratively
 - Assume: at some iteration, (n-1) means have been selected
 - Select the n-th mean to be the data point that maximizes the minimum distance among distances to the current (n-1) means

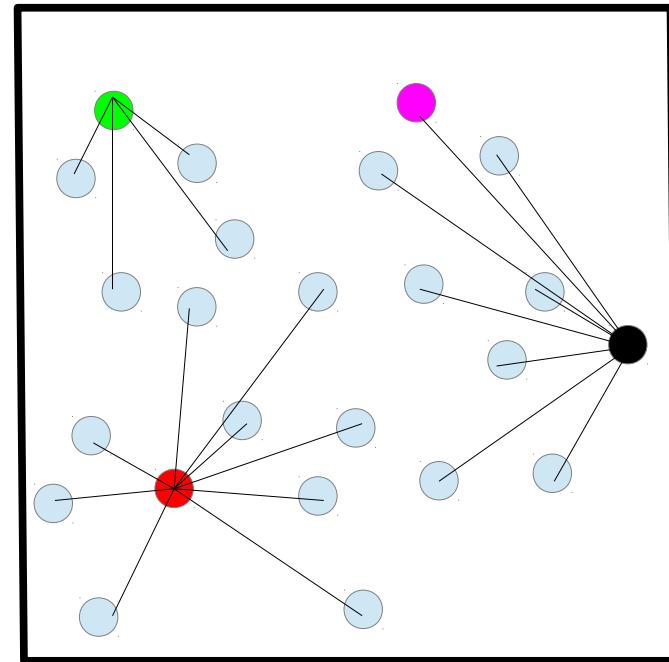
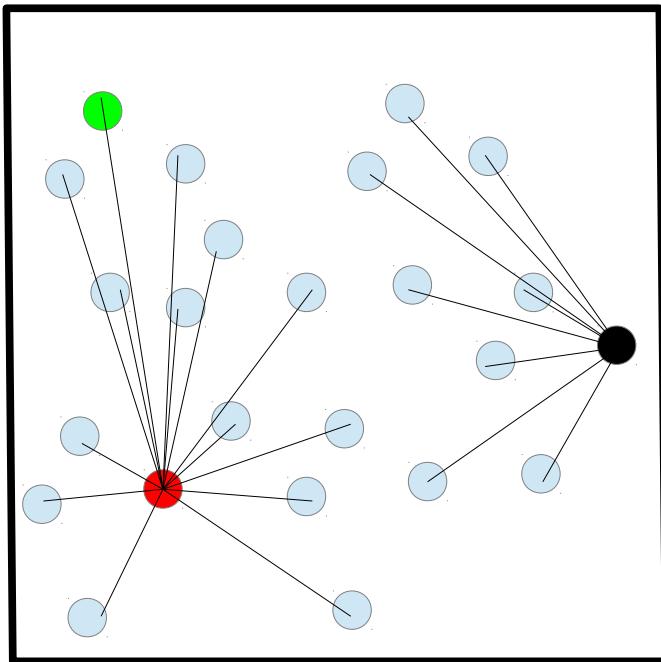
K-means Clustering

- Better initialization
 - Farthest-point clustering



K-means Clustering

- Better initialization
 - Farthest-point clustering
 - Select the point that is farthest from the current means !
 - Spreads out the selected points (roughly) uniformly



K-means Clustering

- Better initialization
 - Farthest-point clustering
 - Sensitive to outliers
 - K-means++
 - Don't pick farthest point
 - Pick point 'x' with probability directly related to minimum distance

pick $x \in S$ uniformly at random and set $T \leftarrow \{x\}$

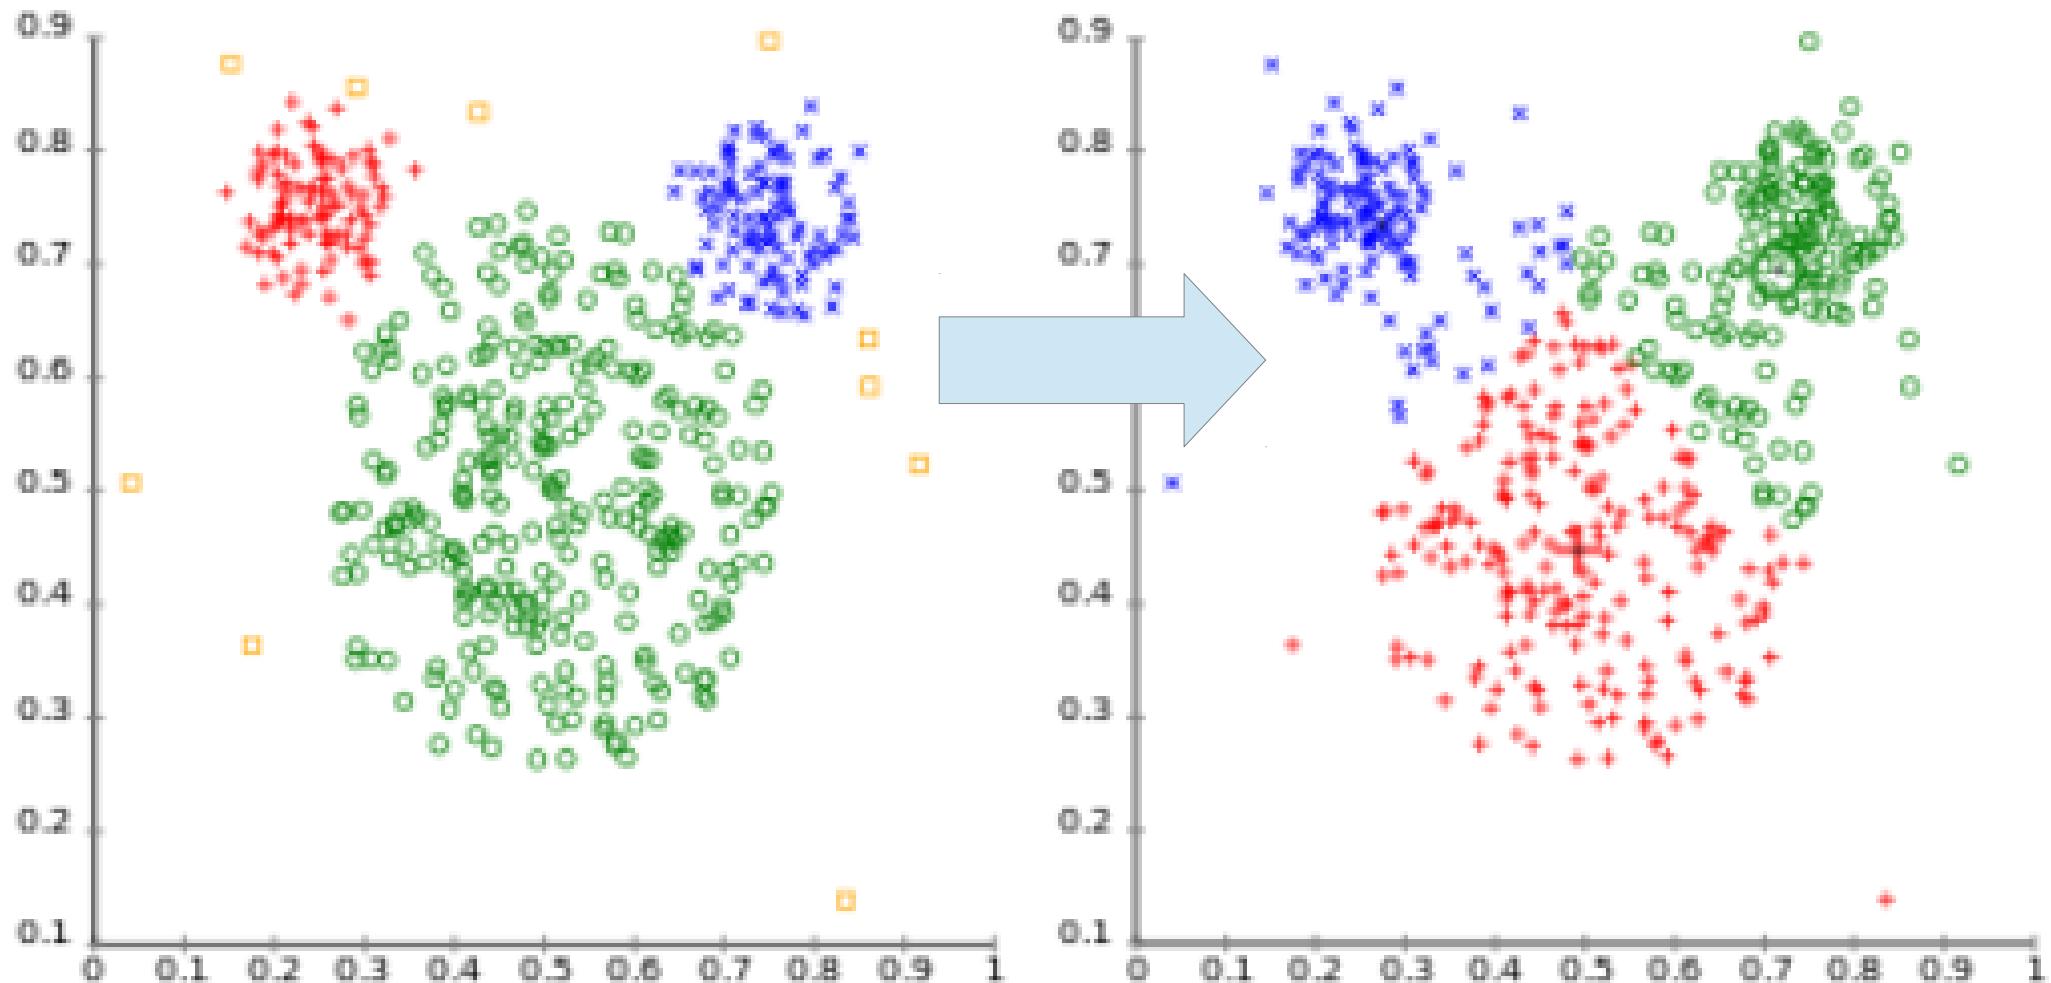
while $|T| < k$:

pick $x \in S$ at random, with probability proportional to $\text{cost}(x, T) = \min_{z \in T} \|x - z\|^2$
 $T \leftarrow T \cup \{x\}$

- Arthur, D. and Vassilvitskii, S. (2007). k-means++: the advantages of careful seeding. In ACM-SIAM Symposium on Discrete Algorithms
- <https://cseweb.ucsd.edu/~dasgupta/291-geom/kmeans.pdf>

K-means Clustering

- Limitations
 - Tendency to produce clusters with equal spreads



K-means Clustering

- Limitations
 - 1) Binary / hard assignments of each data point to cluster
 - What about points near interface of 2 clusters ?
 - 2) Ignore covariance of the cluster
 - Assumes spherical clusters
 - Because relies on simple Euclidean distance
 - Doesn't handle ellipsoidal clusters systematically

Evaluating Quality of Clustering

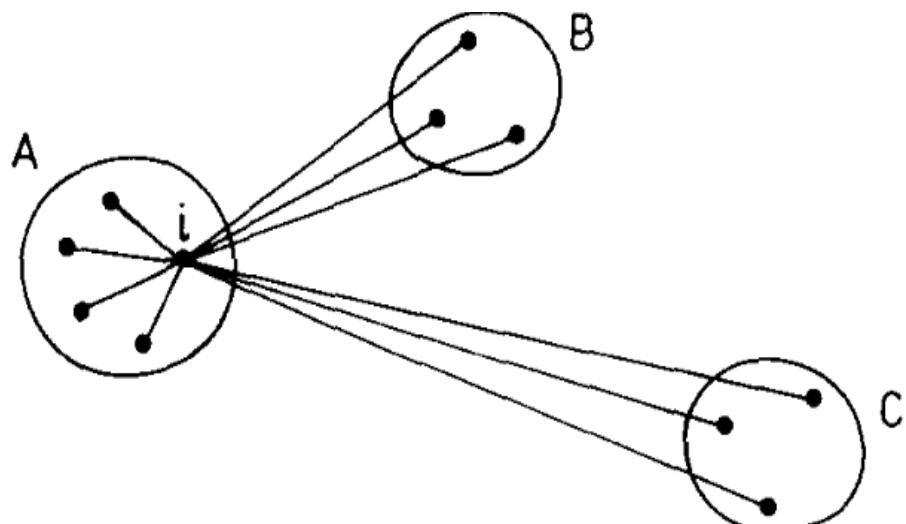
- “Silhouette” analysis

<http://dx.doi.org/10.1016%2F0377-0427%2887%2990125-7>

- Graphical representation
 - Measures how contained is each datum within its cluster

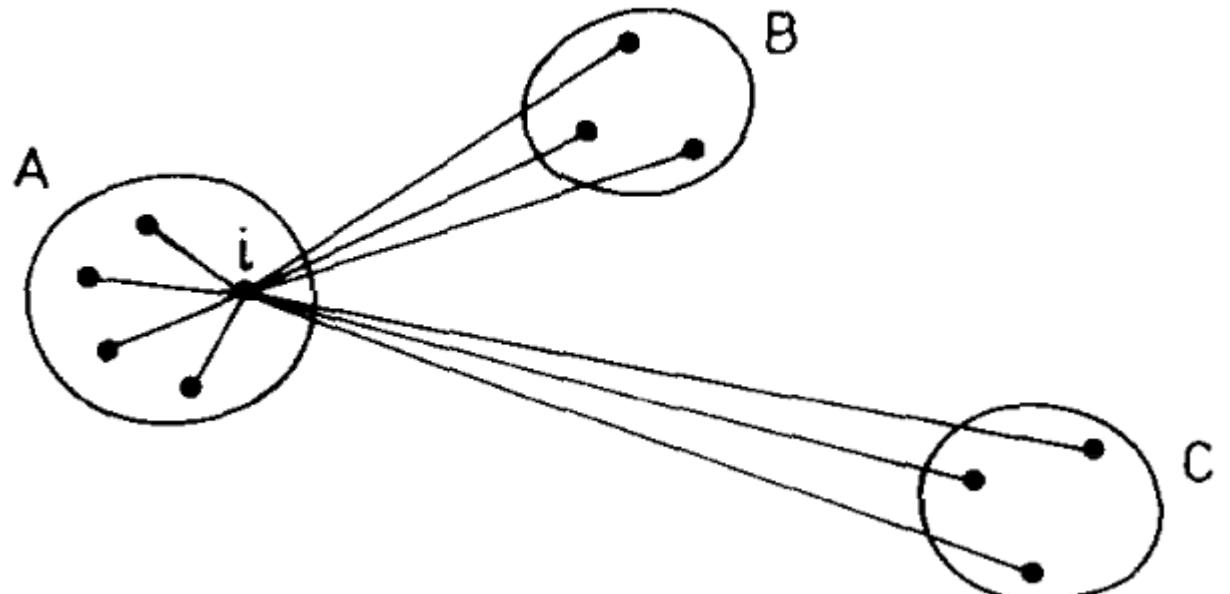
Evaluating Quality of Clustering

- For each datum x_i assigned to cluster A
 - $a_i = \text{avg of distances between } x_i \text{ & } x_j \text{ within cluster A}$
 - $a_i > 0$
 - Small $a_i \rightarrow$ more compact cluster
 - $b_i = ?$
 - 1) avg of distances between x_i & x_j within another cluster B
 - 2) select minimum of these over all other clusters B
 - $b_i > 0$
 - Large $b_i \rightarrow$ more compact cluster A



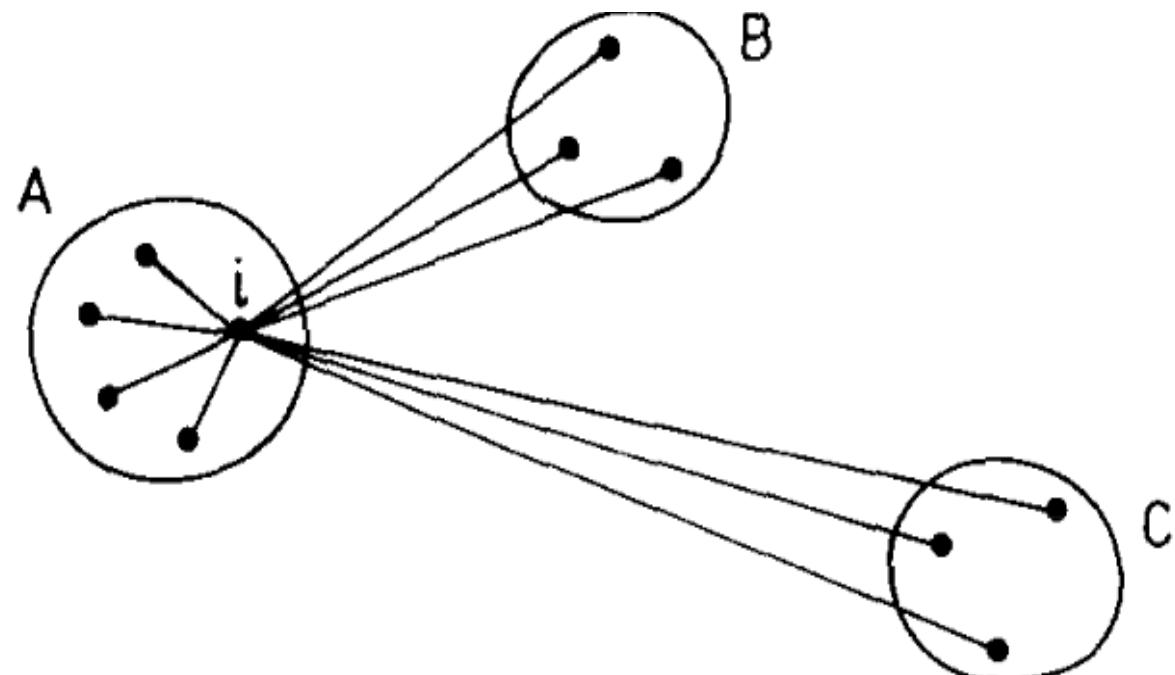
Evaluating Quality of Clustering

- For each datum, define $s_i = (b_i - a_i) / \max(a_i, b_i)$
- When $b_i = a_i$
 - $s_i = 0$
- When $b_i > a_i$
 - $s_i = 1 - a_i / b_i$
 - $0 < s_i \leq 1$
- When $a_i > b_i$
 - $s_i = b_i / a_i - 1$
 - $-1 \leq s_i < 0$



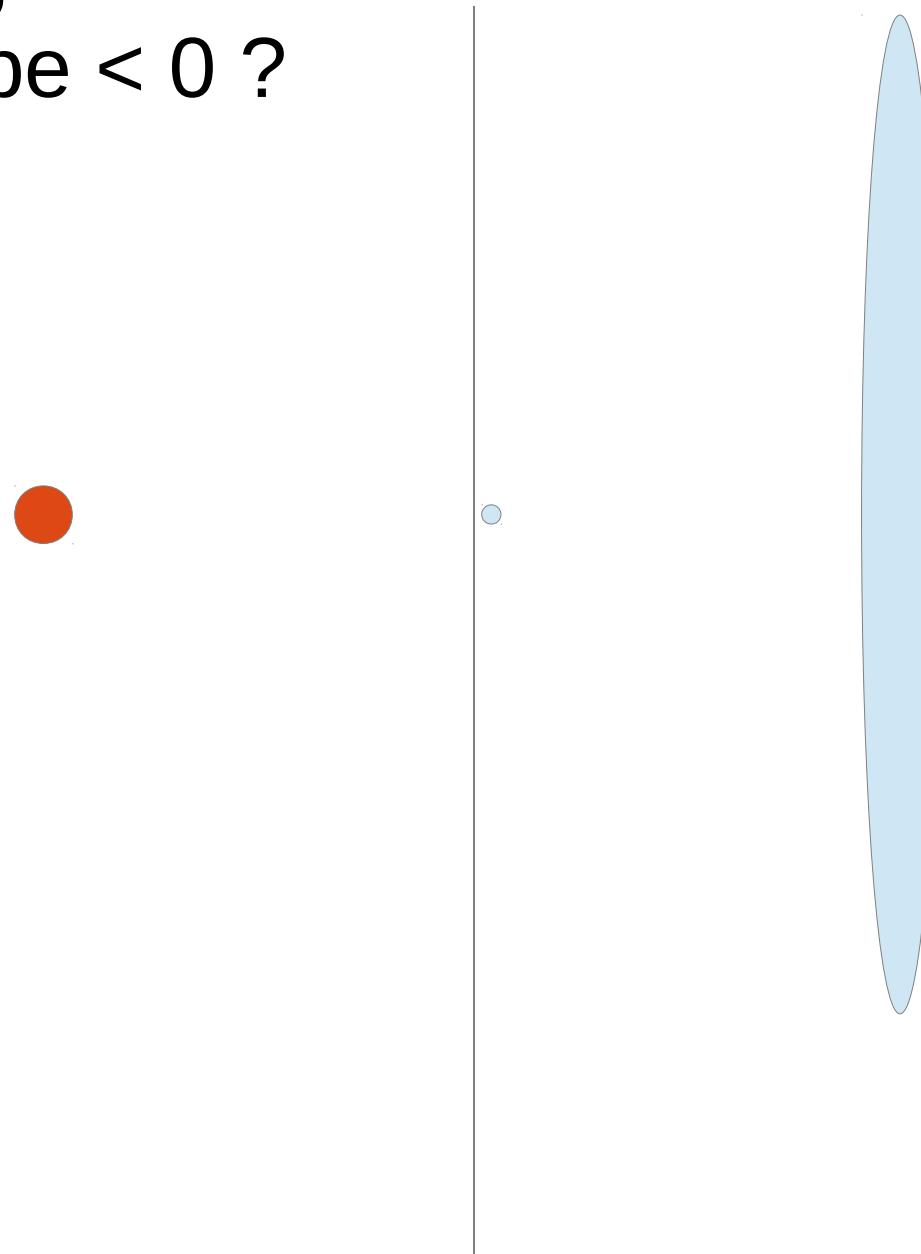
Evaluating Quality of Clustering

- We desire s_i close to +1
 - i.e., $a_i \ll b_i$
- What does it mean to have s_i close to 0 ?
 - Datum x_i lies ...
- What does it mean to have s_i close to -1 ?
 - Datum x_i lies ...



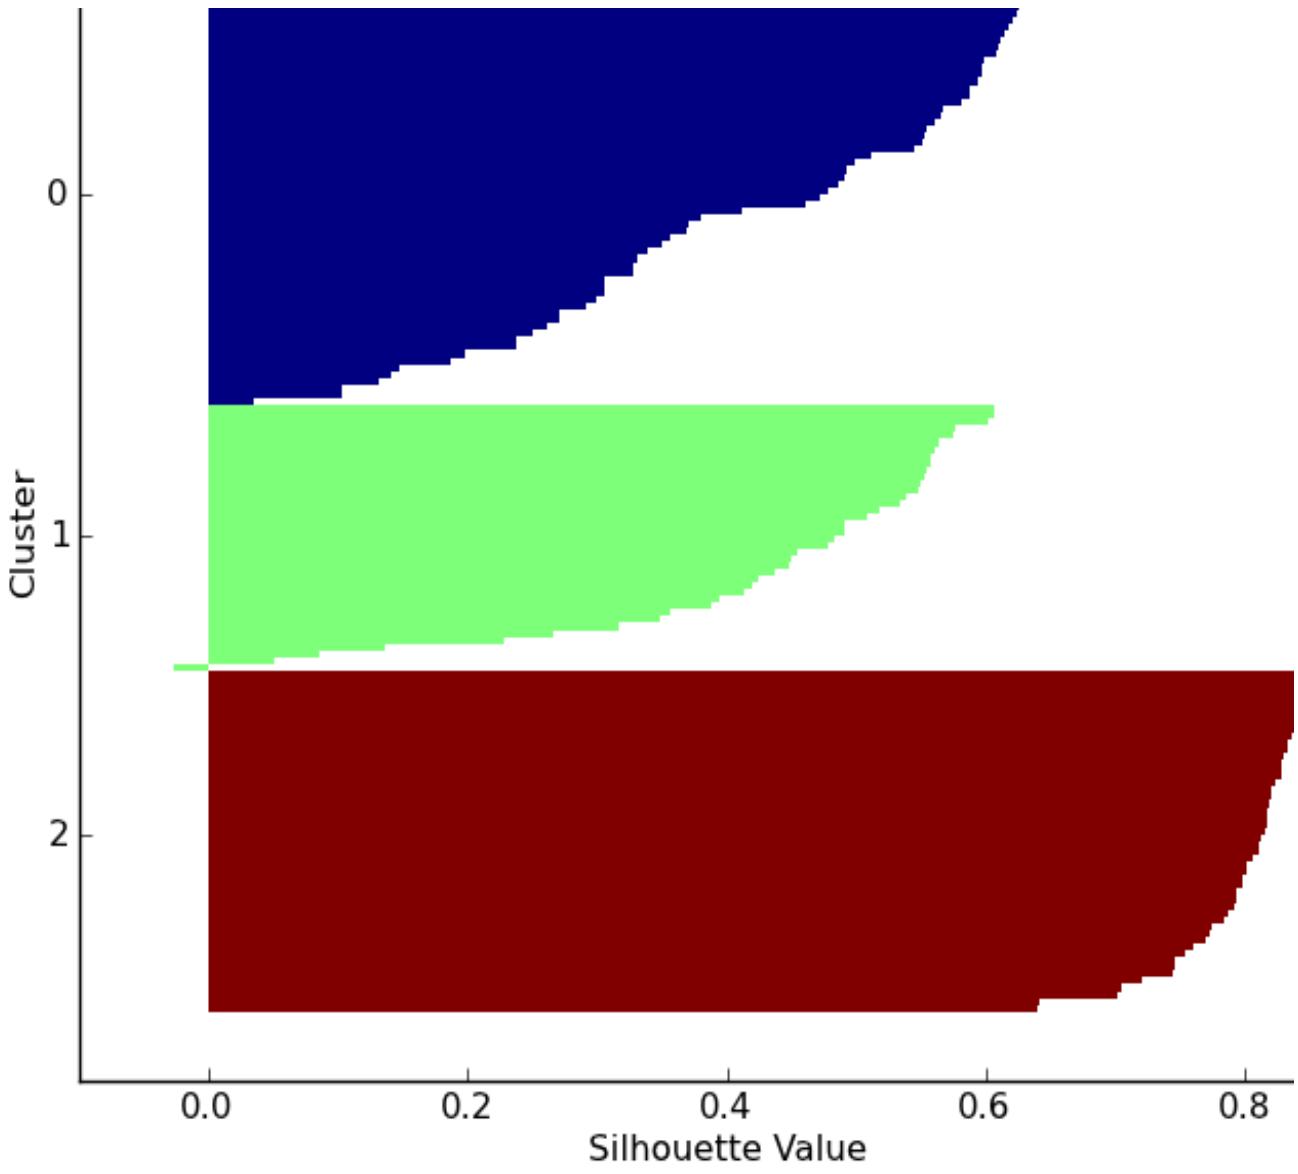
Evaluating Quality of Clustering

- For k-means,
when can s_i be < 0 ?



Evaluating Quality of Clustering

- For each cluster, visualize graph of s_i over i
 - This is the “silhouette” / outline
- May be used to tune # of clusters K or other underlying parameters



Fuzzy-C-Means (FCM) Clustering

- Generalization of K-means clustering
 - Assigns, to each data point, a membership to belong to each cluster
 - Membership in $[0,1]$
 - For each data point, sum of memberships over all clusters = 1
 - Produces a “soft” / “fuzzy” segmentation
 - K-means produces a “hard” / “crisp” segmentation

Fuzzy-C-Means Clustering

- Given
 - Data = { y_j }, $j = 1, \dots, N$
 - Number of clusters = K (known / fixed)
- Memberships
 - u_{jk} = membership of j -th point in k -th cluster
 - (i) memberships are positive: $\forall j, \forall k, u_{jk} \geq 0$ and
 - (ii) for every data point, memberships sum to one: $\forall j, \sum_k u_{jk} = 1$
- “ q ” = user-defined parameter
 - Controls fuzziness of the clusters / memberships
 - $q > 1$

Fuzzy-C-Means Clustering

- Objective function to be minimized

$$\sum_{j=1}^N \sum_{k=1}^K u_{jk}^q (y_j - c_k)^2$$

- Penalize squared distance of point j from mean of class k
- Weight penalty based on membership u_{jk}

- Constraints

$$\forall j, \sum_k u_{jk} = 1$$

- Positivity constraint on memberships:
 - Isn't enforced explicitly
 - Gets satisfied automatically

Fuzzy-C-Means Clustering

- FCM reduces to K-means
 - What happens if you force u_{jk} to be binary ?
 - What happens to the objective function ?

Fuzzy-C-Means Clustering

- FCM reduces to K-means
 - What happens if you force u_{jk} to be binary ?
 - $u_{jk} = 1$ for exactly one cluster k
 - $u_{jk'} = 0$ for all other clusters k'
 - What happens to the objective function ?
 - Objective function becomes exactly the same as that for K means

Fuzzy-C-Means Clustering

- **Constrained Optimization**

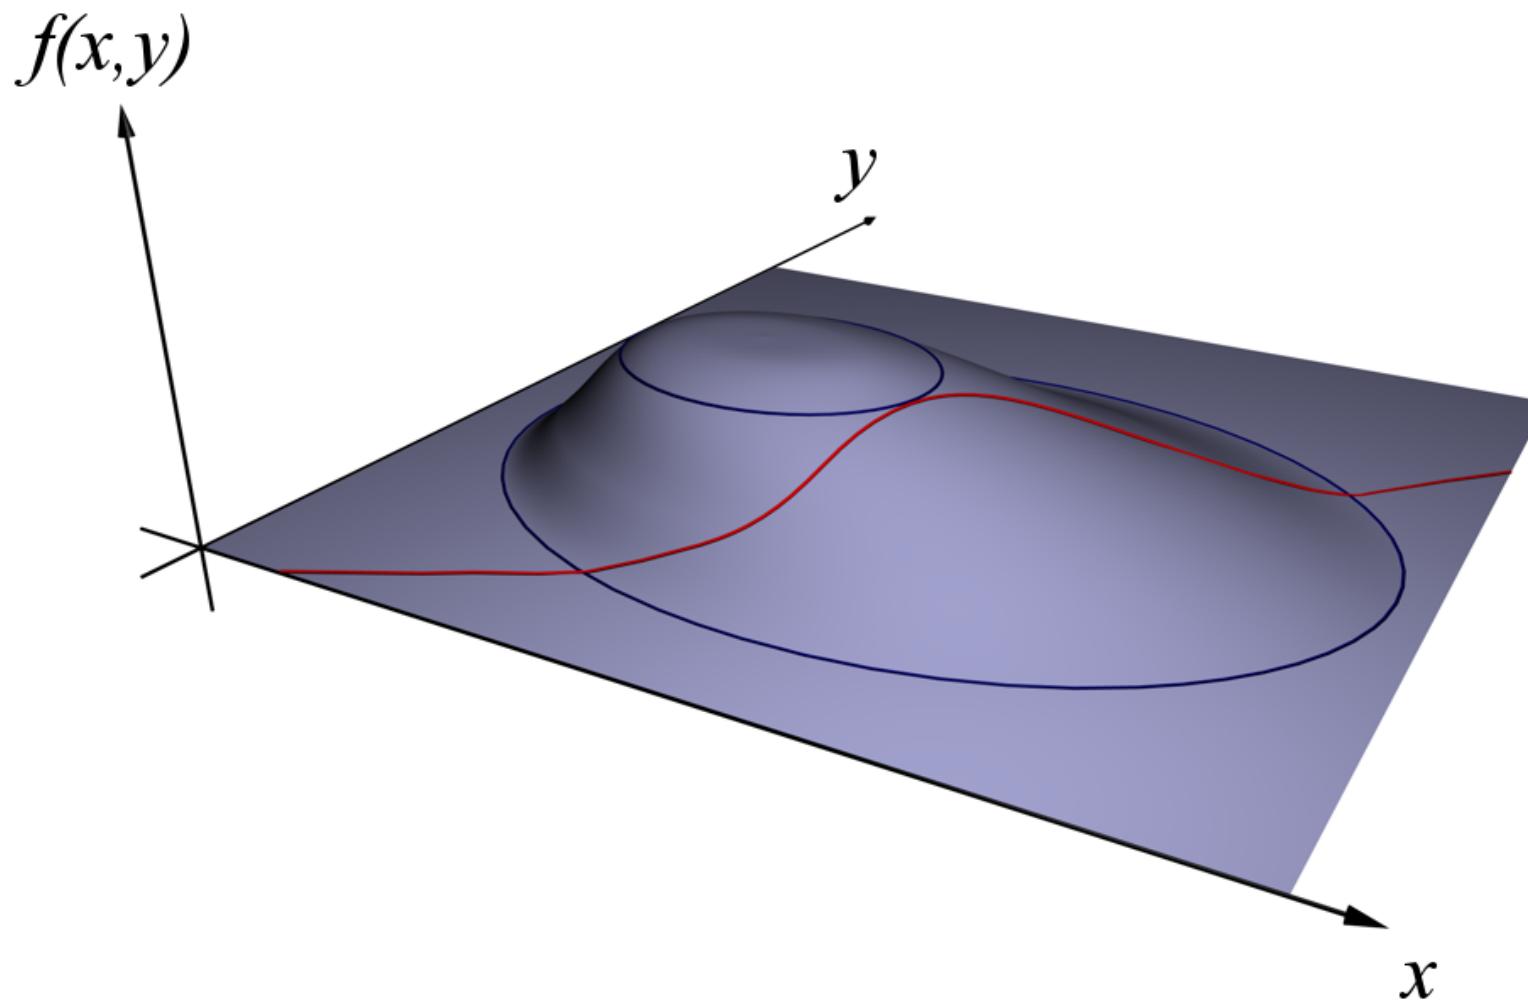
- Objective function type = Non-convex
- Constraint type = Equality
- Approach : Method of Lagrange multipliers

Fuzzy-C-Means Clustering

- Method of Lagrange Multipliers

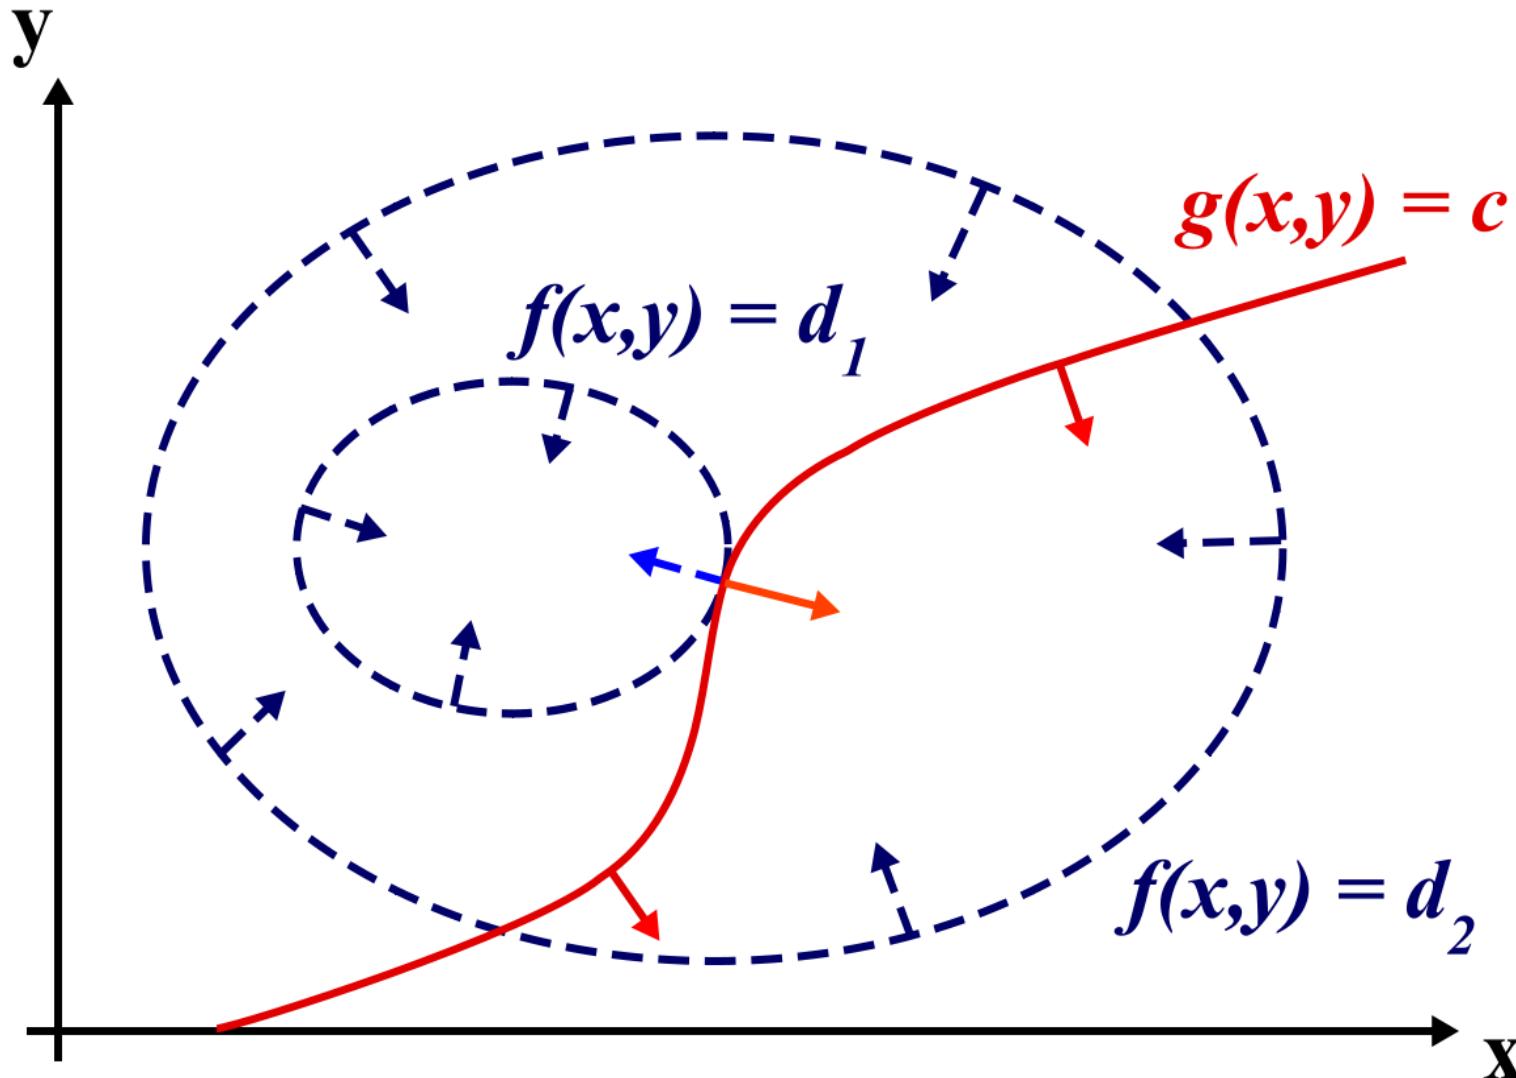
- Optimization problem

- Maximize : $f(x,y)$
 - Subject to (equality / boundary) constraint : $g(x,y) = c$



Fuzzy-C-Means Clustering

- Method of Lagrange Multipliers
 - What should happen at the optimal solution (x^*,y^*) ?



Fuzzy-C-Means Clustering

- Method of Lagrange Multipliers
 - What should happen at the optimal solution (x^*,y^*) ?
 - (1) $g(x^*,y^*) = c$
 - (2) At (x^*,y^*) , gradients of objective function and constraint **must** be parallel to each other

$$\nabla_{x,y} f = -\lambda \nabla_{x,y} g,$$

for some λ

where

$$\nabla_{x,y} f = \left(\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right), \quad \nabla_{x,y} g = \left(\frac{\partial g}{\partial x}, \frac{\partial g}{\partial y} \right)$$

Fuzzy-C-Means Clustering

- Method of Lagrange Multipliers
 - Specify these 2 conditions in the following way
 - Introduce a function called the **Lagrangian**
 - $L(x, y, \lambda) = f(x,y) + \lambda (g(x,y) - c)$
 - λ is called the **Lagrange parameter**
 - Solve for x, y, λ
 - Assign partial derivatives of Lagrangian, w.r.t. x, y, λ , to zero
 - 3 equations, 3 unknowns
 - Partial derivative w.r.t. 'x' and 'y' gives
 - $d f(x,y) / dx = - \lambda d g(x,y) / dx$
 - $d f(x,y) / dy = - \lambda d g(x,y) / dy$
 - Partial derivative w.r.t. ' λ ' gives us back the constraint
 - $g(x,y) = c$

Joseph-Louis Lagrange

- Advisors: Euler + ...
- Students: Fourier, Poisson, ...
- It was not until he was 17 that he showed any taste for math
- For 20 years in Berlin, except for a short time when he was ill, he produced ~ 1 paper / month
- Lagrange is one of the 72 prominent French scientists who were commemorated on plaques at the 1st stage of Eiffel Tower



Joseph-Louis Lagrange

Lagrangian [\[edit\]](#)

- [Lagrangian](#)
- [Lagrangian analysis](#)
- [Lagrangian coordinates](#)
- [Lagrangian derivative](#)
- [Lagrangian drifter](#)
- [Lagrangian foliation](#)
- [Lagrangian Grassmannian](#)
- [Lagrangian intersection Floer homology](#)
- [Lagrangian mechanics](#)
- [Lagrangian mixing](#)
- [Lagrangian point](#)
- [Lagrangian relaxation](#)
- [Lagrangian submanifold](#)
- [Lagrangian subspace](#)
- [Nonlocal Lagrangian](#)
- [Proca lagrangian](#)
- [Special Lagrangian submanifold](#)

Lagrange [\[edit\]](#)

- [Euler–Lagrange equation](#)
- [Green–Lagrange strain](#)
- [Lagrange bracket](#)
- [Lagrange–d'Alembert principle](#)
- [Lagrange error bound](#)
- [Lagrange form](#)
- [Lagrange interpolation](#)
- [Lagrange invariant](#)
- [Lagrange inversion theorem](#)
- [Lagrange multiplier](#)
- [Lagrange number](#)
- [Lagrange point colonization](#)
- [Lagrange polynomial](#)
- [Lagrange property](#)
- [Lagrange reversion theorem](#)
- [Lagrange resolvent](#)
- [Lagrange stream function](#)



Lagrange's [\[edit\]](#)

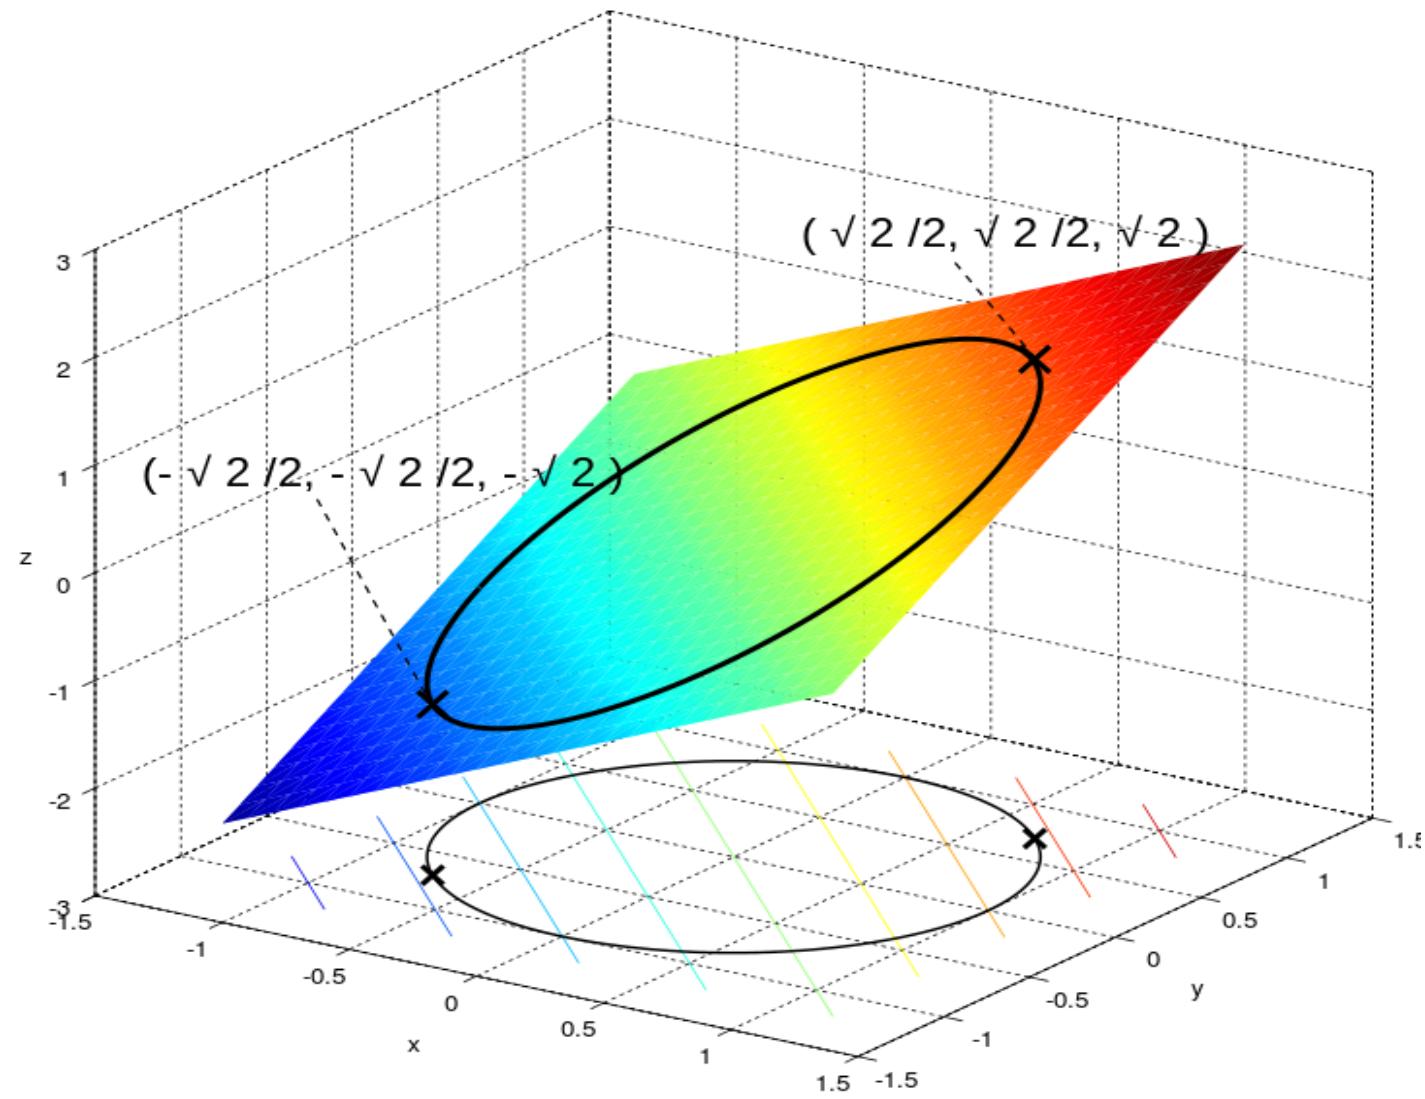
- [Lagrange's approximation theorem](#)
- [Lagrange's formula \(disambiguation\)](#)
- [Lagrange's identity \(disambiguation\)](#)
- [Lagrange's theorem \(group theory\)](#)
- [Lagrange's theorem \(number theory\)](#)
- [Lagrange's four-square theorem](#)
- [Lagrange's trigonometric identities](#)

Fuzzy-C-Means Clustering

- Method of Lagrange Multipliers
 - Lagrangian optimization gives a **necessary** condition
 - Not a sufficient condition
 - Example:
maximize
 $f(x,y) = x + y$
subject to
 $x^2 + y^2 = 1$
 - How does the constraint look like (in 2D XY plane) ?
 - How do the objective-function contours look like (in 2D XY plane) ?

Fuzzy-C-Means Clustering

- Method of Lagrange Multipliers
 - Lagrangian optimization gives a **necessary** condition
 - Not a sufficient condition
 - Example:
maximize
 $f(x,y) = x + y$
subject to
 $x^2 + y^2 = 1$



Fuzzy-C-Means Clustering

- FCM Optimization via Lagrange Multipliers
 - Start with an **initial solution** for **memberships** and **cluster means**
 - Use method of Lagrange multipliers to **iteratively** improve solution
 - Lagrangian

$$L(\{u_{jk}\}, \{c_k\}, \{\lambda_j\}) := \sum_{j=1}^N \sum_{k=1}^K u_{jk}^q (y_j - c_k)^2 + \lambda_j \left(\sum_{k=1}^K u_{jk} - 1 \right)$$

- Denote
 - $d_{jk} := (y_j - c_k)^2$
 - Squared distance of j-th point from k-th cluster mean

Fuzzy-C-Means Clustering

- FCM Optimization via Lagrange Multipliers
 - Keeping d_{jk} fixed, i.e., cluster-means fixed, solve for Memberships and Lagrange Multipliers

$$\frac{\partial L(\{u_{jk}\}, \{\lambda_j\}, \{c_k\})}{\partial u_{jk}} = 0 = qu_{jk}^{q-1}(y_j - c_k)^2 + \lambda_j$$
$$\frac{\partial L(\{u_{jk}\}, \{\lambda_j\}, \{c_k\})}{\partial \lambda_j} = 0 = \sum_k u_{jk} - 1$$

 $u_{jk} = \left(\frac{-\lambda_j}{qd_{jk}} \right)^{\frac{1}{q-1}}$

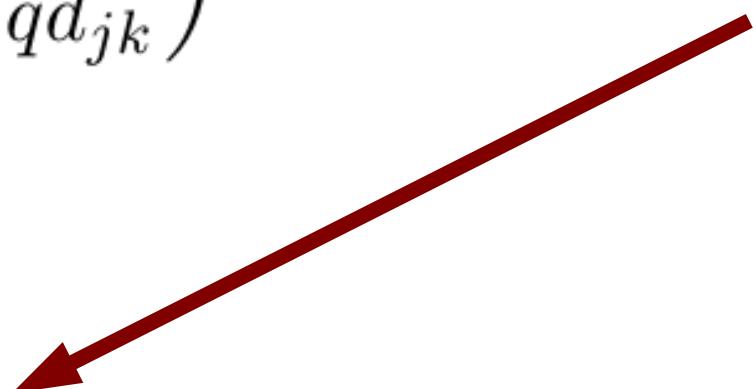
$$L(\{u_{jk}\}, \{c_k\}, \{\lambda_j\}) := \sum_{j=1}^N \sum_{k=1}^K u_{jk}^q (y_j - c_k)^2 + \lambda_j \left(\sum_{k=1}^K u_{jk} - 1 \right)$$

Fuzzy-C-Means Clustering

- FCM Optimization via Lagrange Multipliers
 - Keeping d_{jk} fixed, i.e., cluster-means fixed, solve for Memberships and Lagrange Multipliers

$$u_{jk} = \left(\frac{-\lambda_j}{qd_{jk}} \right)^{\frac{1}{q-1}}$$

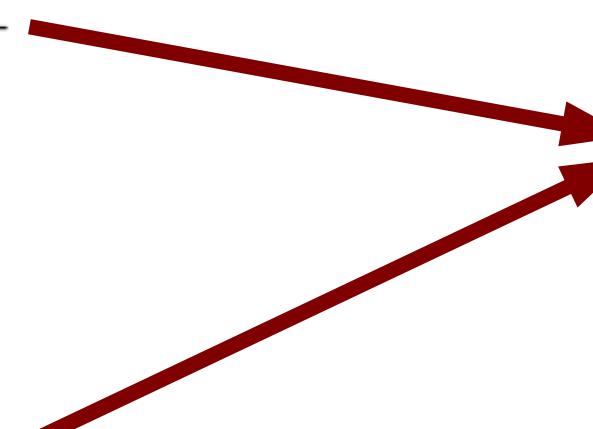
$$0 = \sum_k u_{jk} - 1$$



$$\sum_k \left(\frac{-\lambda_j}{qd_{jk}} \right)^{\frac{1}{q-1}} = 1 \implies (-\lambda_j)^{\frac{1}{q-1}} = \frac{1}{\sum_k \left(\frac{1}{qd_{jk}} \right)^{\frac{1}{q-1}}}$$

Fuzzy-C-Means Clustering

- FCM Optimization via Lagrange Multipliers
 - Keeping d_{jk} fixed, i.e., cluster-means fixed, solve for Memberships and Lagrange Multipliers
 - Substitute lambda value into the expression for u_{jk}

$$(-\lambda_j)^{\frac{1}{q-1}} = \frac{1}{\sum_k \left(\frac{1}{qd_{jk}}\right)^{\frac{1}{q-1}}} \quad u_{jk} = \frac{\left(\frac{1}{d_{jk}}\right)^{\frac{1}{q-1}}}{\sum_k \left(\frac{1}{d_{jk}}\right)^{\frac{1}{q-1}}}$$
$$u_{jk} = \left(\frac{-\lambda_j}{qd_{jk}}\right)^{\frac{1}{q-1}}$$


Fuzzy-C-Means Clustering

- FCM Optimization via Lagrange Multipliers

- Update for memberships

- Positivity constraint met
 - Sum-to-1 constraint met

- Interpretation

- If data point j has a larger distance d_{jk} to cluster mean c_k , then its membership u_{jk} to cluster k ... ???
 - If the user-defined parameter $q \rightarrow \infty$, then $\forall j, \forall k$ memberships u_{jk} are ... ???
 - If the user-defined parameter $q \rightarrow 1$, then membership ... ???

$$u_{jk} = \frac{\left(\frac{1}{d_{jk}}\right)^{\frac{1}{q-1}}}{\sum_k \left(\frac{1}{d_{jk}}\right)^{\frac{1}{q-1}}}$$

Fuzzy-C-Means Clustering

- FCM Optimization via Lagrange Multipliers

- Update for memberships

- Positivity constraint met
 - Sum-to-1 constraint met

$$u_{jk} = \frac{\left(\frac{1}{d_{jk}}\right)^{\frac{1}{q-1}}}{\sum_k \left(\frac{1}{d_{jk}}\right)^{\frac{1}{q-1}}}$$

- Interpretation

- If data point j has a larger distance d_{jk} to cluster mean c_k , then its membership u_{jk} to cluster k reduces
 - If the user-defined parameter $q \rightarrow \infty$, then $\forall j, \forall k$ memberships u_{jk} are constant, i.e., $u_{jk} = 1 / K$
 - If the user-defined parameter $q \rightarrow 1$, then membership $u_{jk} = 1$ iff $d_{jk} < d_{jk'}$ otherwise $u_{jk} = 0$
 - Hard / crisp clustering (equivalent to K means)

Fuzzy-C-Means Clustering

- FCM Optimization via Lagrange Multipliers
 - Keeping memberships and Lagrange multiplier fixed, solve for class means

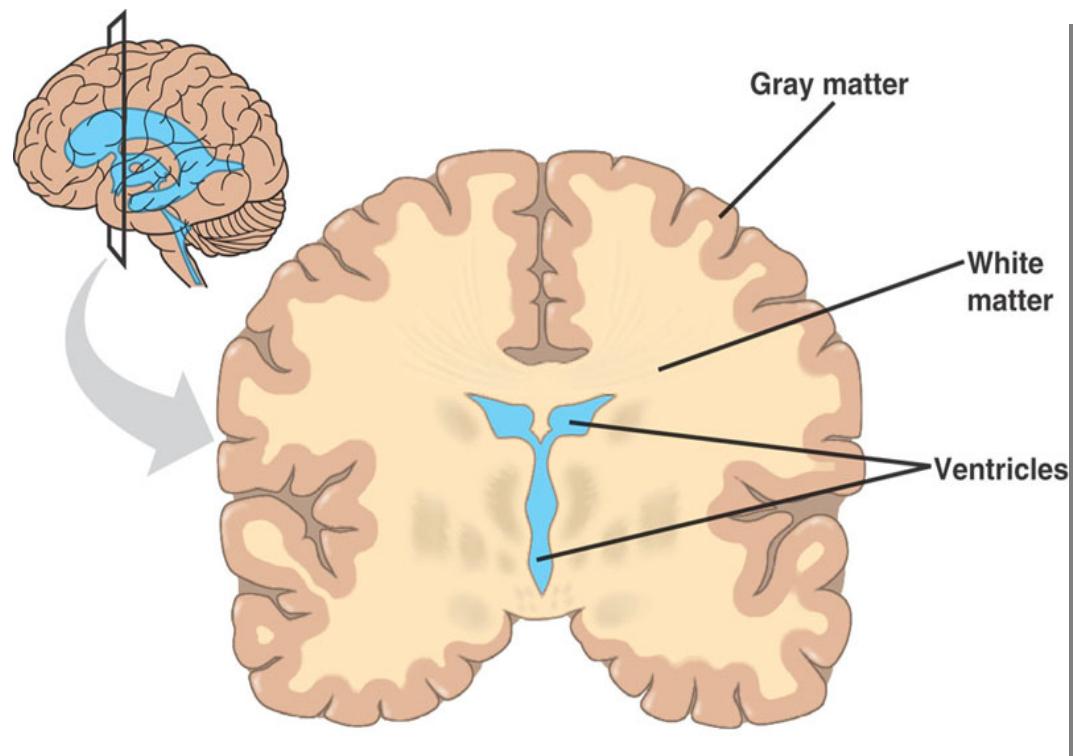
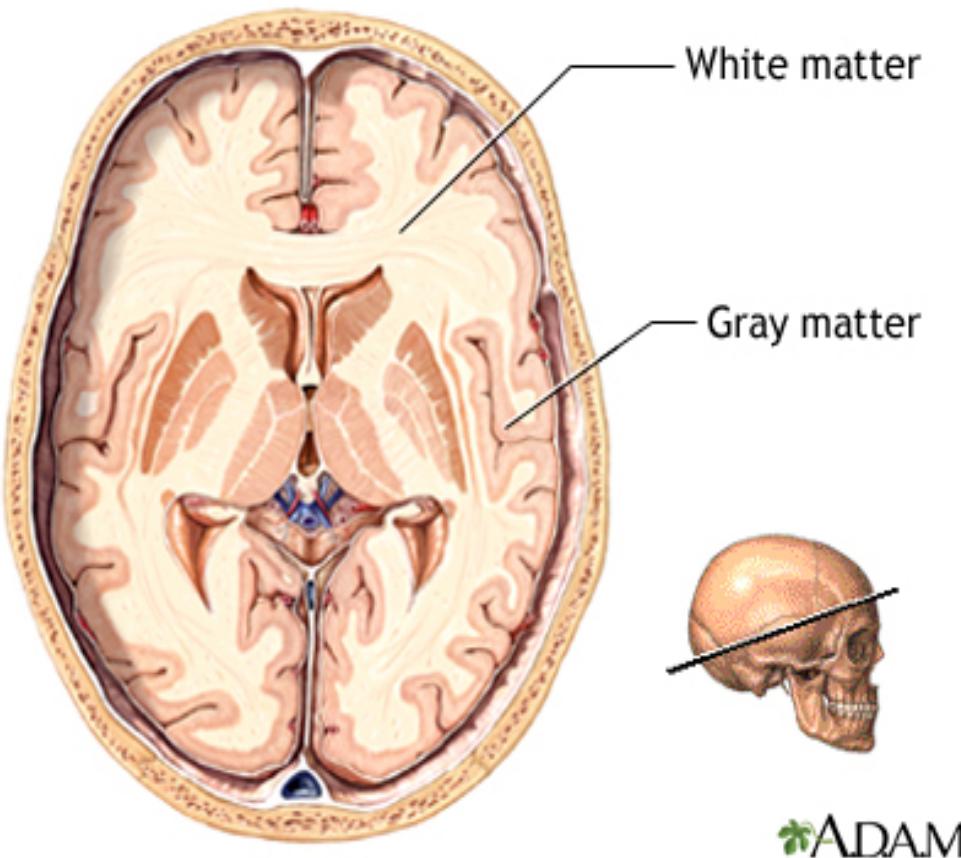
$$\frac{\partial L(\{u_{jk}\}, \{\lambda_j\}, \{c_k\})}{\partial c_k} = 0 = \sum_{j=1}^N u_{jk}^q 2(y_j - c_k) \longrightarrow c_k = \frac{\sum_{j=1}^N u_{jk}^q y_j}{\sum_{j=1}^N u_{jk}^q}$$

- Interpretation
 - Class mean is weighted average of data points
 - Weight larger \leftarrow membership larger
- What happens with memberships are binary ?

$$L(\{u_{jk}\}, \{c_k\}, \{\lambda_j\}) := \sum_{j=1}^N \sum_{k=1}^K u_{jk}^q (y_j - c_k)^2 + \lambda_j \left(\sum_{k=1}^K u_{jk} - 1 \right)$$

Tissue Segmentation in Brain MRI

- Human brain tissues

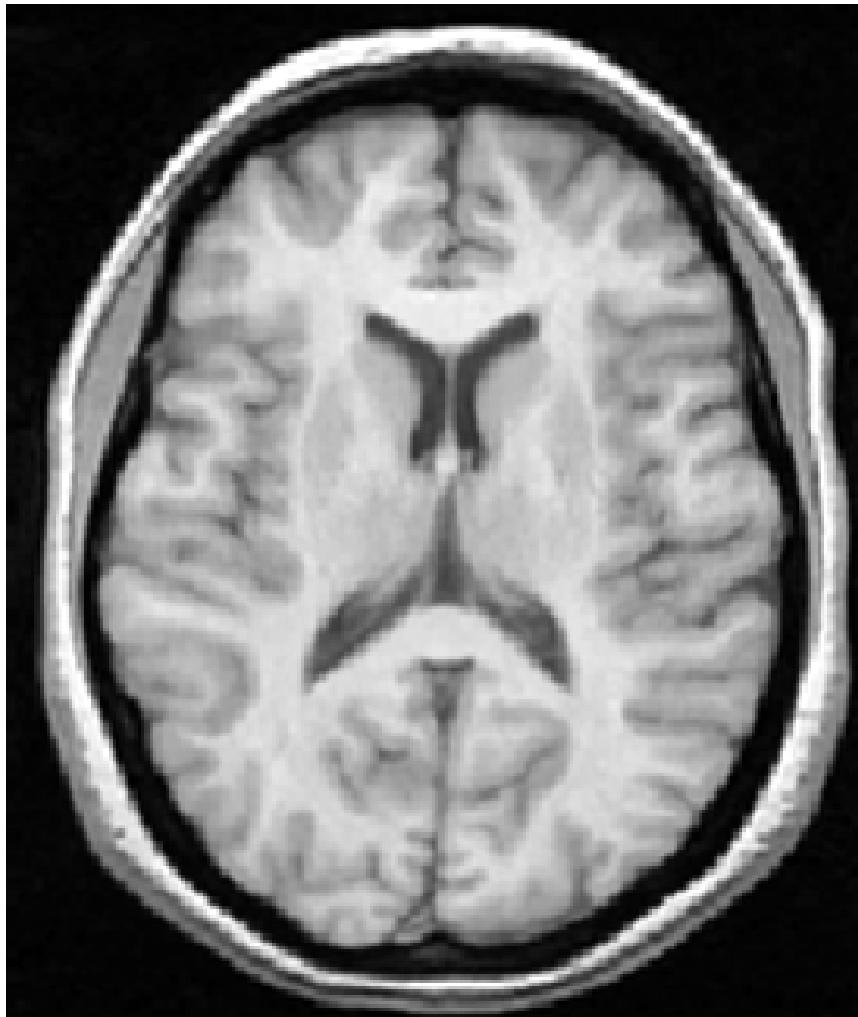


tgesicsebiology.weebly.com/nervous-system.html

www.nlm.nih.gov/medlineplus/ency/imagepages/18117.htm

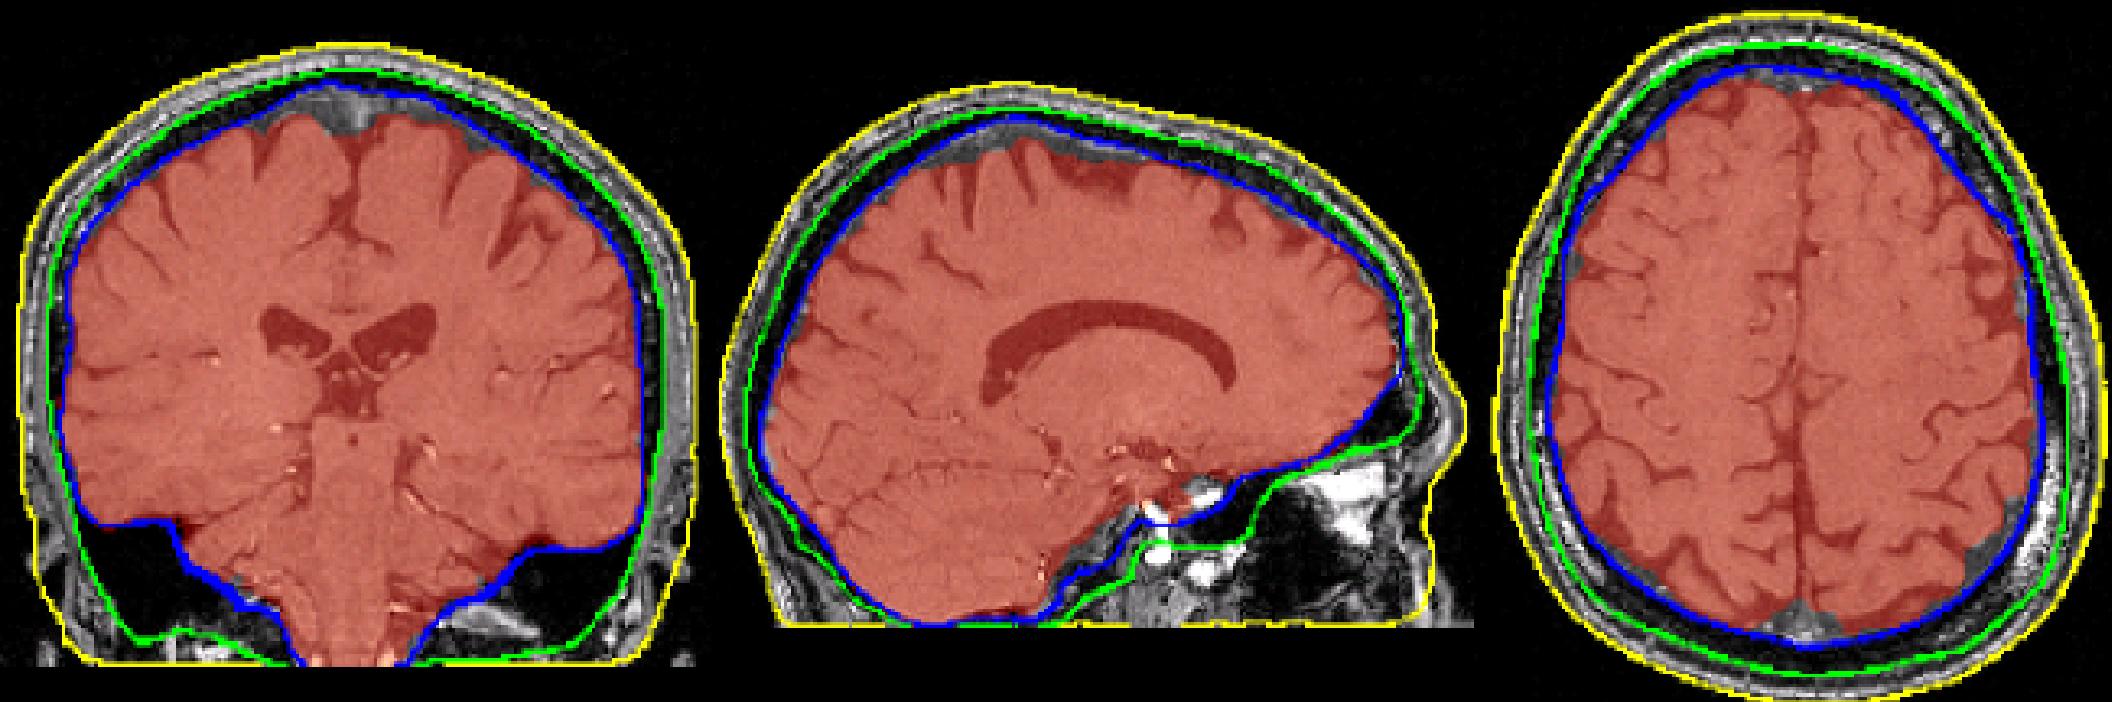
Tissue Segmentation in Brain MRI

- 2 stages
 - (1) Brain extraction from head MR image



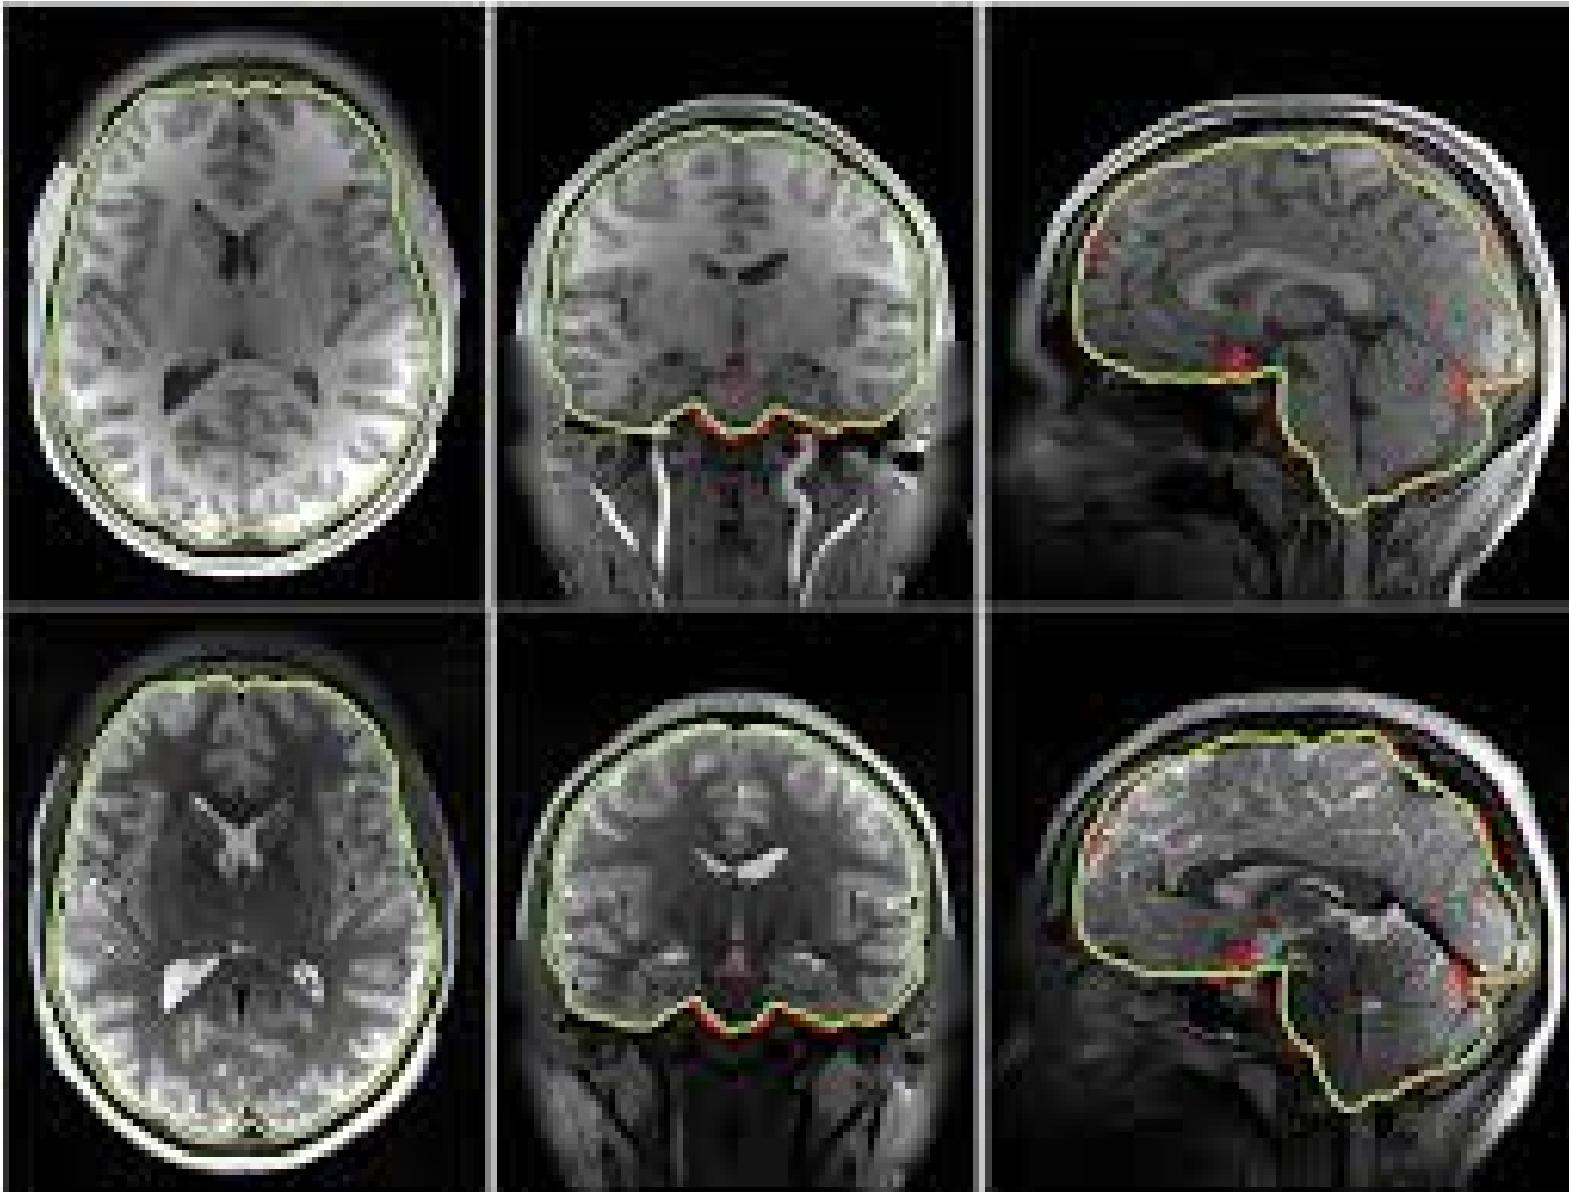
Tissue Segmentation in Brain MRI

- 2 stages
 - (1) Brain extraction from head MR image



Tissue Segmentation in Brain MRI

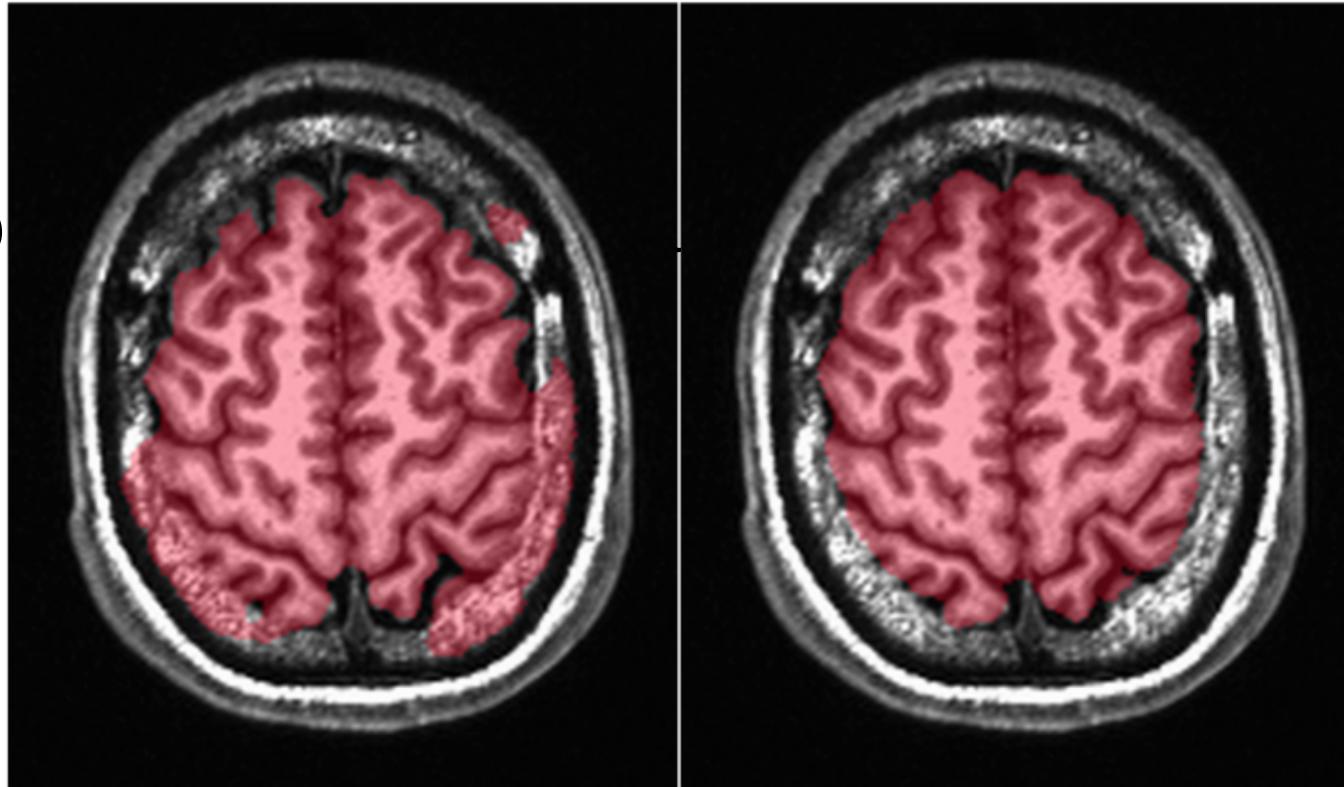
- 2 stages
 - (1) Brain extraction from head MR image



Tissue Seg.

- 2 stages
 - (1) Brain extraction

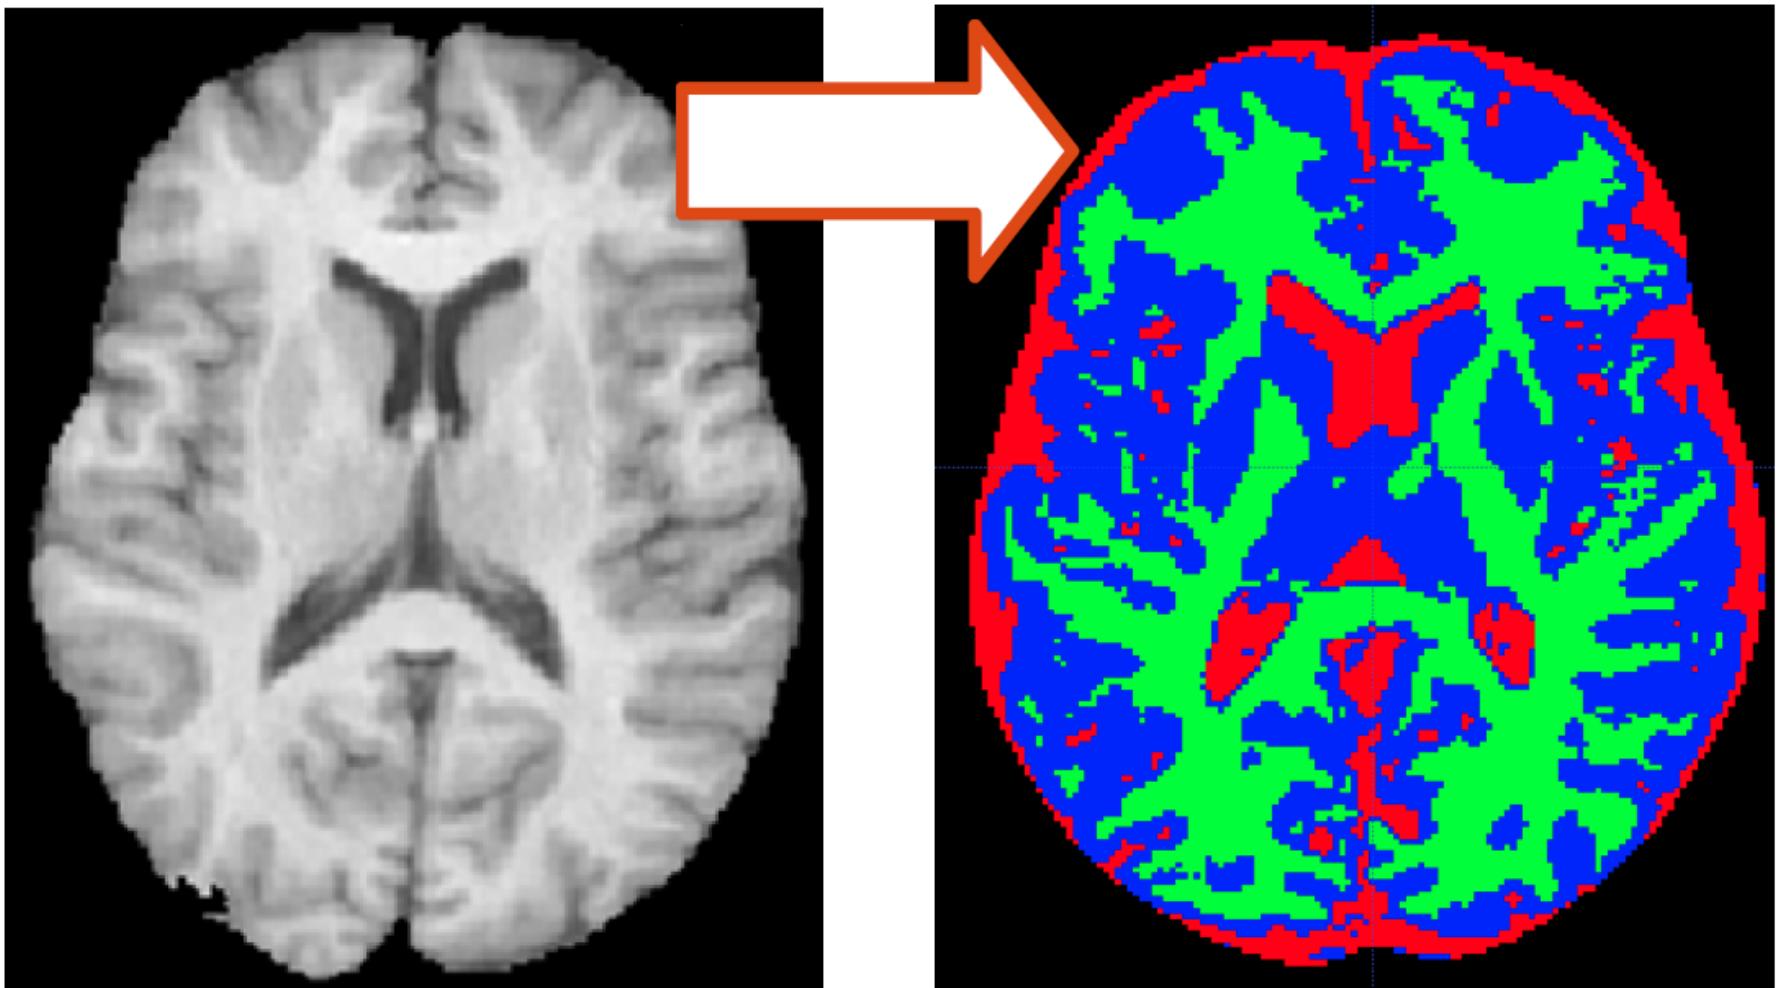
Axial view



Coronal view

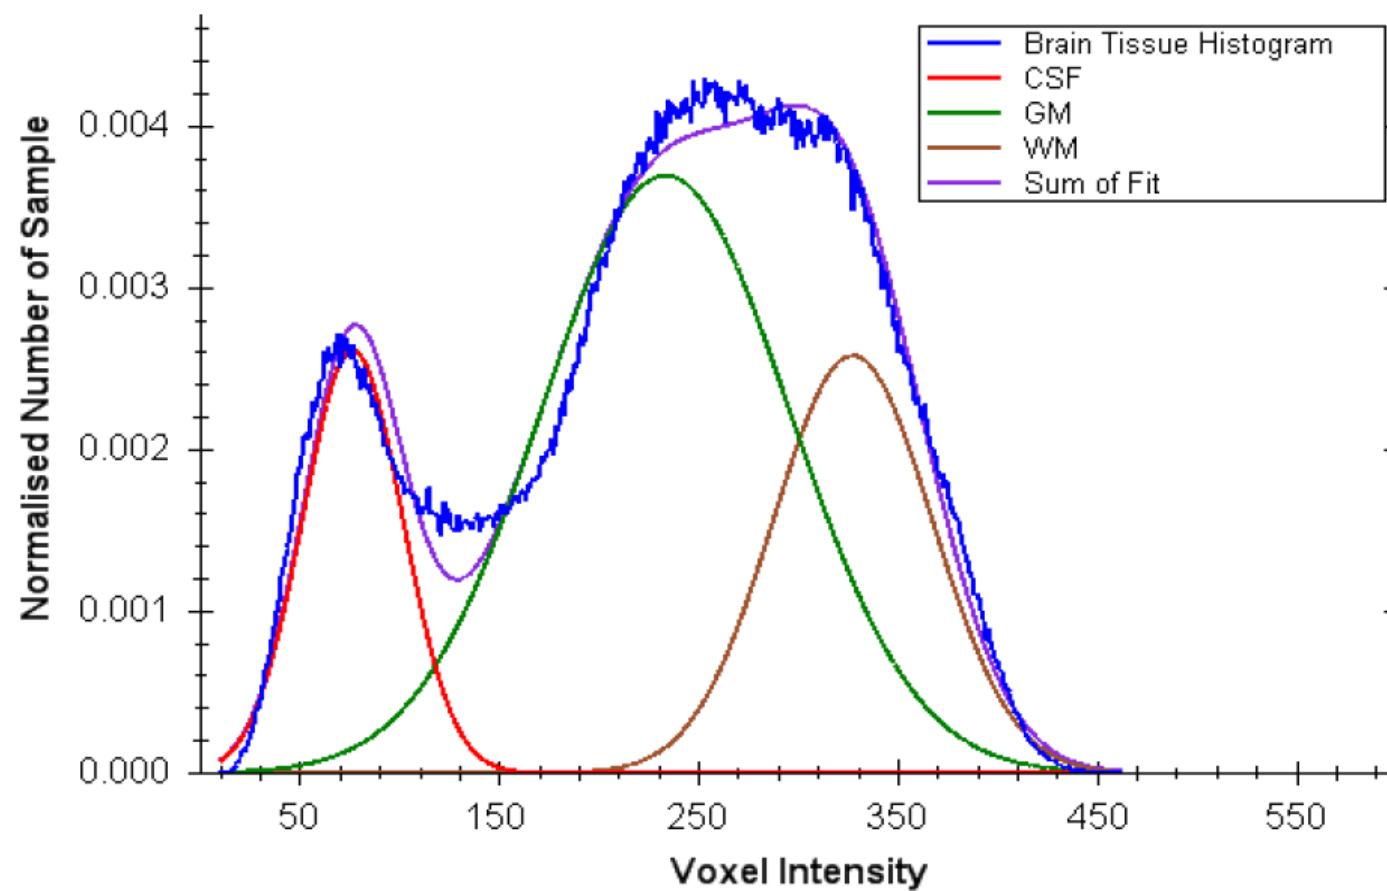
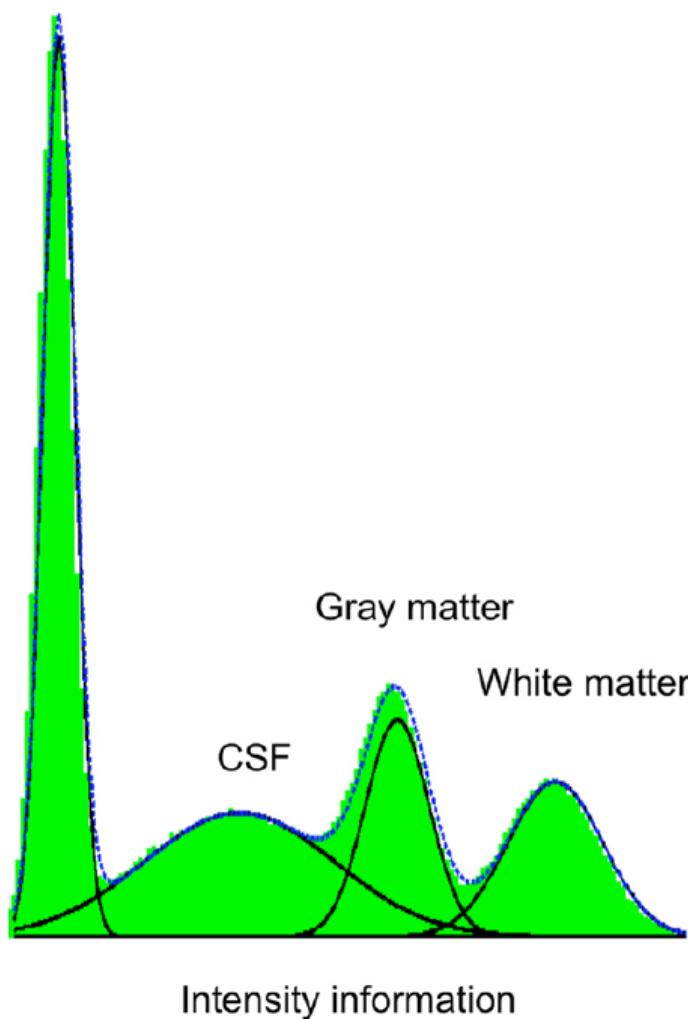
Tissue Segmentation in Brain MRI

- 2 stages
 - (1) Brain extraction from head MR image
 - (2) Tissue segmentation within brain



Tissue Segmentation in Brain MRI

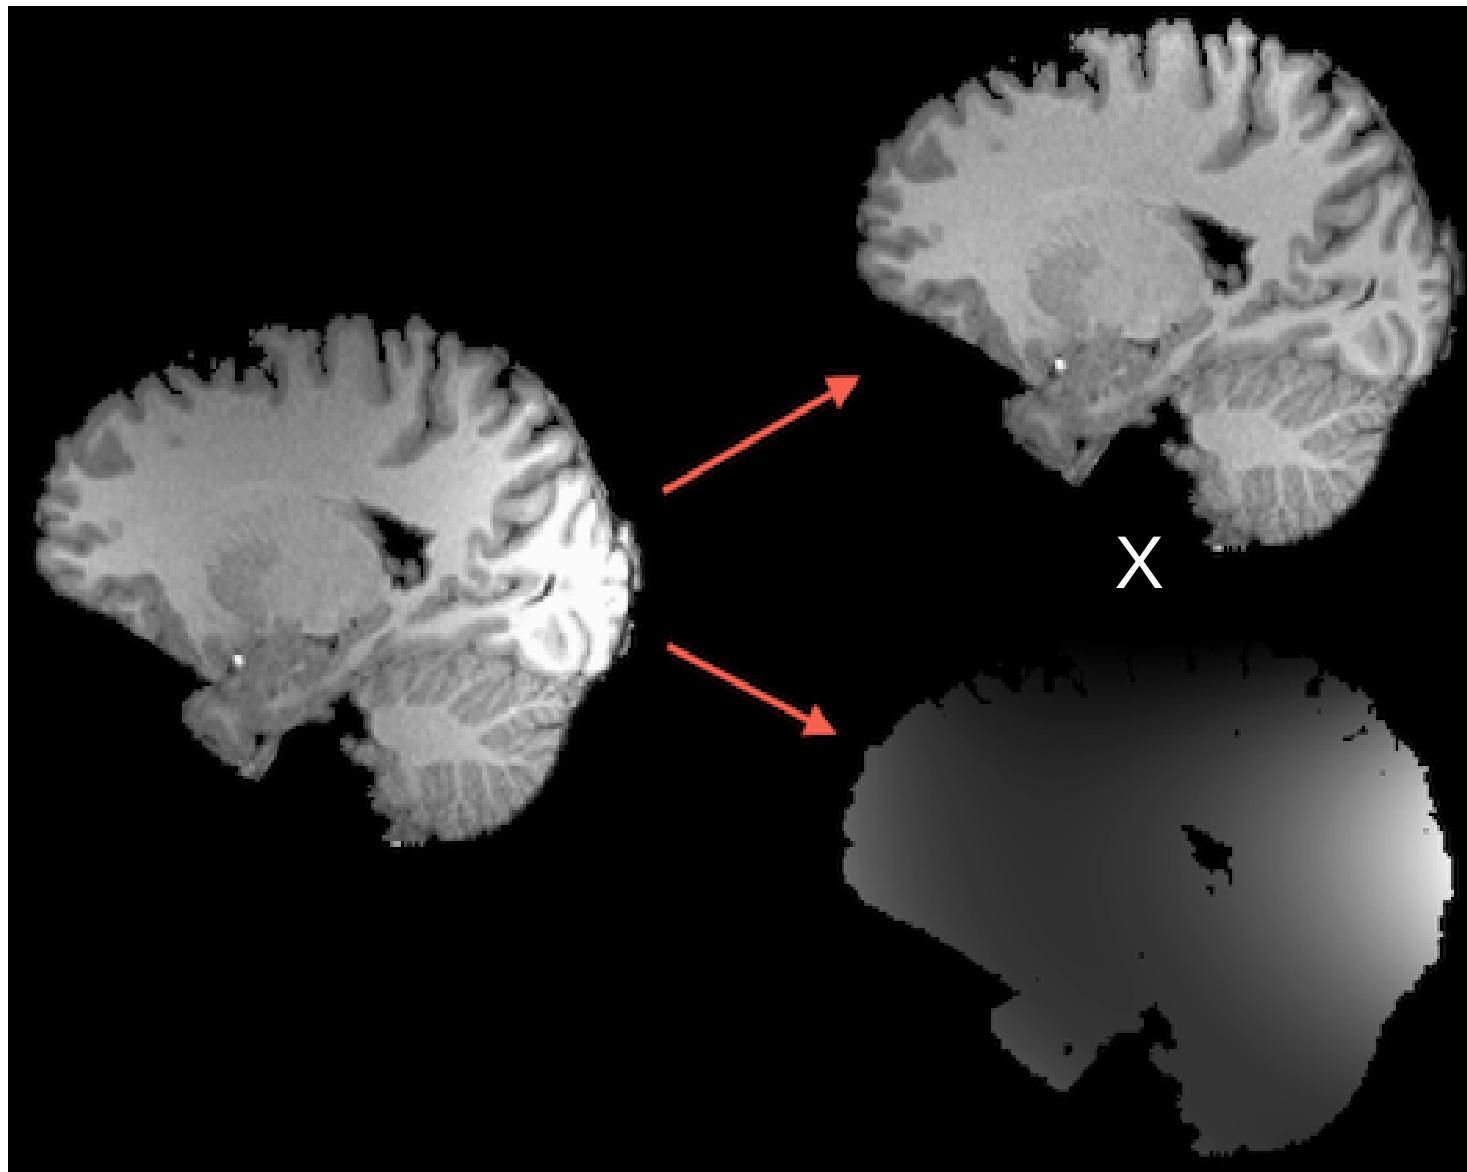
- PDF of intensities in T1-weighted brain MRI



Tissue Segmentation in Brain MRI

- Bias / inhomogeneity field in brain MRI

- Imperfect static b_1 magnetic field (RF pulse)
- Multiplicative corruption on intensities
- Spatially smooth



Vignetting

- Reduction of an image's brightness or saturation at the periphery compared to the image center
 - Often unintended / undesired ← cam. settings / lens limits
 - Sometimes introduced to draw attention to frame center



FCM Seg. + Bias-Field Correction

- MR magnitude image with N voxels
- At voxel i, **intensity** is x_i (unknown)
- At voxel i, **inhomogeneity / bias** is b_i (unknown)
 - Assumption: Varies smoothly over space
- Consider **noise** to be i.i.d. additive Gaussian
 - Approximation valid for voxels with high-SNR
 - Otherwise, can work with complex values
- At voxel i, **observed** intensity is $y_i := x_i b_i + \eta_i$
 - $\eta_i \sim G(0, \sigma^2)$

FCM Seg. + Bias-Field Correction

- 2 more model assumptions
 - Biologically reasonable
 - Image comprises K classes (known K)
 - Each class k has a constant MR (uncorrupted) intensity = c_k
- Modeling assumptions lead to :
 - At voxel i,
if its neighbor j belongs to class k,
then, $x_j = c_k$ (constant intensity of tissue)
 $b_j \approx b_i$ (spatial smoothness of bias)
and, thus, $x_j b_j \approx c_k b_i$

FCM Seg. + Bias-Field Correction

- Strategy
 - At voxel j in neighborhood of voxel i ,
penalize difference between
 - (i) observed neighborhood intensities y_j and
 - (ii) intensities predicted by model $x_j b_j$
 - If neighborhood voxel j belongs to class k ,
then, predicted intensity = $x_j b_j \approx c_k b_i$
 - Because we don't know which class voxel j is in,
so penalize differences for all classes
weighted by membership u_{jk}^q in that class
 - Because bias-constancy is local,
weight differences based on distance between voxels i, j
 - Weight $w_{ij} \rightarrow 0$ when distance > threshold
 - Implement as a neighborhood “mask”. Weights sum to 1.

FCM Seg. + Bias-Field Correction

- Objective function, per voxel i ,

$$J_i := \sum_{j=1}^N w_{ij} \sum_{k=1}^K (u_{jk}^q (y_j - c_k b_i)^2)$$

- Constraints

- Membership, weights

$$\forall j; \sum_{k=1}^K u_{jk} := 1 \quad \sum_{j=1}^N w_{ij} = 1$$

- Objective function

$$J := \sum_{i=1}^N J_i := \sum_{j=1}^N \sum_{k=1}^K u_{jk}^q \left(\sum_{i=1}^N w_{ij} (y_j - c_k b_i)^2 \right)$$

- Compare with standard FCM:

$$\sum_{j=1}^N \sum_{k=1}^K u_{jk}^q (y_j - c_k)^2$$

FCM Seg. + Bias-Field Correction

- Denote

$$d_{kj} := \sum_{i=1}^N w_{ij} (y_j - c_k b_i)^2$$

- d_{kj} depends only on class-number k & voxel-number j

- Rewritten optimization problem

$$\min_{\{u_{jk}\}, \{c_k\}, \{b_i\}} J(\{u_{jk}\}, \{c_k\}, \{b_i\}) := \sum_{j=1}^N \sum_{k=1}^K u_{jk}^q d_{kj}$$

constraints : $\forall j, \sum_k u_{jk} = 1$

FCM Seg. + Bias-Field Correction

- Optimization via Lagrange Multipliers
 - Lagrangian

$$L(\{u_{jk}\}, \{c_k\}, \{b_i\}, \{\lambda_j\}) := \sum_{j=1}^N \sum_{k=1}^K u_{jk}^q d_{kj} + \sum_{j=1}^N \lambda_j \left(1 - \sum_k u_{jk} \right)$$

$$d_{kj} := \sum_{i=1}^N w_{ij} (y_j - c_k b_i)^2$$

FCM Seg. + Bias-Field Correction

- Optimization via Lagrange Multipliers
 - Keeping class means and bias fixed,
solve for **memberships and Lagrange multipliers**

$$\frac{\partial L(\{u_{jk}\}, \{\lambda_j\}, \{c_k\})}{\partial u_{jk}} = 0 = qu_{jk}^{q-1}d_{kj} - \lambda_i$$

$$\frac{dL(\{u_{jk}\}, \{\lambda_j\}, \{c_k\})}{d\lambda_i} = 0 = 1 - \sum_k u_{jk}$$

FCM Seg. + Bias-Field Correction

- Optimization via Lagrange Multipliers
 - Keeping class means and bias fixed,
solve for **memberships and Lagrange multipliers**

$$u_{jk} = \frac{\left(\frac{1}{d_{kj}}\right)^{\frac{1}{q-1}}}{\sum_k \left(\frac{1}{d_{kj}}\right)^{\frac{1}{q-1}}}$$
$$d_{kj} := \sum_{i=1}^N w_{ij} (y_j - c_k b_i)^2$$

- Same as in FCM, except that d_{kj} has a different meaning
 - d_{kj} generalizes / smooths the squared distance over neighborhood
 - How to choose width of mask w ?
 - Smoothness of bias field b (a priori)
 - What happens when no bias ($b_i=1$) leads to choice of wide mask ($w_{ij}=\text{const.}$)?

FCM Seg. + Bias-Field Correction

- Optimization via Lagrange Multipliers
 - Keeping memberships, multipliers, and bias fixed, solve for **class means**

$$\frac{\partial L(\{u_{jk}\}, \{c_k\}, \{b_i\}, \{\lambda_j\})}{\partial c_k} = 0 = \sum_j u_{jk}^q \frac{\partial d_{kj}}{\partial c_k} = \sum_j u_{jk}^q \sum_i w_{ij} 2(y_j - c_k b_i)(-b_i)$$

- Rearranging terms :

$$c_k = \frac{\sum_j u_{jk}^q y_j \sum_i w_{ij} b_i}{\sum_j u_{jk}^q \sum_i w_{ij} b_i^2}$$

$$d_{kj} := \sum_{i=1}^N w_{ij} (y_j - c_k b_i)^2$$

$$L(\{u_{jk}\}, \{c_k\}, \{b_i\}, \{\lambda_j\}) := \sum_{j=1}^N \sum_{k=1}^K u_{jk}^q d_{kj} + \sum_{j=1}^N \lambda_j \left(1 - \sum_k u_{jk} \right)$$

FCM Seg. + Bias-Field Correction

- Optimization via Lagrange Multipliers
 - Keeping memberships, multipliers, and bias fixed, solve for **class means**

$$c_k = \frac{\sum_j u_{jk}^q y_j \sum_i w_{ij} b_i}{\sum_j u_{jk}^q \sum_i w_{ij} b_i^2}$$

- Interpretation
 - Summation over w_{ij} performs convolution with the neighborhood mask (which smooths the bias field)
 - If bias field is constant = b , then c_k = membership-weighted avg of ratios (y_j / b)
 - If bias field varies slowly over neighborhood mask w + memberships are binary, then $c_k \approx \sum_{j \in C^k} y_j b_j / \sum_{j \in C^k} b_j^2$

FCM Seg. + Bias-Field Correction

- Optimization via Lagrange Multipliers
 - Keeping memberships, multipliers, and class means fixed, solve for bias

$$\frac{\partial L(\{u_{jk}\}, \{c_k\}, \{b_i\}, \{\lambda_j\})}{\partial b_i} = 0 = \sum_{j=1}^N \sum_{k=1}^K u_{jk}^q \frac{\partial d_{kj}}{\partial b_i} = \sum_{j=1}^N \sum_{k=1}^K u_{jk}^q w_{ij} 2(y_j - c_k b_i) c_k$$

- Rearranging terms :

$$b_i = \frac{\sum_{j=1}^N w_{ij} y_j \sum_{k=1}^K u_{jk}^q c_k}{\sum_{j=1}^N w_{ij} \sum_{k=1}^K u_{jk}^q c_k^2}$$

$$d_{kj} := \sum_{i=1}^N w_{ij} (y_j - c_k b_i)^2$$

$$L(\{u_{jk}\}, \{c_k\}, \{b_i\}, \{\lambda_j\}) := \sum_{j=1}^N \sum_{k=1}^K u_{jk}^q d_{kj} + \sum_{j=1}^N \lambda_j \left(1 - \sum_k u_{jk} \right)$$

FCM Seg. + Bias-Field Correction

- Optimization via Lagrange Multipliers
 - Keeping memberships, multipliers, and class means fixed, solve for bias

$$b_i = \frac{\sum_{j=1}^N w_{ij} y_j \sum_{k=1}^K u_{jk}^q c_k}{\sum_{j=1}^N w_{ij} \sum_{k=1}^K u_{jk}^q c_k^2}$$

- Interpretation
 - If memberships binary, then $b_i \approx \sum_j w_{ij} y_j c_{k(j)} / \sum_j w_{ij} c_{k(j)}^2$
 - If, in addition, all neighbors are in class k, then bias b_i = spatially-weighted average ratios of (y_j / c_k)
 - Spatial smoothing
 - Helps counter noise in observed data
 - Ensures a smooth bias-field estimate

FCM Seg. + Bias-Field Correction

- Did we formulate the problem right ?
- Is the model identifiable ?
- Can class means & bias field be uniquely identified ?
 - Model: $y_i := x_i b_i + \eta_i$
 - $y_i := c_k b_i + \eta_i$, if voxel i belongs (100%) to class k
 - What happens when, say, noise is absent ?

$$d_{kj} := \sum_{i=1}^N w_{ij} (y_j - c_k b_i)^2$$

$$L(\{u_{jk}\}, \{c_k\}, \{b_i\}, \{\lambda_j\}) := \sum_{j=1}^N \sum_{k=1}^K u_{jk}^q d_{kj} + \sum_{j=1}^N \lambda_j \left(1 - \sum_k u_{jk} \right)$$

K-Means and FCM

- Limitations
 - Can't model clusters with **unequal-spread**
 - How to do that ?
 - Gaussian model (variance)
 - Can't model **non-spherical** clusters (in higher-D)
 - How to do that ?
 - Gaussian model (covariance)
 - Don't enforce **spatial smoothness** on segmentation
 - How to do that ?
 - MRF model on membership image
 - Can't model **multimodal** clusters and '**curved**' clusters
 - How to do that ?

Gauss

- Carl Friedrich Gauss (1777 – 1855)
 - “Prince of mathematicians”
 - “Greatest mathematician since antiquity”
 - Number theory, algebra, statistics, analysis, differential geometry, geophysics, electrostatics, astronomy, optics
 - Child prodigy
 - Published Arithmetical Investigations (age 21), consolidating number theory as a discipline
 - PhD in 1799 (age 22)
 - Provided a new proof of the fundamental theorem of algebra
 - Every non-zero, single-variable, **degree-n polynomial** with complex coefficients has, counted with multiplicity, exactly **n roots**
 - FFT was invented by Gauss in 1805 (age 28)
 - Cooley and Tukey reinvented it in 1965 and popularized it



Gauss



- Carl Friedrich Gauss (1777 – 1855)
 - Refused to publish “unripe” work (“few, but ripe”)
 - Personal diary indicates discoveries before contemporaries
 - Declined to provide intuition behind his elegant proofs
 - Erased all traces of discovery process
 - Wanted proofs to look 'magical'
 - All analysis must be suppressed for the sake of brevity

Bernhard Riemann

- Mathematician
 - Analysis, number theory, differential geometry
 - Advisor: Gauss
- As a child, Riemann showed exceptional math skills, but was shy of public speaking
 - In school, he often outstripped his instructor's knowledge
- Father was a pastor. In high school, he studied Bible intensively, but was often distracted by mathematics
- Went to university to get degree in Theology. Met Gauss, who suggested Riemann to get into math

Riemann

- Clay math institute's Millennium Prize Problems includes 'Riemann hypothesis'



Number theory [edit]

- Riemann-von Mangoldt formula
- Riemann hypothesis
 - Generalized Riemann hypothesis
 - Grand Riemann hypothesis
 - Riemann hypothesis for curves over finite fields
- Riemann theta function
- Riemann Xi function
- Riemann zeta function
- Riemann-Siegel formula
- Riemann-Siegel theta function

Riemannian [edit]

- Pseudo-Riemannian manifold
- Riemannian bundle metric
- Riemannian circle
- **Riemannian cobordism**
- Riemannian connection
- Riemannian connection on a surface
- Riemannian cubic
- Riemannian cubic polynomials
- Riemannian foliation
- Riemannian geometry
 - Fundamental theorem of Riemannian geometry
- **Riemannian graph**
- **Riemannian group**
- Riemannian holonomy
- Riemannian manifold also called Riemannian space
- Riemannian metric tensor
- Riemannian Penrose inequality
- **Riemannian polyhedron**
- Riemannian singular value decomposition
- Riemannian submanifold
- Riemannian submersion
- Riemannian volume form
- **Riemannian wavefield extrapolation**
- Sub-Riemannian manifold
- Riemannian symmetric space

Riemann's [edit]

- Riemann's differential equation
- Riemann's existence theorem
- Riemann's explicit formula
- Riemann's minimal surface
- Riemann's theorem on removable singularities

Free Riemann gas also called primon gas

- Riemann bilinear relations
- Riemann–Cartan geometry
- Riemann conditions
- Riemann curvature tensor also called Riemann tensor
- Riemann form
- Riemann function
- Riemann–Hurwitz formula
- Riemann invariant
- Riemann matrix
- Riemann operator
- Riemann problem
- Riemann–Silberstein vector
- Riemann singularity theorem
- Riemann surface
 - Compact Riemann surface
- Riemann tensor (general relativity)
- The tangential Cauchy–Riemann complex
- Zariski–Riemann space

Analysis [edit]

- Riemann integral
 - Generalized Riemann integral
 - Riemann multiple integral
- Cauchy–Riemann equations
- Riemann mapping theorem
 - Measurable Riemann mapping theorem
- Riemann solver
- Riemann sphere
- Riemann–Hilbert correspondence
- Riemann–Hilbert problem
- Riemann–Lebesgue lemma
- Riemann–Liouville differintegral
- Riemann–Roch theorem
 - Arithmetic Riemann–Roch theorem
 - Riemann–Roch theorem for smooth manifolds
 - Grothendieck–Hirzebruch–Riemann–Roch theorem
 - Hirzebruch–Riemann–Roch theorem
- Riemann–Stieltjes integral
- Riemann series theorem
- Riemann sum

Mixture Model

- What is a mixture model ?
 - Probabilistic model
 - Represents a general multi-modal PDF
 - Useful for observed data that is typically derived from multiple subgroups within a population
- Represent PDF $p(x)$ as a convex combination of standard PDFs $p_k(x)$

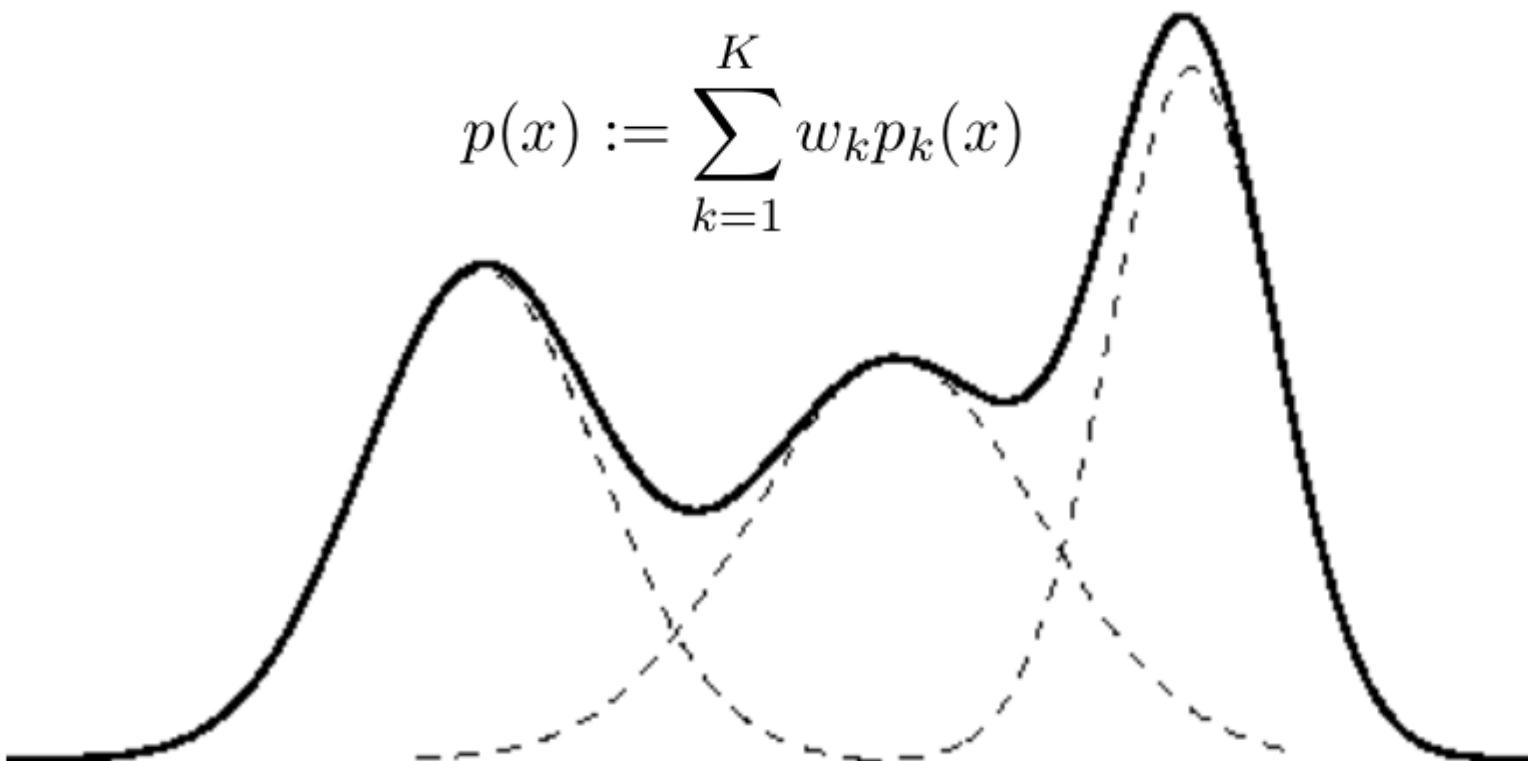
$$p(x) := \sum_{k=1}^K w_k p_k(x)$$

- Weights w_k are positive (why ?) and sum to 1 (why ?)

Mixture Model

- How to interpret weights ?
 - Larger weight $w_k \rightarrow$ data more likely to be derived from k-th Gaussian
- How to generate data from such a model ?

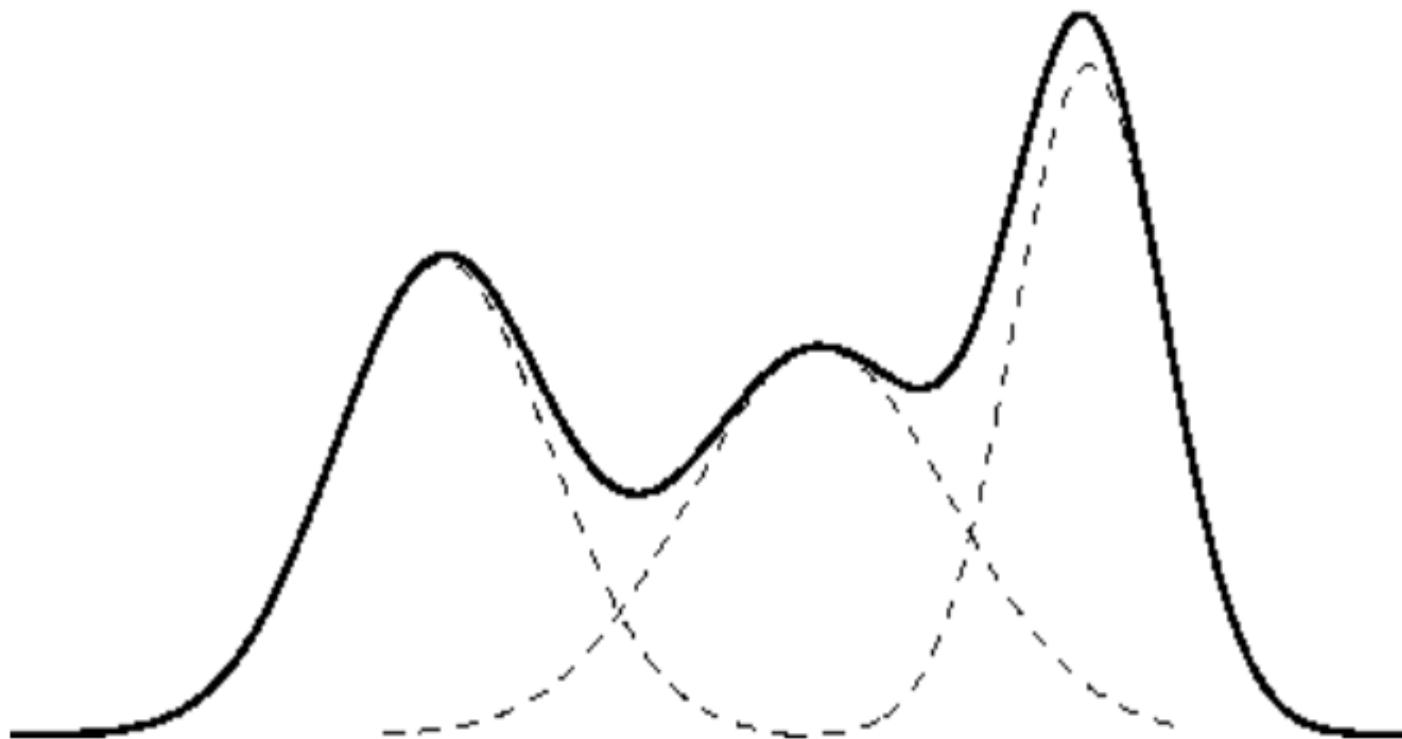
$$p(x) := \sum_{k=1}^K w_k p_k(x)$$



Mixture Model

- When number of PDFs K is finite
 - Finite mixture model
- When each PDF is a Gaussian
 - Gaussian mixture model (GMM)
- Power of the GMM model
 - Can model unimodal / multimodal / curved / complex PDFs

$$p(x) = \sum_{k=1}^K w_k G(x; \mu_k, C_k)$$



Mixture Model

- What happens when we fit a Gaussian to 2D data that lies on a circle ? semi circle ?

Gaussian Mixture Model

- Fitting a GMM to data
 - **Given:** Data $y := \{ y_n \}_{n=1,\dots,N}$
 - Each observation y_n is drawn independently of the others
 - **Goal:** Fit a GMM with K components
 - $P(x) := \sum_{k=1}^K w_k G(x; \mu_k, C_k)$
 - *What are the ways in which datum 'x' could be generated ?*
 - Estimate weights w_k , means μ_k , covariances C_k
 - **Strategy**
 - Maximum-likelihood estimation
 - Parameters $\theta = \{w_k, \mu_k, C_k\}_{k=1}^K$
 - What is the likelihood function ?

Gaussian Mixture Model

- Fitting a GMM to data

- Likelihood function

$$L(\theta|y) := P(y|\theta) := \prod_{n=1}^N P(y_n|\theta) := \prod_{n=1}^N \left(\sum_{k=1}^K w_k G(y_n; \mu_k, C_k) \right)$$

- Maximizing the log-likelihood

$$\max_{\theta} L(\theta|y) = \max_{\theta} \sum_{n=1}^N \log \left(\sum_{k=1}^K w_k G(y_n; \mu_k, C_k) \right)$$

- Log sum can't be simplified further !

Gaussian Mixture Model

- Fitting a GMM to data
 - (1) No closed-form updates possible
 - (2) Gradient-ascent optimization
 - Converges very slowly in practice
 - (3) Is there a better alternative that alleviates both above problems ?
 - Yes
 - **Expectation maximization (EM)** optimization algorithm

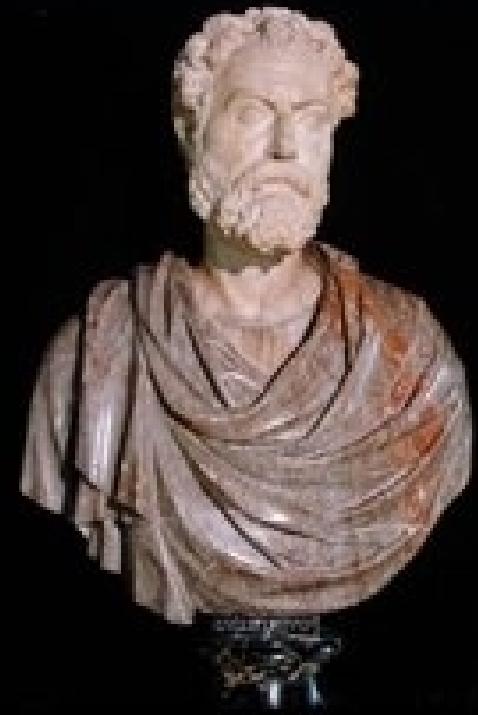


**‘There is nothing
more practical than
a good theory’.**

*Phrase attributed to Kurt Lewin, German-American psychologist, known as one of the modern pioneers of social, organizational, and applied psychology.

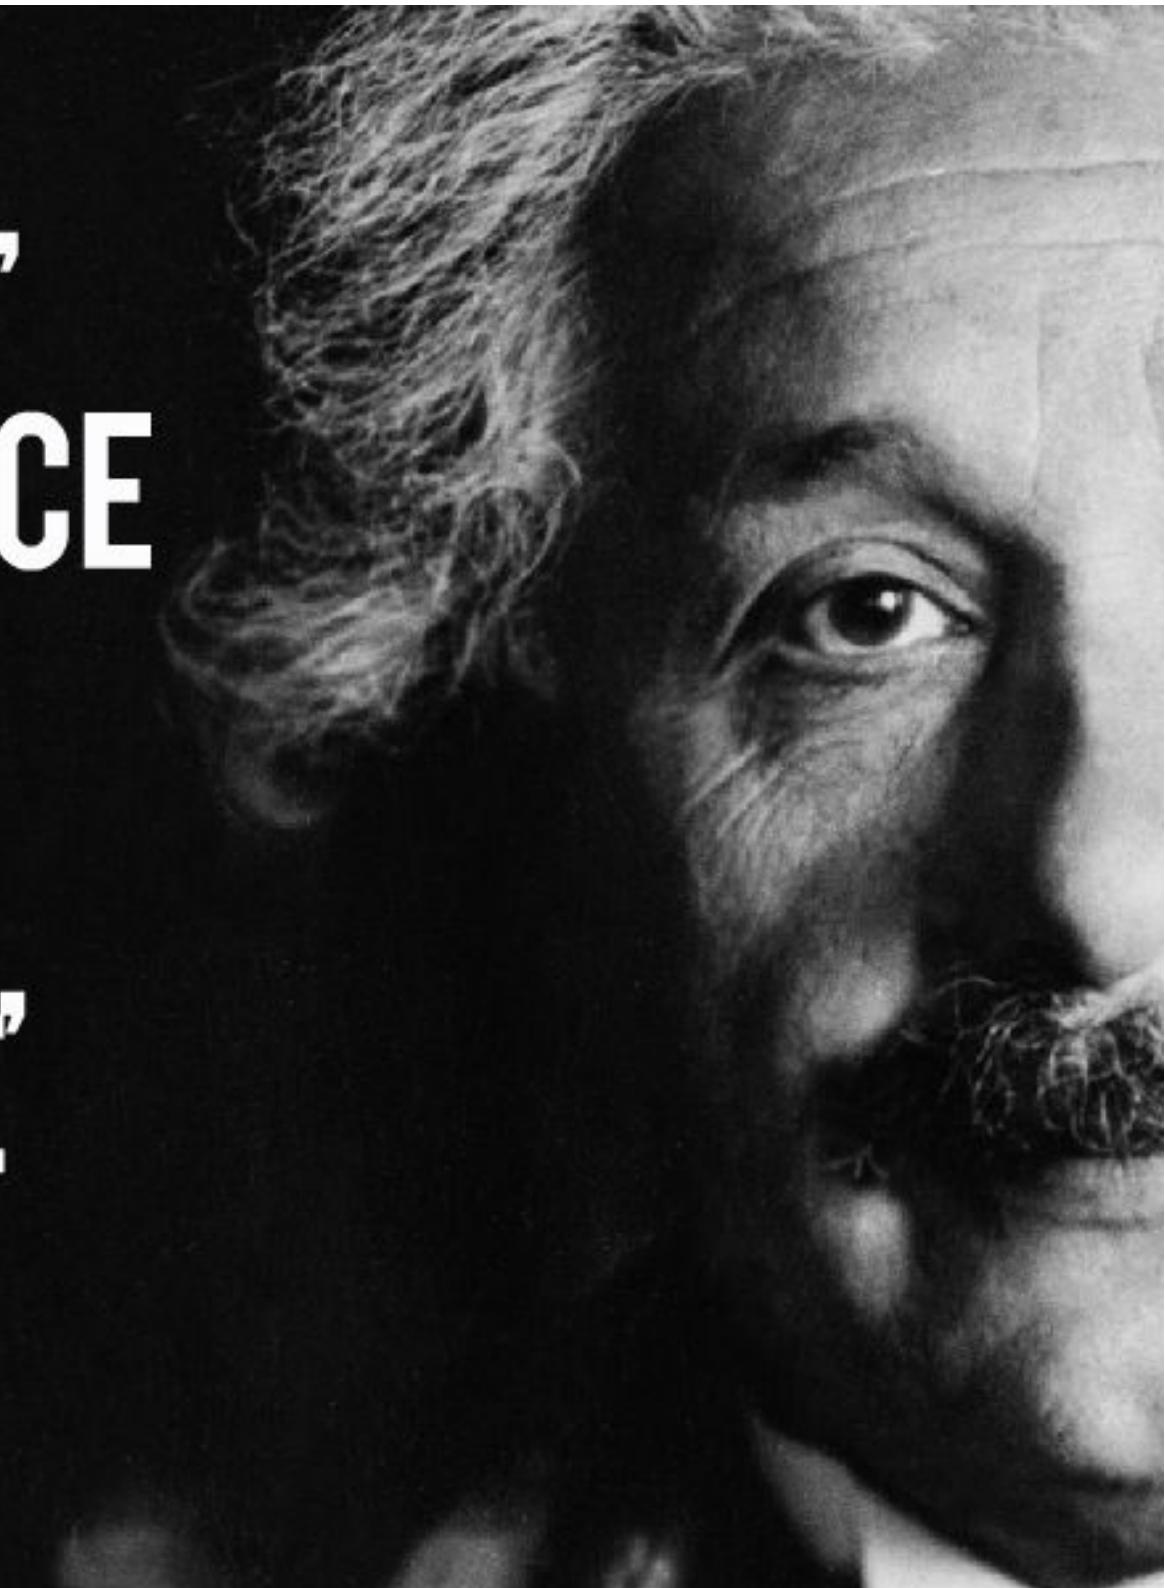
Practice without theory is
more valuable than
a theory without
practice

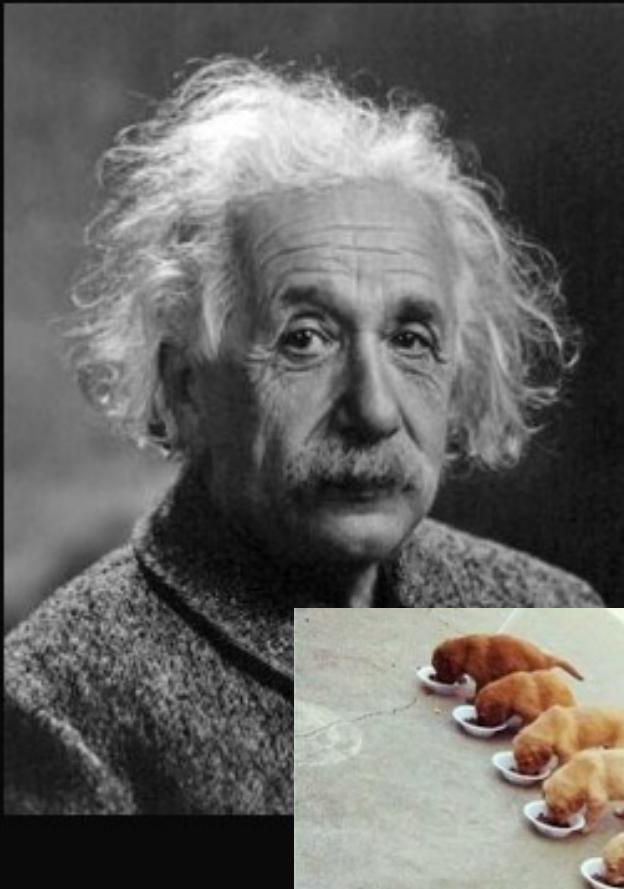
~ Marcus Quintilianus ~



**“IN THEORY,
THEORY
AND PRACTICE
ARE THE
SAME. IN
PRACTICE,
THEY ARE
NOT.”**

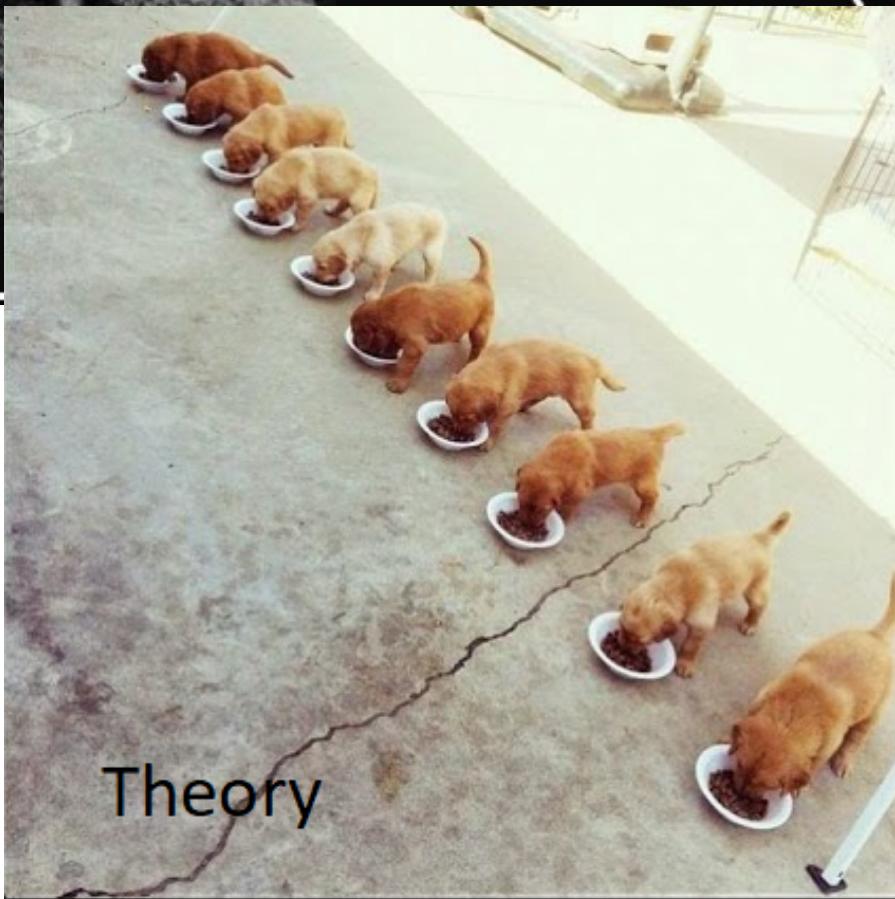
Albert Einstein





As far as the laws of mathematics refer to reality, they are not certain; and as far as they are certain, they do not refer to reality.

(Albert Einstein)



**“Theory without
practice is empty.
Practice without
theory is blind.”**

J.R. Kidd

Statistical / Machine



Theory is when you know everything but nothing works.

Practice is when everything works but no one knows why.

In our lab, theory and practice are combined: nothing works and no one knows why.

THEORY

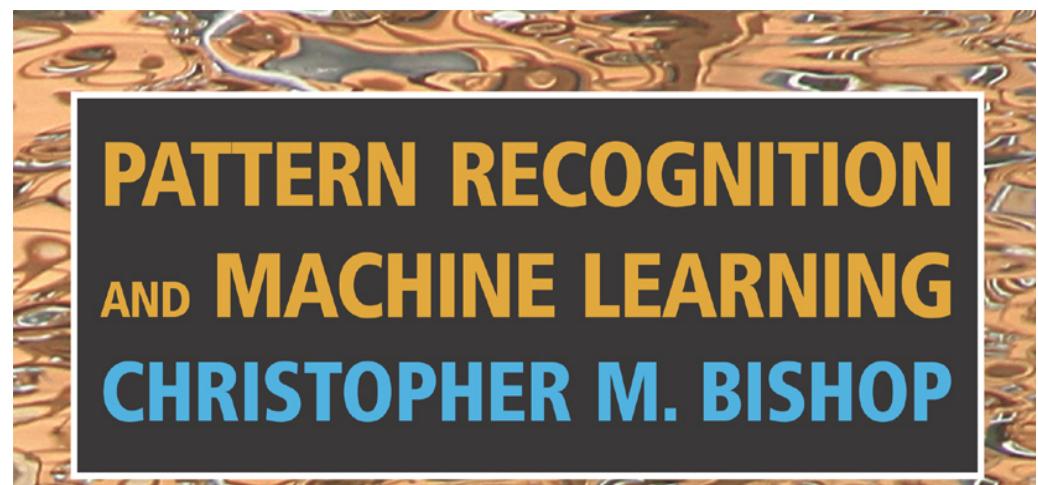
FUNDAMENTAL OR ABSTRACT
PRINCIPLES UNDERLYING
A SCIENCE OR AN ART.

AFTER ALL THOSE YEARS
OF THEORY...



EM Optimization

- Observed random variable $Y := \{ Y_n \}_{n=1,\dots,N}$
 - Leads to data $y := \{ y_n \}_{n=1,\dots,N}$
- Parameters: θ that needs to be estimated
- Hidden / latent random variable X
 - “Data” x (corresponding to instances of X) are missing
- Complete Data
 - = combination of observed + missing data
 - = $(x, y) = \{ (x_n, y_n) \}_{n=1,\dots,N}$



EM Optimization

- Key Idea
 - Hidden variable X is designed (creatively) s.t.
if missing data x were available,
then, it would become easy to solve for parameters
- Example : Segmentation
 - Design **hidden variable = cluster label** corresponding to each data point
 - Let x_n be label associated with data point y_n s.t.
 x_n = cluster label that y_n belongs to
 - If we **know cluster labels** (data points that belong to cluster k), then it is easy to find optimal values of μ_k , C_k
 - If we **know parameters** w_k , μ_k , C_k , then ...

EM Optimization

- EM optimization performs ML estimation
 - ML estimate is :
$$\max_{\theta} P(y|\theta) = \max_{\theta} \int_x P(y, x|\theta) dx$$
 - EM introduces hidden variable, then marginalizes over it
- What is an alternative ?
 - Treat 'x' has parameter
 - Optimize for a value of 'x'
 - Why is EM preferred ?

EM Optimization

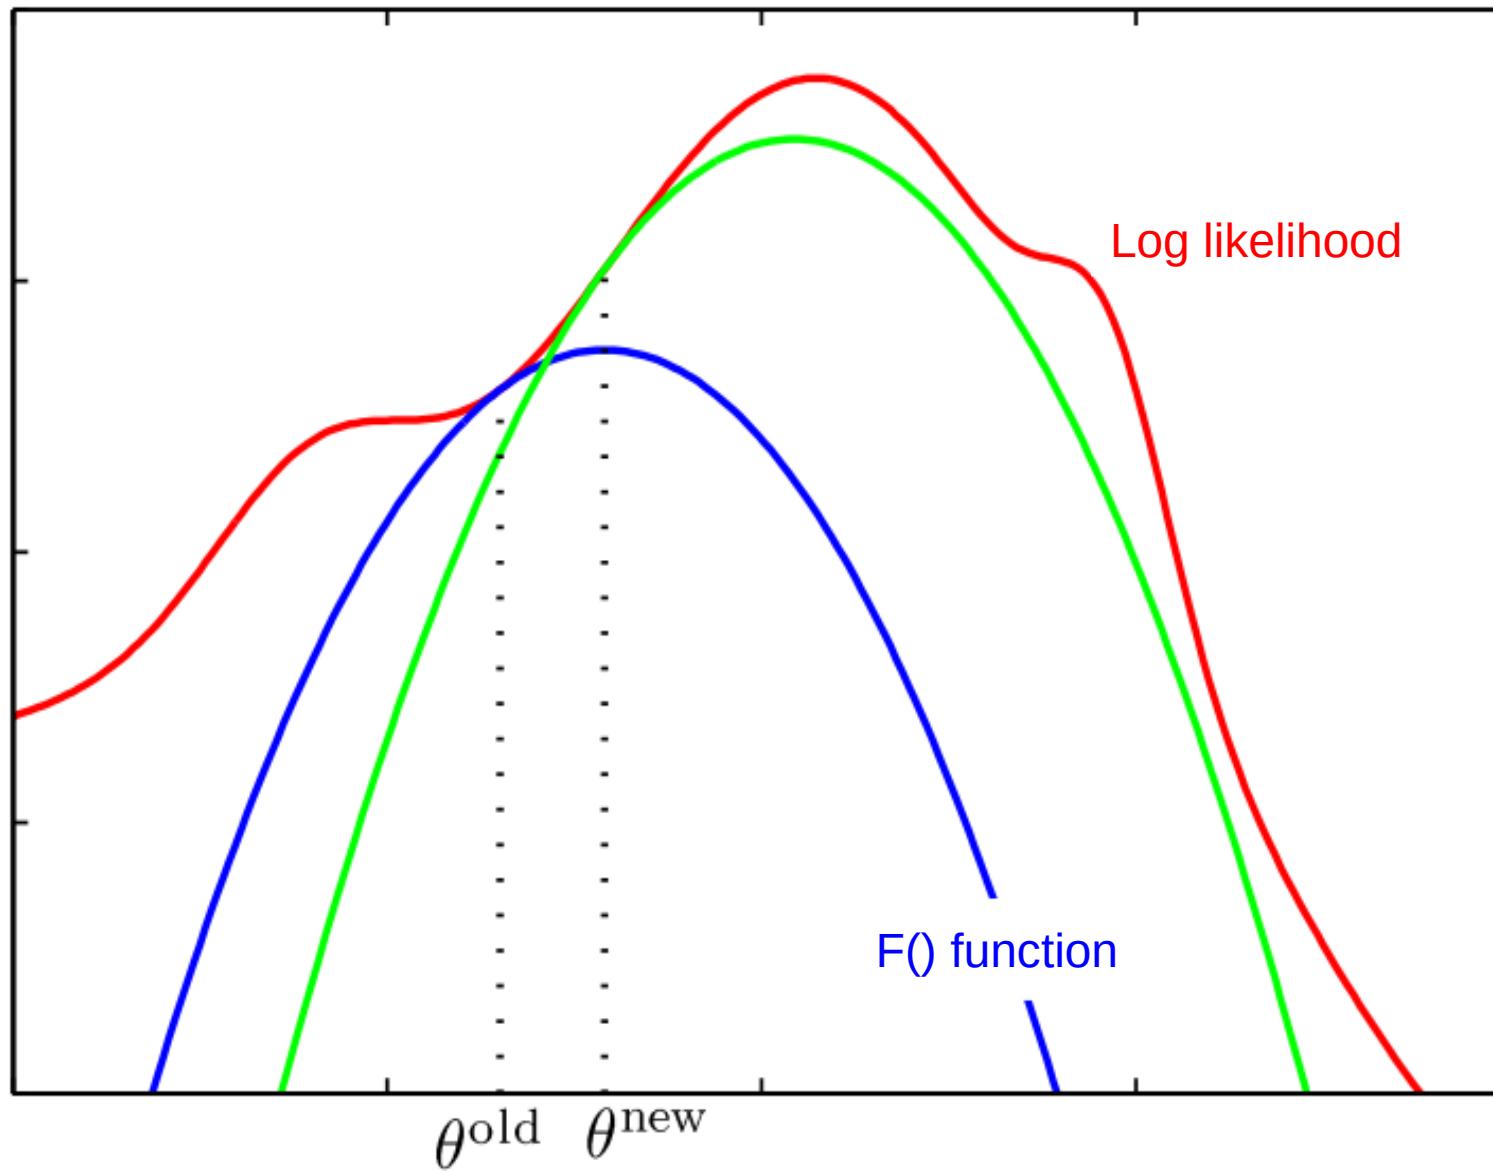
- EM performs iterative optimization
- Each EM iteration comprises 2 steps :
 - (1) E step
 - (2) M step

EM Optimization

- Assume
 - Iteration i
 - Parameter estimates are θ_i
- E step
 - Designs a function $F(\theta; \theta_i)$ that :
 - (1) is a lower bound for (log) likelihood function
 - (2) touches likelihood function at current estimate θ_i

EM Optimization

- EM: E step and M step within 1 iteration

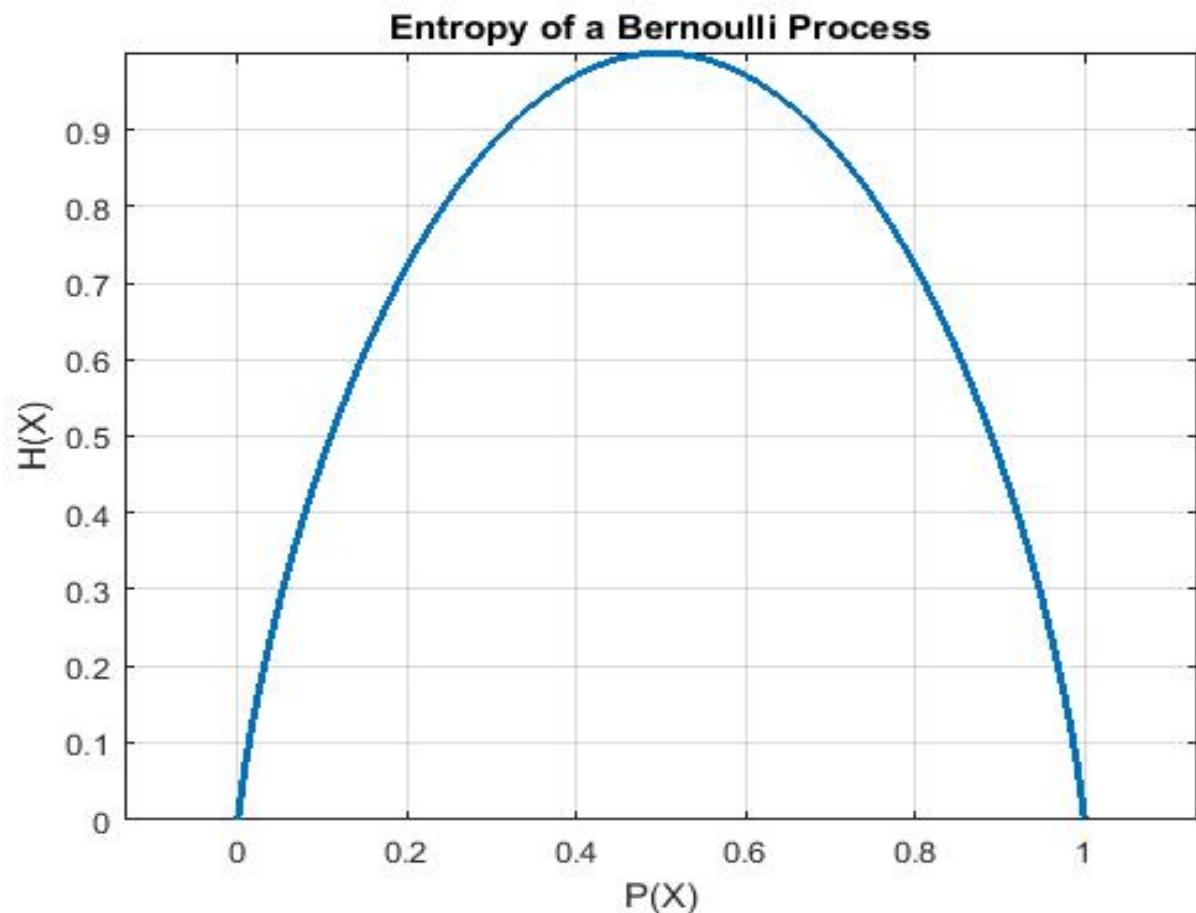


EM Optimization

- Assume
 - Iteration i
 - Parameter estimates are θ_i
- E step
 - Designs a function $F(\theta; \theta_i)$ that :
 - (1) is a lower bound for (log) likelihood function
 - (2) touches likelihood function at current estimate θ_i
- M step
 - Maximizes $F(\theta; \theta_i)$ over θ

Entropy

- Entropy (X) = $E [- \log P(X)]$
 - Shannon
- A measure of:
 - Uncertainty of random variable
 - Spread of PDF



Kullback–Leibler divergence

- Measure of dissimilarity between PDFs / PMFs

$$D_{\text{KL}}(P\|Q) = \sum_i P(i) \ln \frac{P(i)}{Q(i)}$$

$$D_{\text{KL}}(P\|Q) = \int_{-\infty}^{\infty} p(x) \ln \frac{p(x)}{q(x)} dx$$

- Expectation of log ratios of PDFs
 - Expectation taken over one of the PDFs
- Defined only when support of P subset of support of Q
 - $q(x) = 0 \rightarrow p(x) = 0$
- Positive
- Asymmetric
- NOT a distance metric

EM Optimization

- Consider the log likelihood = $\log P(y|\theta)$
- Log likelihood can be written as
 - $F(q,\theta) + KL(q \parallel p(x|y,\theta))$, where
- $F(q,\theta)$ is a:
 - Functional of the PDF $q(x)$
 - Function of parameters θ
 - $F(q,\theta) = \sum_x q(x) \log \{ P(y,x|\theta) / q(x) \}$
- $KL(q \parallel p(x|y,\theta))$ = KL divergence between PDFs
 - = $-\sum_x q(x) \log \{ P(x|y,\theta) / q(x) \}$
 - $KL(p_1, p_2)$ always ≥ 0
 - $KL(p_1, p_2) = 0$ iff $p_1 = p_2$

EM Optimization

- So, $F(q, \theta) = \log P(y|\theta) - KL(q \parallel p(x|y, \theta))$
- So, $F(q, \theta)$ is a **lower bound** of likelihood $\log P(y|\theta)$
 - For any chosen $q(x)$
 - $F(q, \theta^*)$ equals $\log P(y|\theta^*)$ iff $q(x)$ = true posterior $p(x|y, \theta^*)$
- We don't know true (ML) parameter values θ^*
- At iteration i , parameter estimates = θ_i
- **E step** = choose $q(\cdot)$ to maximize $F(q, \theta_i)$
 - $q(x) = p(x | y, \theta_i)$ = posterior PDF of labels given data
 - Thus, at θ_i , $F(q, \theta_i) = \log P(y | \theta_i)$
 - Thus, $F(\cdot)$ touches log-likelihood function at θ_i

EM Optimization

- So, $F(q, \theta) = \log P(y|\theta) - KL(q \parallel p(x|y, \theta))$
- So, $F(q, \theta)$ is a **lower bound** of likelihood $\log P(y|\theta)$
- We don't know true parameter values θ^*
- At iteration i, parameter estimates = θ_i
- E step fixes $q(x) := p(x | y, \theta_i)$
- M step = choose θ to maximize $F(q, \theta)$
 - Observe (rewrite RHS):
$$F(q, \theta) = \mathbb{E}_{q(\cdot)} [\log P(x, y|\theta)] + H(q), \text{ where}$$
 - $\mathbb{E}_{q(\cdot)} [\log P(x, y|\theta)]$ is our $Q(\theta, \theta_i)$ function
 - $H(q)$ = entropy of $q(\cdot)$ = NOT function of θ

$$F(q, \theta) = \sum_x q(x) \log \{ P(y, x|\theta) / q(x) \}$$

EM Optimization

- E step

$$Q(\theta; \theta^i) := E_{P(x|y, \theta^i)} [\log P(x, y|\theta)]$$

= expectation of the complete-data log likelihood
w.r.t. a posterior PDF of hidden variable x,
where

posterior PDF approximated via current parameter
estimates θ^i

- M step

$$\theta^{i+1} \leftarrow \arg \max_{\theta} Q(\theta; \theta^i)$$

- Convergence

- Guaranteed to converge to local maximum of likelihood function $P(y|\theta)$

EM Optimization

- What does the graph of $Q(\theta^{i+1}; \theta^i)$ w.r.t. 'i' look like ?
- Termination criterion ?
 - Relative change between parameter values between 2 consecutive iterations falls below a user-defined threshold
- What if we don't solve for maximum in M step ?
 - What else can we do ?
 - It should converge
 - It should be useful, e.g., optimizes likelihood

GMM Optimization via EM

- Observed random variable: $Y := \{ Y_n \} \quad n=1, \dots, N$
 - Leads to data $y := \{ y_n \} \quad n=1, \dots, N$
- PDF model: $P(x) := \sum_{k=1}^K w_k G(x; \mu_k, C_k)$
- Parameters: $\theta = \{ w_k, \mu_k, C_k \} \quad k=1, \dots, K$
- Hidden random variable: $Z := \{ Z_n \} \quad n=1, \dots, N$
 - $Z_n = z_n$ is label associated with data point y_n
 - z_n takes values $1, \dots, K$
- Optimization strategy
 - ML estimation of parameters using EM algorithm
 - Constraint: $\sum_k w_k = 1$

GMM Optimization via EM

- Assume, at iteration i , parameter estimates = θ^i
- E step

$$Q(\theta; \theta^i) := E_{P(z|y, \theta^i)} [\log P(y, z|\theta)]$$

$= E_{P(z|y, \theta^i)} [\log \prod_n P(y_n, z_n|\theta)]$ Independent observations, Independent labels

$= E_{P(z|y, \theta^i)} [\log \prod_n P(y_n|z_n, \theta) P(z_n|\theta)]$ Conditional probability

$$= E_{P(z|y, \theta^i)} \left[\sum_n \log \left(P(y_n|z_n, \theta) P(z_n|\theta) \right) \right]$$

GMM Optimization via EM

- E step

$$\begin{aligned} &= E_{P(z|y, \theta^i)} \left[\sum_n \log \left(P(y_n | z_n, \theta) P(z_n | \theta) \right) \right] \\ &= \sum_n E_{P(z|y, \theta^i)} \left[\log \left(P(y_n | z_n, \theta) P(z_n | \theta) \right) \right] \text{ Linearity of expectation} \\ &= \sum_n E_{P(z_n|y_n, \theta^i)} \left[\log \left(P(y_n | z_n, \theta) P(z_n | \theta) \right) \right] \text{ Conditional independence} \\ &= \sum_n \sum_k P(z_n = k | y_n, \theta^i) \log \left(P(y_n | z_n = k, \theta) P(z_n = k | \theta) \right) Z_n \text{ is a discrete RV} \end{aligned}$$

$$\text{GM} : \sum_n \sum_k P(z_n = k | y_n, \theta^i) \log \left(P(y_n | z_n = k, \theta) P(z_n = k | \theta) \right)$$

- E step

- Note that $w_k := P(z_n = k | \theta)$
 - Thus, w_k is non-negative
- Denote **membership** of y_n in k -th cluster by :

$$\begin{aligned}\gamma_{nk} &:= P(z_n = k | y_n, \theta^i) \\ &= \frac{P(y_n | z_n = k, \theta^i) P(z_n = k | \theta^i)}{P(y_n | \theta^i)} \text{ Bayes rule} \\ &= \frac{G(y_n | \mu_k^i, C_k^i) w_k^i}{\sum_k G(y_n | \mu_k^i, C_k^i) w_k^i}\end{aligned}$$

- Each membership is non-negative
- For each point y_n , sum of memberships over K classes is 1

$$\textbf{GM} : \sum_n \sum_k P(z_n = k | y_n, \theta^i) \log \left(P(y_n | z_n = k, \theta) P(z_n = k | \theta) \right)$$

- E step

- Finally, we have the Q(.) function as :

$$Q(\theta; \theta^i) = \sum_n \sum_k \gamma_{nk} \left(-0.5 \log |C_k| - 0.5(y_n - \mu_k)' C_k^{-1} (y_n - \mu_k) + \log w_k \right)$$

- Gammas are fixed
 - Mahalanobis distance

P C Mahalanobis

- Statistician (1893-1972)
- Founded Indian Statistical Institute
 - In 1932
- Sample surveys
 - Using sampling
 - Consumer expenditure, tea-drinking habits, public opinion, crop acreage, plant disease



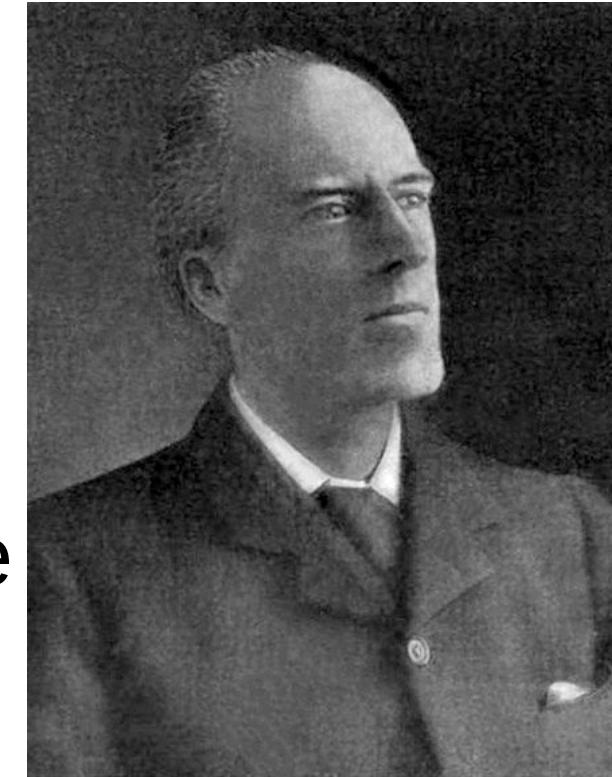
P V Sukhatme



- Statistician
- Pioneering work in the 1940s by applying random sampling methods in agricultural statistics and in biometry
- Made significant contributions in the statistical theory of sampling
 - PhD advisors: J Neyman, ES Pearson
 - ES Pearson ← son of Karl Pearson
 - Karl Pearson founded the world's first university statistics department at University College London in 1911
 - Established the field of mathematical statistics
 - ESP also later became Head of Applied Statistics dept at UCL

Karl Pearson

- Correlation coefficient
- Method of moments
- Chi distance: precursor and special case of Mahalanobis distance
- P value
- Foundations of the statistical hypothesis testing theory and the statistical decision theory
- Pearson's chi-squared test
- PCA
- First introduction of the histogram



GMM Optimization via EM

- M step $Q(\theta; \theta^i) = \sum_n \sum_k \gamma_{nk} \left(-0.5 \log |C_k| - 0.5(y_n - \mu_k)' C_k^{-1} (y_n - \mu_k) + \log w_k \right)$
 - Update parameter estimates by solving :
$$\arg \max_{\theta} Q(\theta; \theta^i)$$
under the constraint $\sum_k w_k = 1$
 - Solve for the Gaussian means μ_k and covariances C_k
 - Take partial derivatives and assign them to 0

GMM Optimization via EM

- M step $Q(\theta; \theta^i) = \sum_n \sum_k \gamma_{nk} \left(-0.5 \log |C_k| - 0.5(y_n - \mu_k)' C_k^{-1} (y_n - \mu_k) + \log w_k \right)$
 - Solve for **mean**
 - Derivative of a quadratic form
 $\partial (x' A x) = x' (A + A') \partial x$, where
 - $x' A x$ is a scalar
 - $x' (A + A')$ is a row vector
 - ∂x is a column vector
 - <http://www.ee.ic.ac.uk/hp/staff/dmb/matrix/calculus.html>

$$\frac{\partial Q(\theta; \theta^i)}{\partial \mu_k} = 0 = \sum_n \gamma_{nk} C_k^{-1} (y_n - \mu_k)$$
$$\Rightarrow \mu_k = \frac{\sum_n \gamma_{nk} y_n}{\sum_n \gamma_{nk}}$$

GMM Optimization via EM

- M step $Q(\theta; \theta^i) = \sum_n \sum_k \gamma_{nk} \left(-0.5 \log |C_k| - 0.5(y_n - \mu_k)' C_k^{-1} (y_n - \mu_k) + \log w_k \right)$
 - Solve for covariance
 - Derivative of a determinant
 $\partial |\mathbf{A}| = |\mathbf{A}| (\mathbf{A}^{-T}) :' \partial \mathbf{A}$
 - where “ : ” corresponds to vectorization
 - Derivative of an inverse
 $\partial (x' \mathbf{A}^{-1} y) = -(A^{-T} x y' A^{-T}) :' \partial \mathbf{A}$

$$\frac{\partial Q(\theta; \theta^i)}{\partial C_k} = 0 = - \sum_n \gamma_{nk} C_k^{-T} + \sum_n \gamma_{nk} C_k^{-T} (y_n - \mu_k) (y_n - \mu_k)' C_k^{-T}$$
$$\Rightarrow C_k = \frac{\sum_n \gamma_{nk} (y_n - \mu_k) (y_n - \mu_k)'}{\sum_n \gamma_{nk}}$$

GMM Optimization via EM

- M step $Q(\theta; \theta^i) = \sum_n \sum_k \gamma_{nk} \left(-0.5 \log |C_k| - 0.5(y_n - \mu_k)' C_k^{-1} (y_n - \mu_k) + \log w_k \right)$
 - Solve for **weights** w_k
 - Constrained optimization problem

$$\arg \max_{\{w_k\}} Q(\theta; \theta^i) \text{ such that } \sum_k w_k = 1$$

$$\arg \max_{\{w_k\}} \sum_n \sum_k \gamma_{nk} \log w_k \text{ such that } \sum_k w_k = 1$$

- Use method of Lagrange multipliers

GMM Optimization via EM

- M step
 - Solve for **weights** w_k
 - Constrained optimization problem
 - Lagrangian is :

$$\begin{aligned} L(\{w_k\}) &:= \sum_n \sum_k \gamma_{nk} \log w_k + \lambda \left(\sum_k w_k - 1 \right) \\ &= \sum_k \log w_k \gamma_k + \lambda \left(\sum_k w_k - 1 \right) \end{aligned}$$

where $\gamma_k := \sum_n \gamma_{nk}$ is the **total membership** associated with the k -th Gaussian

GMM Optimization via EM

- M step
 - Solve for **weights** w_k
 - Constrained optimization problem

$$\frac{\partial L(\{w_k\})}{\partial w_k} = 0 = \frac{\gamma_k}{w_k} + \lambda \implies w_k = \frac{-\gamma_k}{\lambda}$$

$$\frac{\partial L(\{w_k\})}{\partial \lambda} = 0 = \sum_k w_k - 1$$

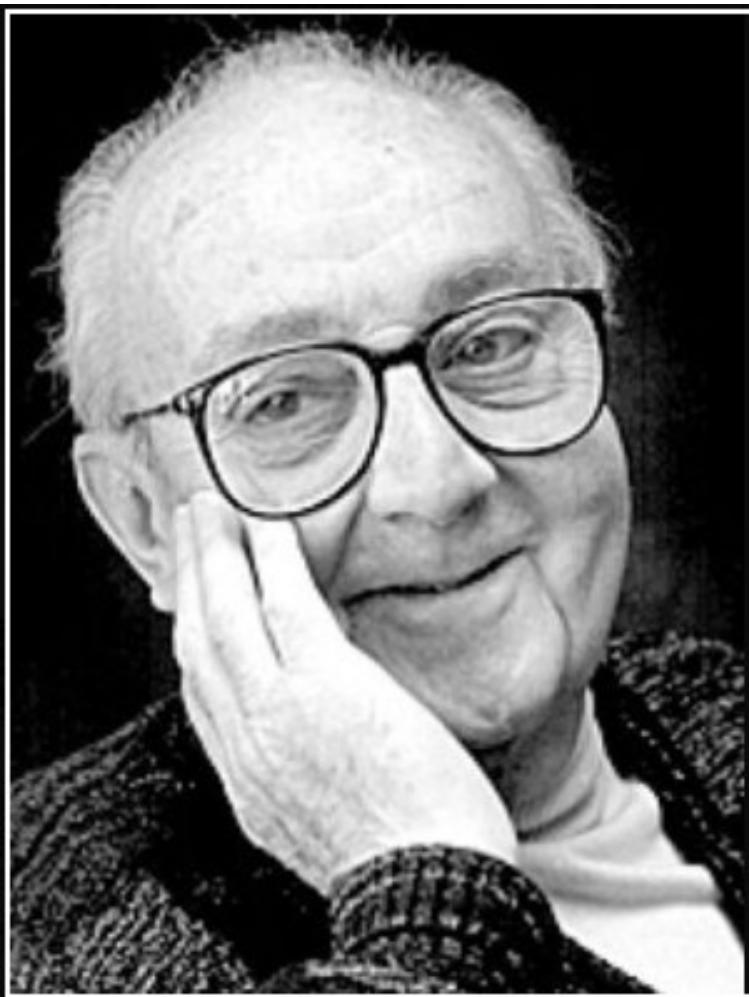
Thus,

$$\lambda = - \sum_k \gamma_k \text{ and } w_k = \frac{\gamma_k}{\sum_k \gamma_k} = \frac{\gamma_k}{N}$$

where we note that $\sum_k \gamma_k = \sum_k \sum_n \gamma_{nk} = \sum_n (\sum_k \gamma_{nk}) = \sum_n (1) = N$

GMM-EM Segmentation: Limitation

- Doesn't enforce spatial smoothness on segmentation
 - How to do that ?
 - MRF model on label images



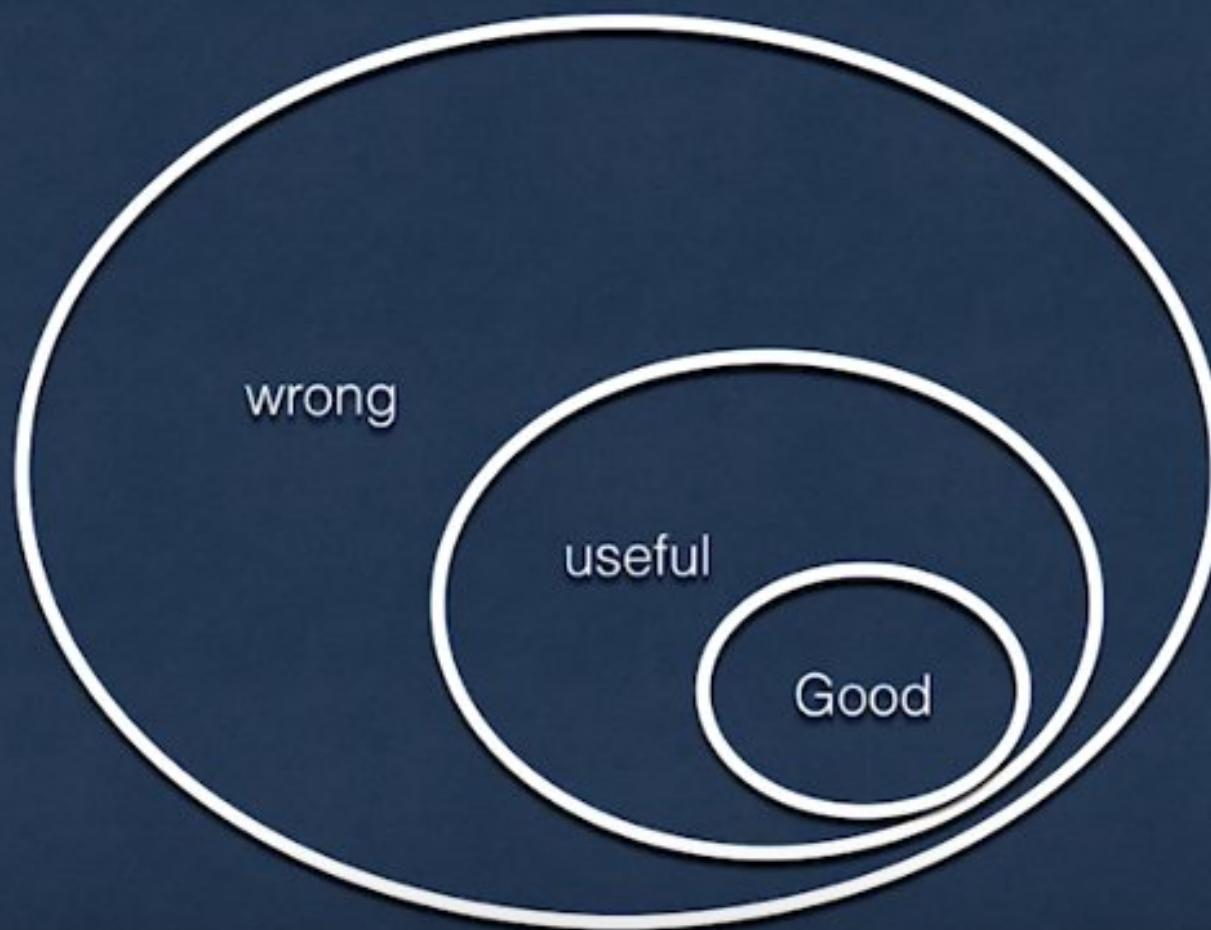
MODELS:

“Essentially, all models are wrong, but some are useful.”

“Remember that all models are wrong; the practical question is how wrong do they have to be to not be useful.”

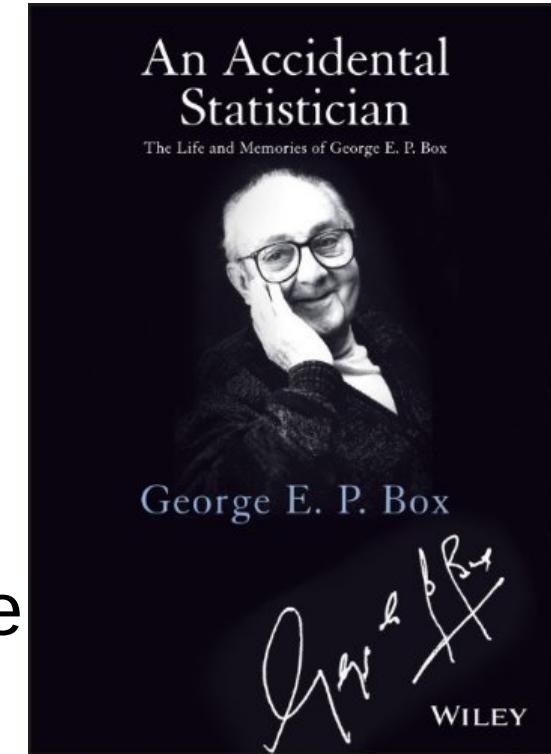
George E.P. Box

Models



George E. P. Box

- Statistician (1919 - 2013)
 - “One of the great statistical minds of the 20th century”
 - Quality control, time-series analysis, design of experiments, Bayesian inference
- Advisor:
ES Pearson, son of Karl Pearson
- During World War II, performed experiments for the British Army exposing small animals to poison gas. To analyze the results of his experiments, he taught himself statistics from available texts
- After war, got PhD from UCL



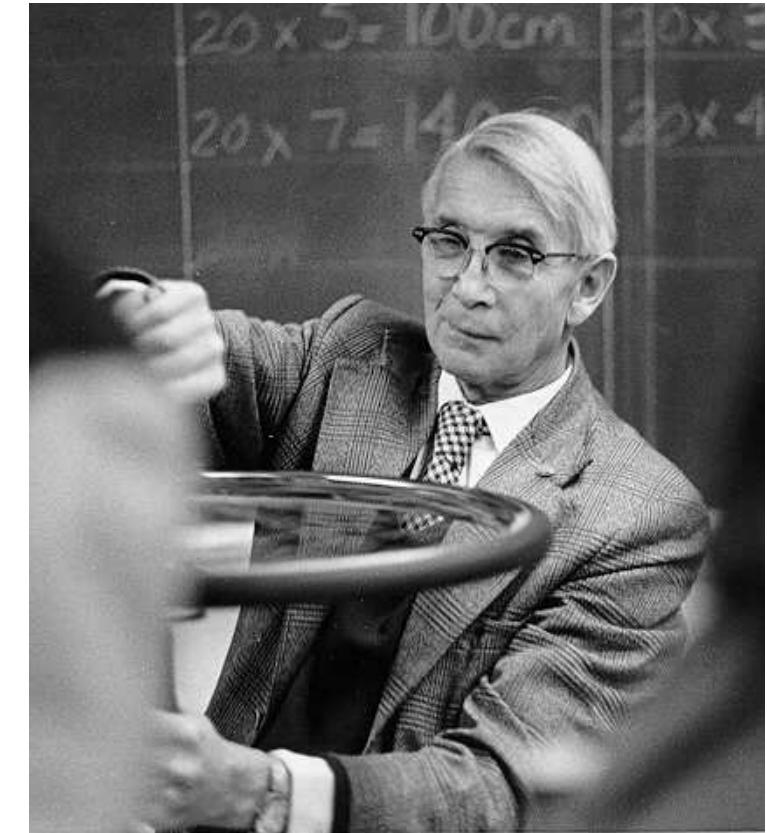
Priors on Label Images

- **MRF on Labels**

- Binary labels
 - Ising model
- Labels taking multiple values
 - Potts model
- MRF where potential function on 2-cliques
 - Gives a constant penalty if (neighboring) labels are unequal
 - Otherwise, gives a zero penalty
 - Does a quadratic error function make sense ?

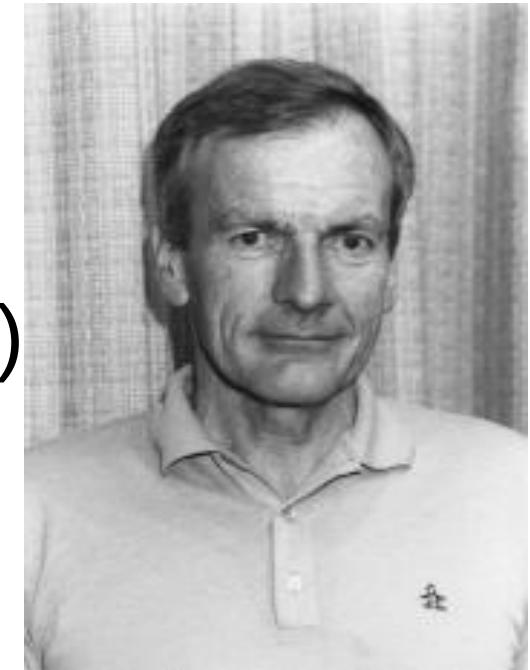
Ernst Ising

- Physicist (1900 – 1998)
- Advisor: Wilhelm Lenz
 - Studied ferromagnetism with Wolfgang Pauli (Nobel Laureate)
- Finished PhD in 1924
 - PhD work = Ising model
 - Perhaps, start of probabilistic graphical models
- Life
 - Forbidden to teach / research in Hitler's regime (1933+)
 - For some time, worked as a shepherd, railroad worker
 - Immigrated to US in 1947, never published later
 - Realized success of his PhD work in 1949



Renfrey Potts

- Mathematician (1925 – 2005)
- Proposed Potts model in his PhD (1951)
- Switched fields to Operations Research
- Potts model got recognition 20 years after it was first published



EM - 1977

- Arthur P. Dempster (1929 -)
 - Joined Harvard statistics in 1957 (4 faculty in dept.)
 - Advisor: John Tukey (FFT fame)
 - Founding member of Princeton's statistics department
- NM Laird (1943 -) ← PhD student of Dempster
- DB Rubin (1943 -)
 - Worked for ETS
 - In Harvard psychology grad school.
 - Asked to take intro stats courses ← insufficient backgrnd.
 - Felt insulted, given background in physics @ Princeton.
 - Transferred to applied math → PhD in stats @ Harvard.
 - 271K+ citations
 - Rubin ← Cochran ← Wishart ← Karl Pearson

Bayesian Image Segmentation

- **MAP-MRF Segmentation (Hard Seg.)**
 - **Hidden MRF + GMM**
 - Consider **label MRF $X = x$** , with labels $k \in 1, 2, \dots, K$
 - Consider observed image **data $Y = y$**
 - Generative model for intensities given class label :
 $P(Y | X) := \prod_i P(Y_i | X_i)$
 - e.g., for each class k , assume that the model is a Gaussian $P(Y_i | X_i = k) = G(Y_i | \mu_k, \sigma_k)$
 - Let θ = parameters underlying noise model and MRF model (in general)
 - No weight parameters w_k
 - Why not ? What is a 'procedure' of generating data ?

Bayesian Image Segmentation

- MAP-MRF Segmentation (Hard Seg.)
 - Optimization problem
 - Obtain segmentation x and parameter estimates θ by optimizing $\max_{x,\theta} P(x | y, \theta)$
 - Optimization strategy
 - Alternate between :
 - (i) finding optimal parameters $\max_{\theta} P(x | y, \theta)$ and
 - (ii) finding optimal segmentation $\max_x P(x | y, \theta)$

Bayesian Image Segmentation

- MAP-MRF Segmentation (Hard Seg.)
 - Optimal parameters
 - Assume MRF model parameters to be fixed
 - Given hard segmentation x ,
if noise model is Gaussian,
then,
optimal parameters for each class = ?
 - Sample mean & variance (over pixels labeled to that class)

Bayesian Image Segmentation

- MAP-MRF Segmentation **via ICM** (Hard Seg.)

- MAP segmentation

- $\max_x P(x | y, \theta)$

- Rewrite $P(X | y, \theta) =$

$$= P(X_i, X_{\sim i} | y, \theta)$$

$$= P(X_i | X_{\sim i}, y, \theta)P(X_{\sim i} | y, \theta) \text{ Conditional probability}$$

$$= P(X_i | X_{N_i}, y, \theta)P(X_{\sim i} | y, \theta) \text{ Markov assumption on } X$$

$$= P(X_i | X_{N_i}, y_i, \theta)P(X_{\sim i} | y, \theta) \text{ Conditional independence in noise model}$$

Bayesian Image Segmentation

- MAP-MRF Segmentation via ICM (Hard Seg.)
 - At each voxel i , perform :

$$\begin{aligned}\max_{x_i} P(X|y, \theta) &= \max_{x_i} P(X_i|X_{N_i}, y_i, \theta) P(X_{\sim i}|y, \theta) \\&= \max_{x_i} P(X_i|X_{N_i}, y_i, \theta) \text{ Second term doesn't depend on } x_i \\&= \max_{x_i} P(y_i|X_i, X_{N_i}, \theta) P(X_i|X_{N_i}, \theta) / P(y_i|X_{N_i}, \theta) \text{ Bayes Rule} \\&= \max_{x_i} P(y_i|X_i, X_{N_i}, \theta) P(X_i|X_{N_i}, \theta) \text{ Denominator doesn't depend on } x_i \\&= \max_{x_i} P(y_i|X_i, \theta) P(X_i|X_{N_i}, \theta) \text{ Conditional independence assumption in }\end{aligned}$$

- $P(y_i | X_i, \theta)$ = label likelihood
- $P(X_i | X_{N_i}, \theta)$ = label prior
- How to perform this maximization ?

Bayesian Image Segmentation

- MAP-MRF Segmentation (Hard Seg.)
 - Order of Label Updates
 - Sequentially: Column by column, and then row by row
 - Can lead to artifacts
 - Sequentially: Randomized order each iteration
 - Need to generate random sequence each iteration
 - Are artifacts eliminated ?
 - In Parallel: Doesn't guarantee increase in image probability
 - No need to generate random sequence
 - Are artifacts eliminated ?
 - In Parallel: Update only if image probability increases !!
 - Typically, it will !
 - May terminate too soon !

Bayesian Image Segmentation

- **MAP-MRF Segmentation (Soft Seg.)**
 - **Hidden MRF + GMM + EM**
 - <https://www.ncbi.nlm.nih.gov/pubmed/11293691>

IEEE TRANSACTIONS ON MEDICAL IMAGING, VOL. 20, NO. 1, JANUARY 2001

Segmentation of Brain MR Images Through a Hidden Markov Random Field Model and the Expectation-Maximization Algorithm

Yongyue Zhang*, Michael Brady, and Stephen Smith

$$\text{Bayesian Ir} \sum \sum_{n \quad k} P(z_n = k | y_n, \theta^i) \log \left(P(y_n | z_n = k, \theta) P(z_n = k | \theta) \right)$$

- **Hidden MRF + GMM + EM (Soft Seg)**

- Number of voxels N
- Observed Data Variable = $Y = \{ Y_i \} \quad n=1, \dots, N$
- Label MRF (Hidden) Variable = $X = \{ X_i \} \quad n=1, \dots, N$
- Number of classes = L
- Parameters $\theta = \{ \mu_l, \sigma_l \} \quad l=1, \dots, L$
 - No weight parameters
 - Weight parameters $w_k \rightarrow$ priors on number of points derived from class k, while computing memberships
 - We replace by MRF prior
- For voxel i, neighbors are denoted by N_i

Bayesian Image Segmentation

- Hidden MRF + GMM + EM (Soft Seg)
 - Optimization Strategy
 - Find ML estimates for parameters
 $\max_{\theta} P(y|\theta)$
 - Using EM

Bayesian Image Segmentation

- Hidden MRF + GMM + EM (Soft Seg)
 - E step

$$\begin{aligned} Q(\theta|\theta^t) &:= E_{P(X|y,\theta^t)}[\log P(X, y|\theta)] \\ &= E_{P(X|y,\theta^t)}[\log(P(y|X, \theta)P(X))] \text{ Conditional Probability} \\ &= E_{P(X|y,\theta^t)}[\log P(X) + \log \Pi_i P(y_i|X_i, \theta)] \text{ i.i.d. noise model} \\ &= E_{P(X|y,\theta^t)}[\log P(X) + \sum_i \log P(y_i|X_i, \theta)] \end{aligned}$$

Bayesian Image Segmentation

- Hidden MRF + GMM + EM (Soft Seg)

- E step

$$\begin{aligned} &= E_{P(X|y, \theta^t)}[\log P(X)] + E_{P(X|y, \theta^t)}\left[\sum_i \log P(y_i|X_i, \theta)\right] \text{ linearity of expectation} \\ &= E_{P(X|y, \theta^t)}[\log P(X)] + \sum_i E_{P(X|y, \theta^t)}[\log P(y_i|X_i, \theta)] \text{ linearity of expectation} \\ &= E_{P(X|y, \theta^t)}[\log P(X)] + \sum_i E_{P(X_i, X_{\sim i}|y, \theta^t)}[\log P(y_i|X_i, \theta)] \end{aligned}$$

- We can ignore the first term
 - Why ?
 - What is the purpose of the Q(.) function ?
 - E step requires integration w.r.t. posterior prob on X
 - How to handle that ?

Bayesian Image Segmentation

- Hidden MRF + GMM + EM (Soft Seg)

- E step

- Approximate expectation :

$$E_{P(x_i, x_{\sim i} | y, \theta^t)} [\log P(y_i | x_i, \theta)] \approx E_{P(x_i | x_{\sim i}, y, \theta^t)} [\log P(y_i | x_i, \theta)]$$

- Left hand side
 - Sum over possible labels at i-th voxel,
accounting for all possibilities of neighbor labels (& their neighbors, & ...)
 - Right hand side
 - Sum over possible labels at i-th voxel, keeping rest of image fixed !
 - This somewhat underestimates sampling variability !
This also introduces some bias !
 - Label image X fixed to what ? X is hidden !
 - Fix to MAP segmentation, given current parameters. Works for brain seg.

Bayesian Image Segmentation

- Hidden MRF + GMM + EM (Soft Seg)

- E step

Calling $C := \log E_{P(x|y,\theta^t)}[P(x)]$ as a constant independent of θ , we get

$$\begin{aligned} Q(\theta|\theta^t) - C &\approx \sum_i E_{P(x_i|x_{\sim i}^{\text{MAP}}, y, \theta^t)}[\log P(y_i|x_i, \theta)] \\ &= \sum_i E_{P(x_i|x_{N_i}^{\text{MAP}}, y, \theta^t)}[\log P(y_i|x_i, \theta)] \\ &= \sum_i \sum_{l=1}^L P(x_i = l|x_{N_i}^{\text{MAP}}, y, \theta^t) \log P(y_i|x_i = l, \theta) \end{aligned}$$

- Below is what we had for GMM + EM (without MRF):

$$\sum_n \sum_k P(z_n = k|y_n, \theta^i) \log \left(P(y_n|z_n = k, \theta) P(z_n = k|\theta) \right)$$

Bayesian Image Segmentation

- Hidden MRF + GMM + EM (Soft Seg)
 - E step
 - Lets see how the MRF prior comes in
 - M step
 - Remains very similar to that of GMM fitting via EM (without MRF)
 - What changes ? Only memberships changes. Lets see how.

Bayesian Image Segmentation

- Hidden MRF + GMM + EM (Soft Seg)

- E step

- **Memberships** = $P(x_i = l | \overset{\text{MAP}}{x_{N_i}}, y, \theta^t) =$

$$= P(x_i = l | y_i, \overset{\text{MAP}}{x_{N_i}}, \theta^t) \text{ Conditional Independence of } X_i, Y_{\sim i}, \text{ Given } Y_i, \overset{\text{MAP}}{X_{N_i}}$$

$$= \frac{P(y_i | x_i = l, \overset{\text{MAP}}{x_{N_i}}, \theta^t) P(x_i = l | \overset{\text{MAP}}{x_{N_i}}, \theta^t)}{P(y_i | \overset{\text{MAP}}{x_{N_i}}, \theta^t)} \text{ Bayes Rule}$$

$$= \frac{G(y_i | \mu_l, \sigma_l) P(x_i = l | \overset{\text{MAP}}{x_{N_i}})}{\sum_{l=1}^L G(y_i | \mu_l, \sigma_l) P(x_i = l | \overset{\text{MAP}}{x_{N_i}})} \text{ Conditional Independence of } Y_i, X_{N_i}, \text{ Given } X_i$$

- What is the prior here ? What is its effect ?
 - Below is what we had for GMM + EM (without MRF):

$$\begin{aligned}\gamma_{nk} &:= P(z_n = k | y_n, \theta^i) \\ &= \frac{P(y_n | z_n = k, \theta^i) P(z_n = k | \theta^i)}{P(y_n | \theta^i)} \text{ Bayes rule} \\ &= \frac{G(y_n | \mu_k^i, C_k^i) w_k^i}{\sum_k G(y_n | \mu_k^i, C_k^i) w_k^i}\end{aligned}$$

Bayesian Image Segmentation

- Hidden MRF + GMM + EM (Soft Seg)

- E step
 - Conditional label probability is:

$$P(X_i|X_{S-\{i\}}) = \frac{\exp\left(-\sum_{a \in A_i} V_a(X_a)\right)}{\sum_{x'_i} \exp\left(-\sum_{a \in A_i} V_a(X_a)\right)} = \frac{\exp\left(-\sum_{a \in A_i} V_a(X_a)\right)}{Z_i}$$

- What is this probability in case all the following hold:
 - 2 classes + 4 neighborhood
 - 3 neighbors are L1, 1 neighbor is L2
 - $V(L1, L2) = \text{beta}$ when $L1 \approx L2$, where parameter beta > 0
 - $V(L1, L2) = 0$ otherwise

Bayesian Image Segmentation

- Hidden MRF + GMM + EM (Soft Seg)

- M step

- If we call memberships $\rightarrow \gamma_{nk}$
 - Update for mean and (co)variance

$$\mu_k = \frac{\sum_n \gamma_{nk} y_n}{\sum_n \gamma_{nk}}$$

$$C_k = \frac{\sum_n \gamma_{nk} (y_n - \mu_k)(y_n - \mu_k)'}{\sum_n \gamma_{nk}}$$

- How will you make segmentation smoother ?
 - What parameter will you modify ?
 - What happens if the MRF prior is removed ?

Bayesian Image Segmentation

- Hidden MRF + GMM + EM (Soft Seg)
- Algorithm
 - Initialize parameters: means, covariances
 - **E step (approximation)**
 - 1) Compute MAP label image, given parameters
 - How to compute this ?
 - 2) Evaluate memberships
 - Spatially smooth
 - **M Step**
 - 3) Update means and covariances, for all classes
 - Repeat E and M steps, until convergence
 - **Output:** Memberships (soft, spatially smooth)

Bayesian Image Segmentation

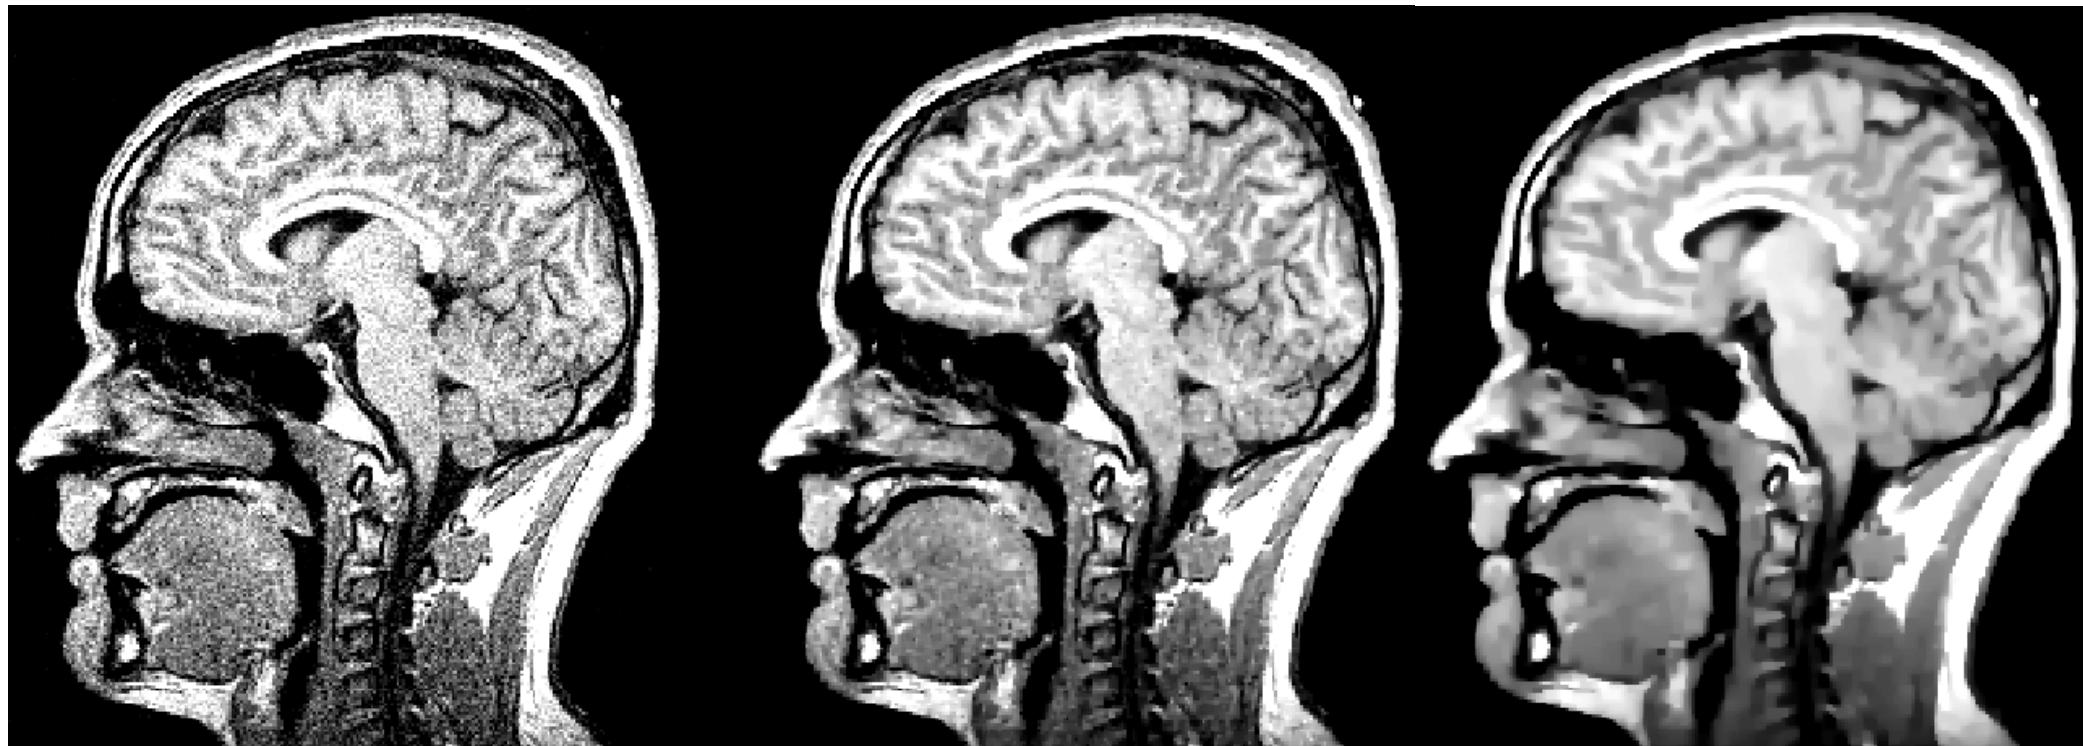
- How to get initial segmentation ?
 - e.g., for brain tissue segmentation

Validation

- How to validate an algorithm for real-world applications ?
 - (A) real-world phantoms
 - (B) human clinical data
- What about segmentation ?
 - What is “ground truth” ?
 - Kinds of variability
- What about denoising, reconstruction ?
- What about segmentation ?

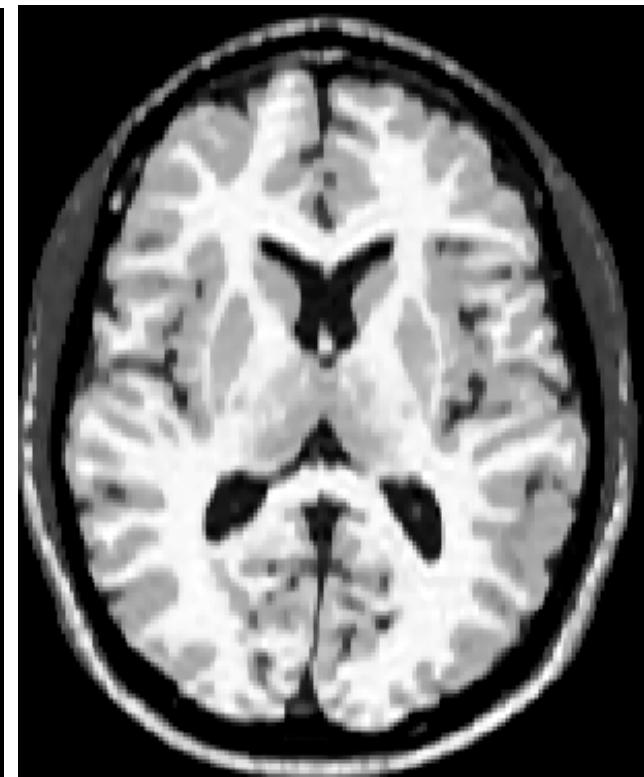
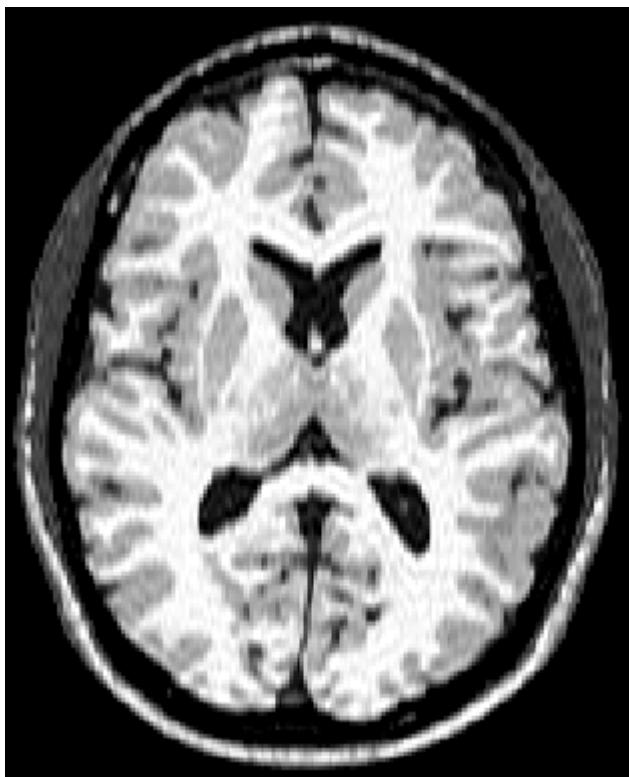
Validation

- Examples of denoised images using different values of algorithm parameters



Validation

- Examples of reconstructed images using different values of algorithm parameters





Normal Distribution



Paranormal Distribution

Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Number of voxels N
 - Observed Data Variable = $Y = \{ Y_i \} \quad n=1, \dots, N$
 - Label MRF Variable = $X = \{ X_i \} \quad i=1, \dots, N$
 - Number of classes = 2
 - Let **label values x_i be binary, i.e., 0 or 1**
 - Generative model for intensities given class label
 $P(Y | X) := \prod_i P(Y_i | X_i)$
 - Let θ = parameters underlying noise model
 - Assume MRF model parameters to be fixed

Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Optimization Problem
 - Obtain segmentation x and parameter estimates θ by maximizing posterior
 - Optimization Strategy
 - Alternate between
 - (i) finding optimal parameters $\max_{\theta} P(x | y, \theta)$ and
 - (ii) finding optimal segmentation $\max_x P(x | y, \theta)$
 - Optimizing parameters is easy
 - If noise model is Gaussian, given the segmentation, optimal parameters for each class are sample mean and sample variance over voxels in that class

Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Optimal segmentation, given parameters

$$\max_x P(x|y, \theta)$$

$$= \max_x \log P(x|y, \theta)$$

$$= \max_x (\log P(y|x, \theta) + \log P(x))$$

$$= \max_x \left(\log \prod_i P(y_i|x_i, \theta) + \log \frac{\exp(0.5 \sum_{(i,j) \in \mathcal{N}} \beta_{ij} V(x_i, x_j))}{Z} \right)$$

Consider only cliques \mathcal{N}

where

- Z is a constant that depends only on the MRF parameters that are fixed/known
- $\beta_{ij} \geq 0$: non-negative (convention; relative to how the potential function $V()$ is defined)
- $\beta_{ij} := \beta_{ji}$: neighbor interactions are symmetric
- $\beta_{ii} := 0$: interaction with self isn't allowed

Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Assume that potential function $V(\cdot, \cdot)$ is defined, for the case of binary label values, as
$$V(a, b) := ab + (1 - a)(1 - b)$$
 - When neighbor labels are same :
 - $a=b=0$ or $a=b=1$, then $V(a,b) = 1$
 - Leads to higher-prob state because $\beta_{ij} > 0$
 - When neighbor labels are different :
 - $a=1-b=0$ or $a=1-b=1$, then $V(a,b) = 0$
 - Leads to lower-prob state because $\beta_{ij} > 0$

Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Rewrite the likelihood function

$$\Pi_i P(y_i|x_i, \theta) = \Pi_i P(y_i|x_i = 1, \theta)^{x_i} P(y_i|x_i = 0, \theta)^{1-x_i}$$

- Simplifying the objective function (log Posterior)

$$\begin{aligned} \max_x P(x|y, \theta) &= \max_x \left(\sum_i x_i \log P(y_i|x_i = 1, \theta) + (1 - x_i) \log P(y_i|x_i = 0, \theta) \right. \\ &\quad \left. + 0.5 \sum_i \sum_j \beta_{ij} (x_i x_j + (1 - x_i)(1 - x_j)) \right) \end{aligned}$$

- Factor 0.5 ensures that pairs aren't considered twice

Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Simplifying the objective function (log Posterior)

$$\max_x P(x|y, \theta) = \max_x \left(\sum_i \lambda_i x_i + 0.5 \sum_i \sum_j \beta_{ij} (2x_i x_j - x_i - x_j) \right)$$

where

$$\lambda_i := \log P(y_i|x_i = 1, \theta) - \log P(y_i|x_i = 0, \theta) = \log \frac{P(y_i|x_i=1, \theta)}{P(y_i|x_i=0, \theta)}$$

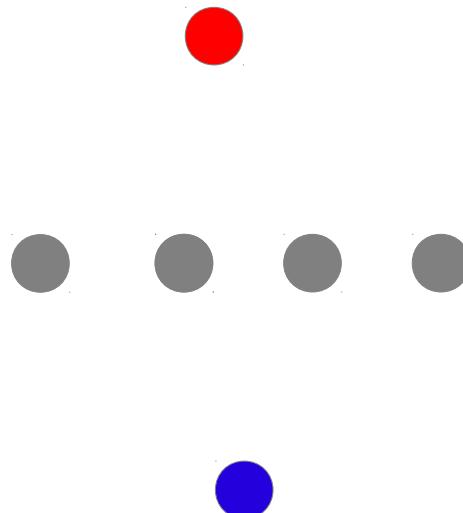
= log of likelihood ratio

(which is a constant independent of the segmentation)

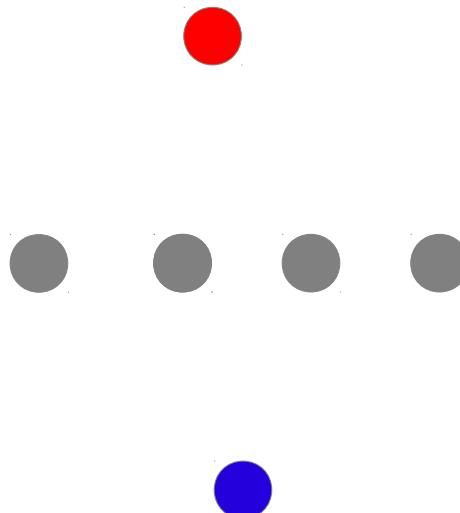
Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Construct an s-t graph as follows :
 - (1) Add a vertex for each voxel i
 - (2) Add two additional vertices called s and t

Example 1



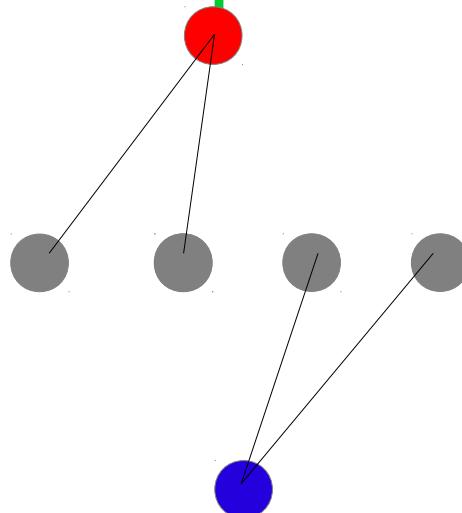
Example 2



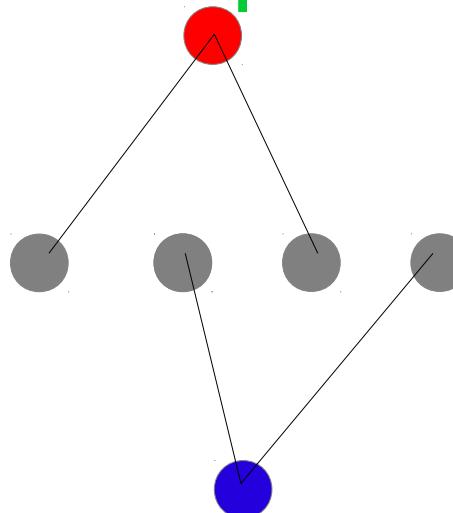
Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Construct an s-t graph as follows :
 - (3) For each voxel i , if $\lambda_i > 0$, then add an edge from vertex s to vertex i with cost $c_{si} := \lambda_i > 0$
 - (4) For each voxel i , if $\lambda_i \leq 0$, then add an edge from vertex i to vertex t with cost $c_{it} := -\lambda_i > 0$

Example 1



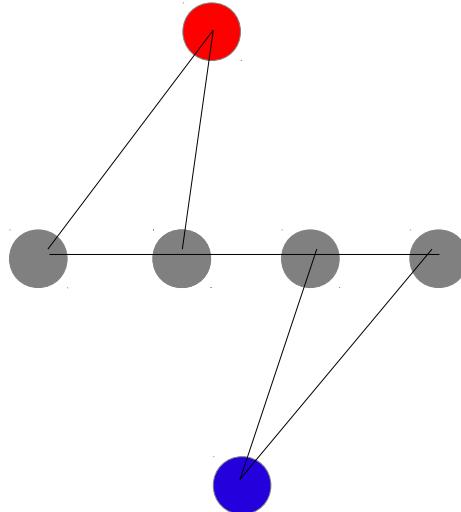
Example 2



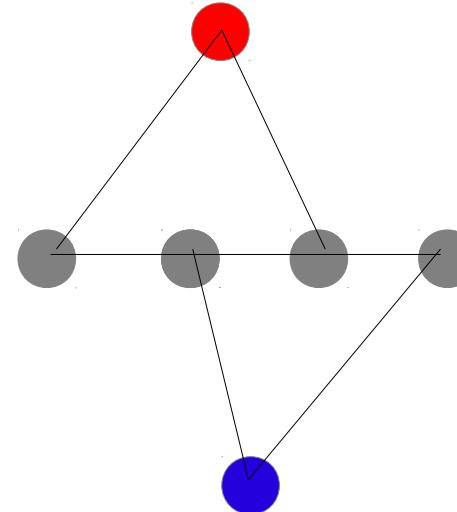
Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Construct an s-t graph as follows :
 - (5) For every pair of neighboring voxels i, j in image, add an edge with cost $c_{ij} := \beta_{ij} > 0$

Example 1



Example 2



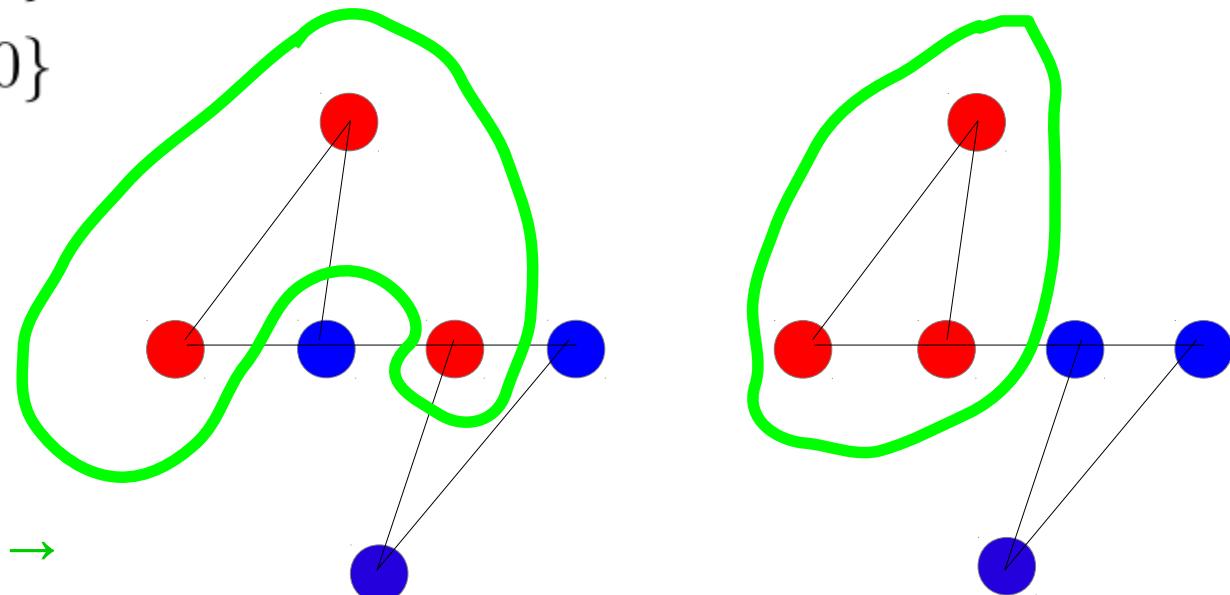
Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - A **cut** of the s-t graph is a **partition** of the graph into **2 mutually-exclusive & exhaustive sets** of vertices s.t.
 - one of the sets contains “s” (& all vertices in set have label 1)
 - the other set contains “t” (& all vertices in set have label 0)

$$S := \{s\} \cup \{i : x_i = 1\}$$

$$T := \{t\} \cup \{i : x_i = 0\}$$

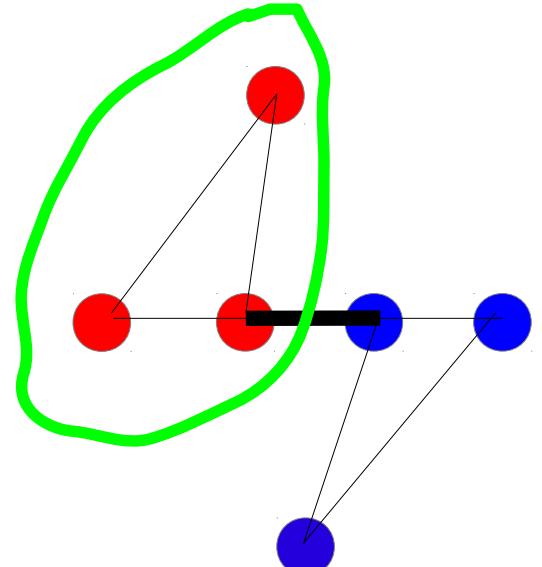
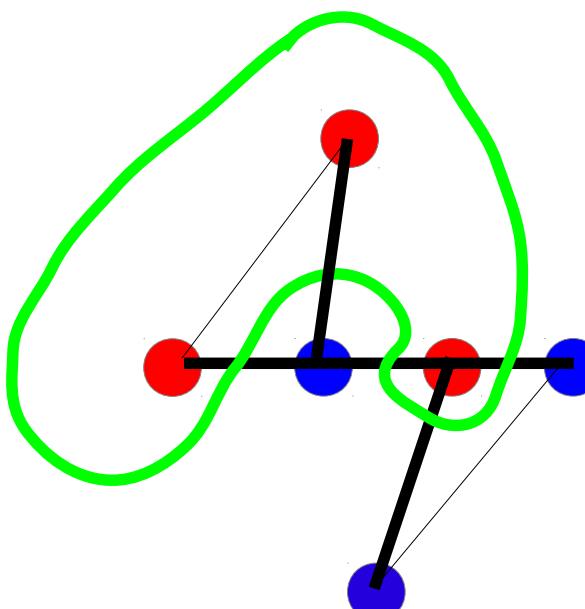
- *Cut separates vertex s from vertex t*
- *2 example cuts for graph in Example 1 →*



Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - **Capacity** of the cut = ?
 - Capacity of the cut (S, T)
= sum of costs of each edge with one of the vertices in S and the other vertex in T

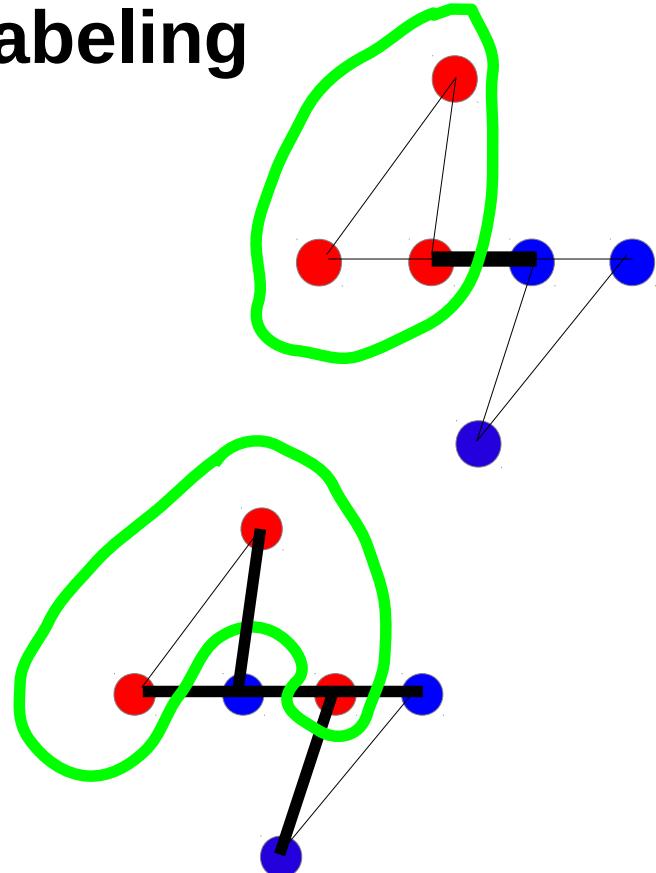
$$C(S, T) := \sum_{u \in S, v \in T} c_{uv}$$



Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - **Claim : Min s-t Cut equals MAP labeling**
 - Is this true ? Lets analyze.
 - For some cut (S, T) :
 - Or equivalently, for some labeling x :

$$\begin{aligned}C(S, T) &:= \sum_{u \in S, v \in T} c_{uv} \\&= \sum_{i \in T} c_{si} \text{ Edges between vertex } s \text{ and vertices in set } T \\&\quad + \sum_{i \in S} c_{it} \text{ Edges between vertices in set } S \text{ and vertex } t \\&\quad + \sum_{i \in S - \{s\}, j \in T - \{t\}} c_{ij} \text{ Edges between vertices in set } S - \{s\} \text{ and vertices in set } T - \{t\}\end{aligned}$$



Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Claim : Min s-t Cut equals MAP labeling
 - (1) Edges $s-i$ between “s” and vertices in set T **exist iff**
 - i belongs to T (i.e., label $x_i = 0$) AND
 - an edge exists between s and i (i.e., $\lambda_i > 0$)
 - If edge exists, its **cost = λ_i**
 - For any voxel i,
Edge Existence + Edge Cost are captured by :
$$(1 - x_i) \max(0, \lambda_i)$$

Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Claim : Min s-t Cut equals MAP labeling
 - (2) Edges $i \rightarrow t$ between vertices in set S and “t” exist iff
 - i belongs to S (i.e., label $x_i = 1$) AND
 - an edge exists between i and t (i.e., $\lambda_i \leq 0$)
 - If edge exists, its cost = $-\lambda_i$
 - For any voxel i,
Edge Existence + Edge Cost are captured by :
 $x_i \max(0, -\lambda_i)$

Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Claim : Min s-t Cut equals MAP labeling
 - (3) Edges $i-j$ between vertices in sets S-s and T-t exist iff
 - labels at the two vertices are different (i.e., $x_i \neq x_j$)
 - If edge exists, its cost = β_{ij}
 - For any voxels i,j
Edge Existence + Edge Cost are captured by :
$$(x_i - x_j)^2 \beta_{ij}$$

Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Claim : Min s-t Cut equals MAP labeling
 - Capacity of the cut reduces to :

$$C(S, T) = \sum_i (1 - x_i) \max(0, \lambda_i) \text{ Edges between vertex } s \text{ and vertices in set } T$$

$$+ \sum_i x_i \max(0, -\lambda_i) \text{ Edges between vertices in set } S \text{ and vertex } t$$

$$+ \sum_i \sum_j 0.5(x_i - x_j)^2 \beta_{ij} \text{ Edges between vertices in set } S - \{s\} \text{ and vertices in set } T - \{t\}$$

$$C(S, T) := \sum_{u \in S, v \in T} c_{uv}$$

$$= \sum_{i \in T} c_{si} \text{ Edges between vertex } s \text{ and vertices in set } T$$

$$+ \sum_{i \in S} c_{it} \text{ Edges between vertices in set } S \text{ and vertex } t$$

$$+ \sum_{i \in S - \{s\}, j \in T - \{t\}} c_{ij} \text{ Edges between vertices in set } S - \{s\} \text{ and vertices in set } T - \{t\}$$

Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Claim : Min s-t Cut equals MAP labeling
 - Capacity of the cut reduces to :

$$\begin{aligned} &= \sum_i \max(0, \lambda_i) + \sum_i x_i \left(\max(0, -\lambda_i) - \max(0, \lambda_i) \right) + \sum_i \sum_j 0.5(x_i + x_j - 2x_i x_j) \beta_{ij} \\ &= \sum_i \max(0, \lambda_i) + \sum_i x_i \left(-\lambda_i \right) + 0.5 \sum_i \sum_j \beta_{ij} (x_i + x_j - 2x_i x_j) \end{aligned}$$

Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Proof : Min s-t Cut equals MAP labeling

$$\min_{(S,T)} C(S, T) = \min_x C(S(x), T(x))$$

$$\begin{aligned} &= \min_x \left(\sum_i \max(0, \lambda_i) + \sum_i x_i (-\lambda_i) + 0.5 \sum_i \sum_j \beta_{ij} (x_i + x_j - 2x_i x_j) \right) \\ &= \min_x \left(- \sum_i x_i \lambda_i - 0.5 \sum_i \sum_j \beta_{ij} (2x_i x_j - x_i - x_j) \right) \\ &= \max_x \left(\sum_i x_i \lambda_i + 0.5 \sum_i \sum_j \beta_{ij} (2x_i x_j - x_i - x_j) \right) \\ &= \max_x P(x|y, \theta) \end{aligned}$$

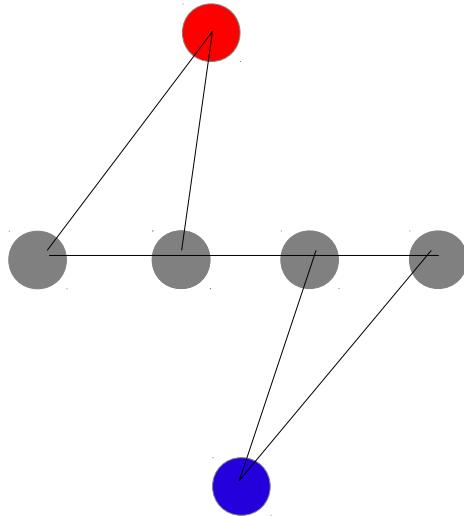
Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)
 - Strengths
 - Global optimum in polynomial time !
 - Various algorithms : polynomial in number of vertices or edges or both
 - <http://heim.ifi.uio.no/~geird/bergen.pdf>
 - <http://www.cs.princeton.edu/courses/archive/spr05/cos423/lectures/07maxflow.pdf>
 - NOT obvious because number of possible cuts grows exponentially with number of vertices
 - Any likelihood (noise) model can be built in
 - Limitations
 - Doesn't produce soft memberships
 - Handles only 2 classes
 - Min cut with > 2 labels is NP hard
(but fast approximations with bounded sub-optimality exist)
 - <http://www.cs.cornell.edu/courses/cs5540/2010sp/lectures/Lec18-graph-cuts.v1.pdf>

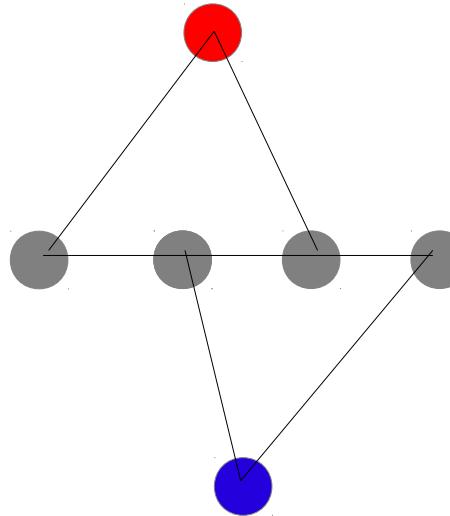
Bayesian Image Segmentation

- MAP-MRF Segmentation via s-t Cuts (Hard Seg)

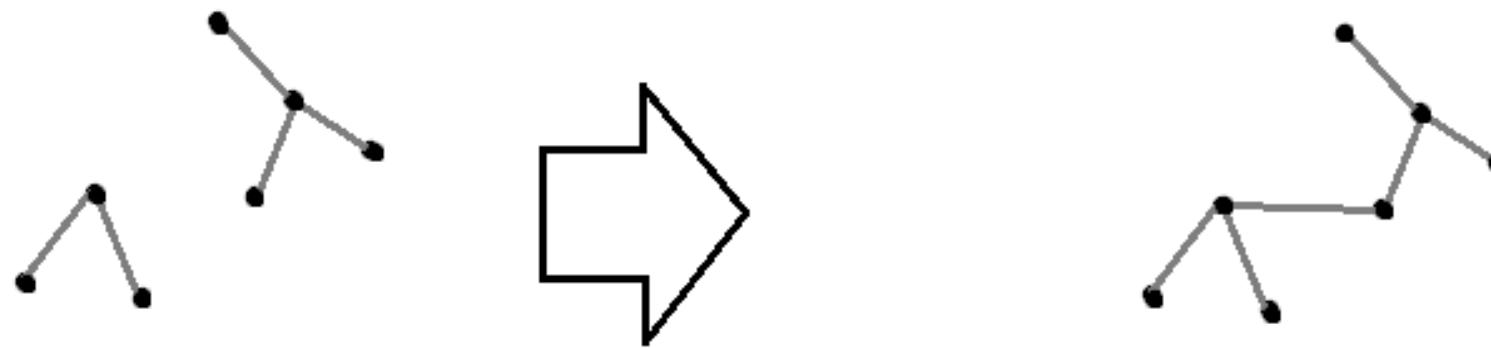
- Example 1



- Example 2



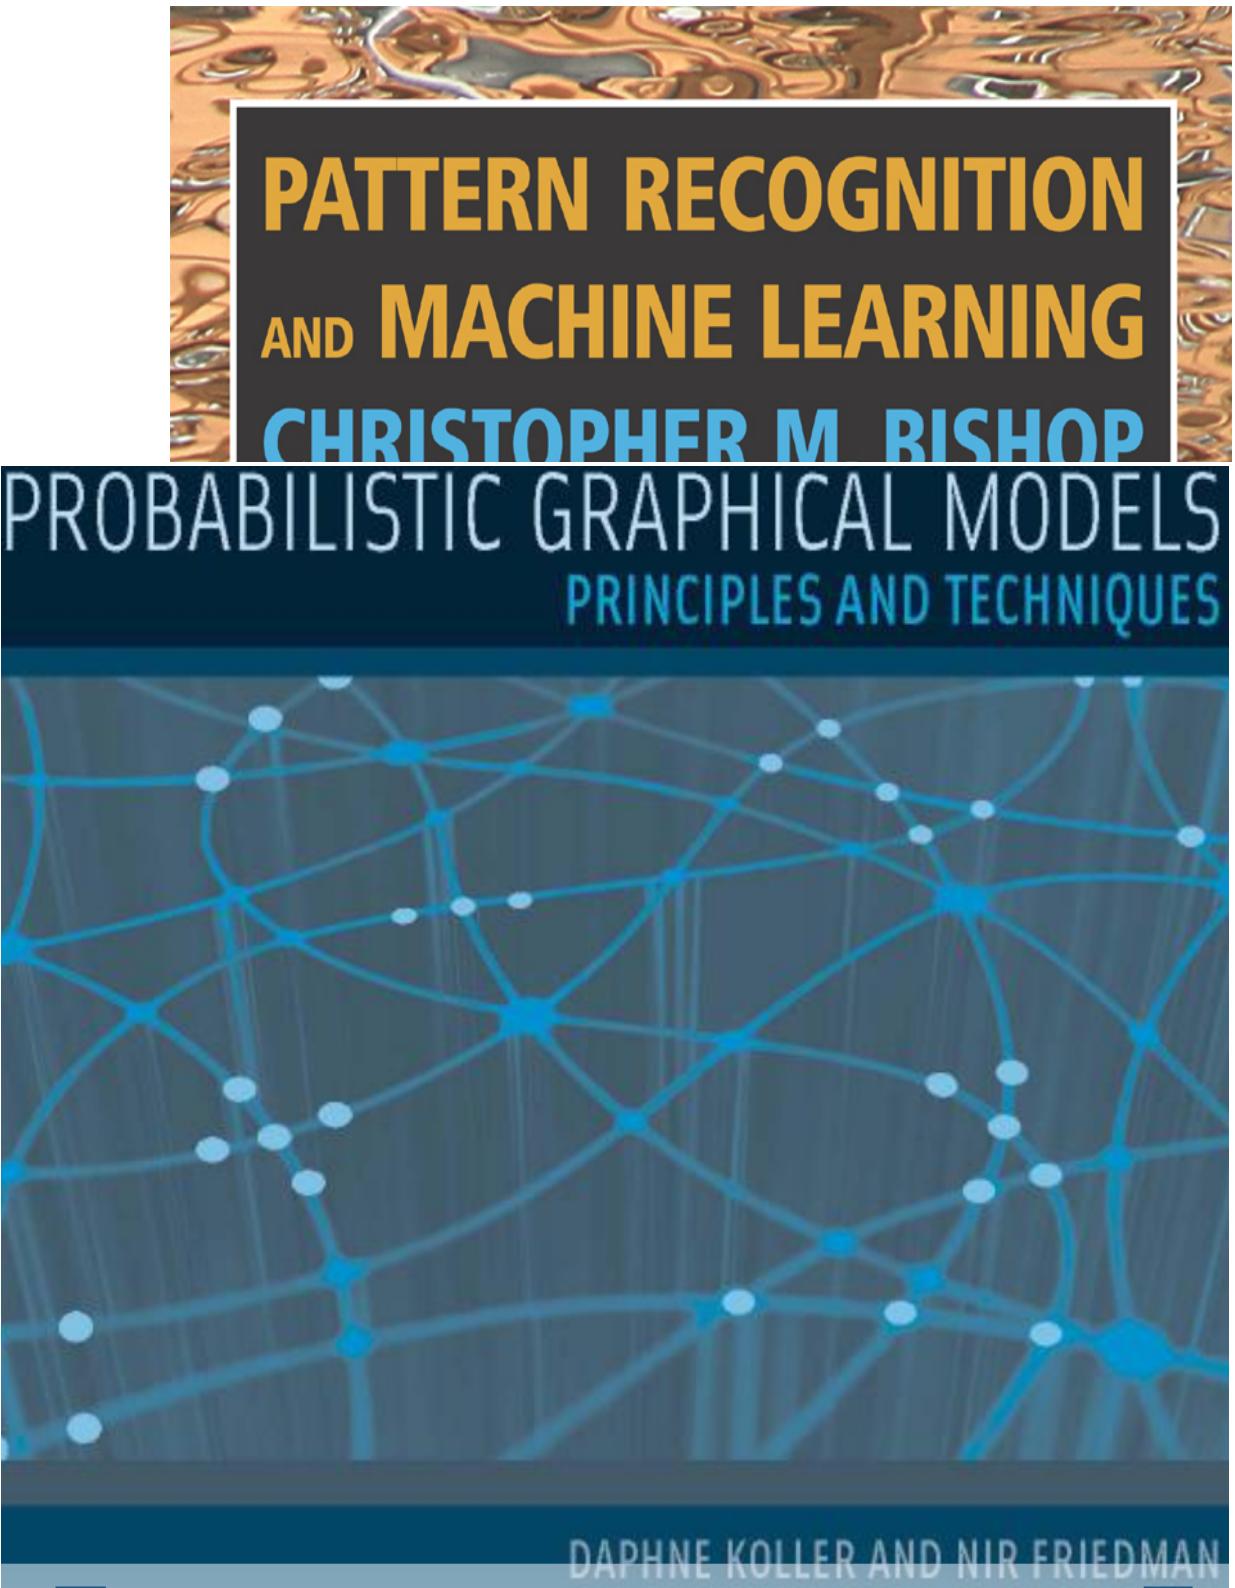
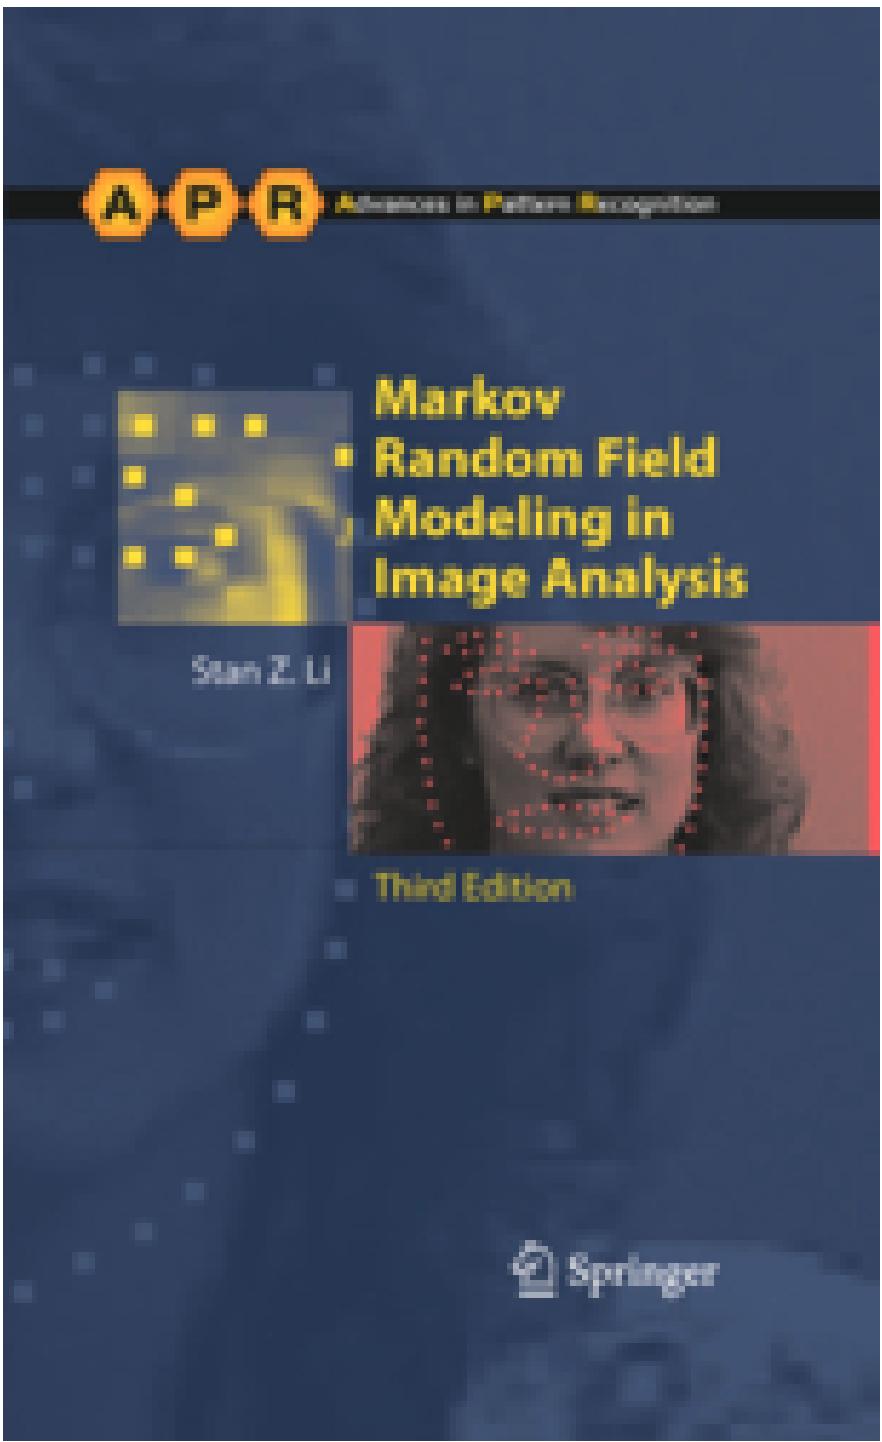
- (1) What are the cuts if all $\beta_{ij} = 0$?
- (2) What are possible cuts if β_{ij} values $\gg |\lambda_i|$ values ?



Deforestation:

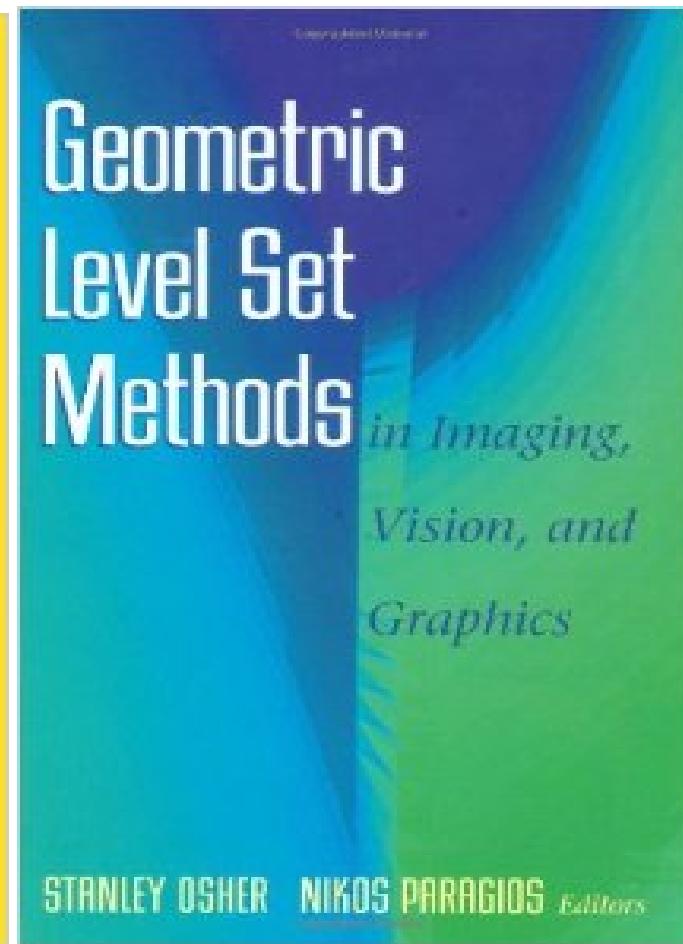
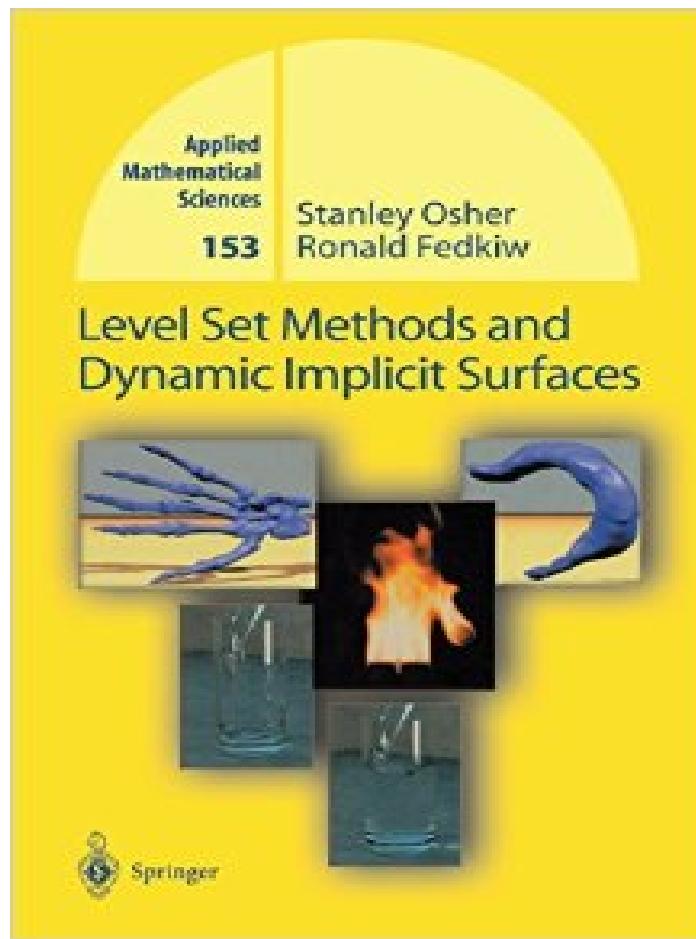
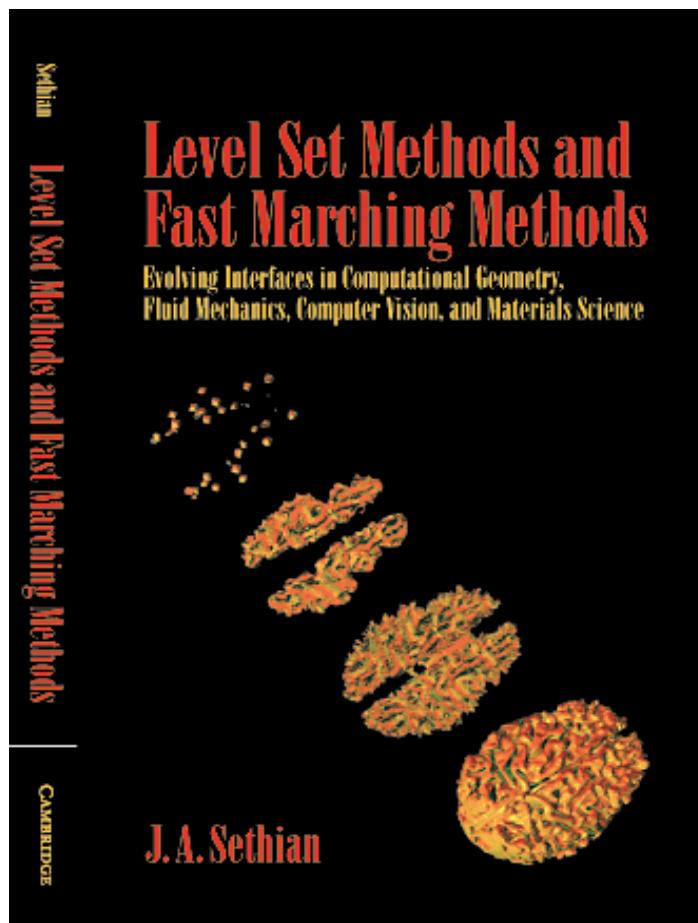
When adding a branch gives you fewer trees.

For More ...



Level-Set Based Segmentation

- A brief overview
 - https://math.berkeley.edu/~sethian/2006/level_set.html
- Books



Level-Set Based Segmentation

- Level set = contour = iso-surface
 - Of some function

Level-Set Based Segmentation

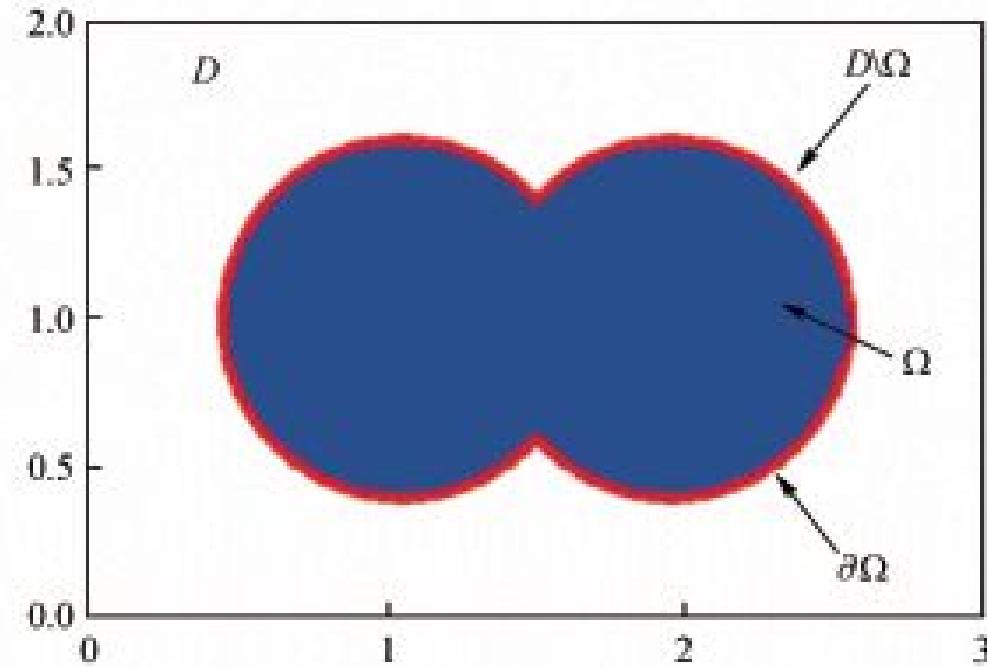
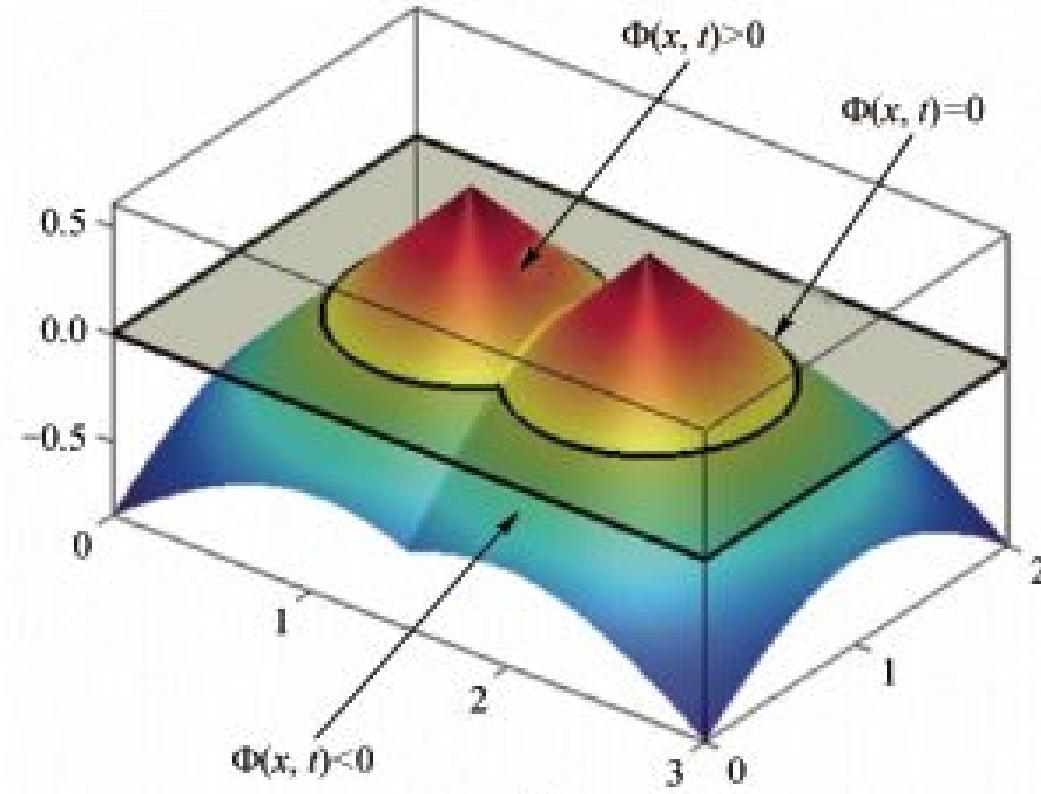
- Key ideas

- (1) Object boundary / surface representation

- (1) Represent object boundary / surface as **zero-level set** of ...

- (2) ... **boundary's (signed) distance transform**

- 1D Curve (2D surface) represented using a 2D (3D) image



Level-Set Based Segmentation

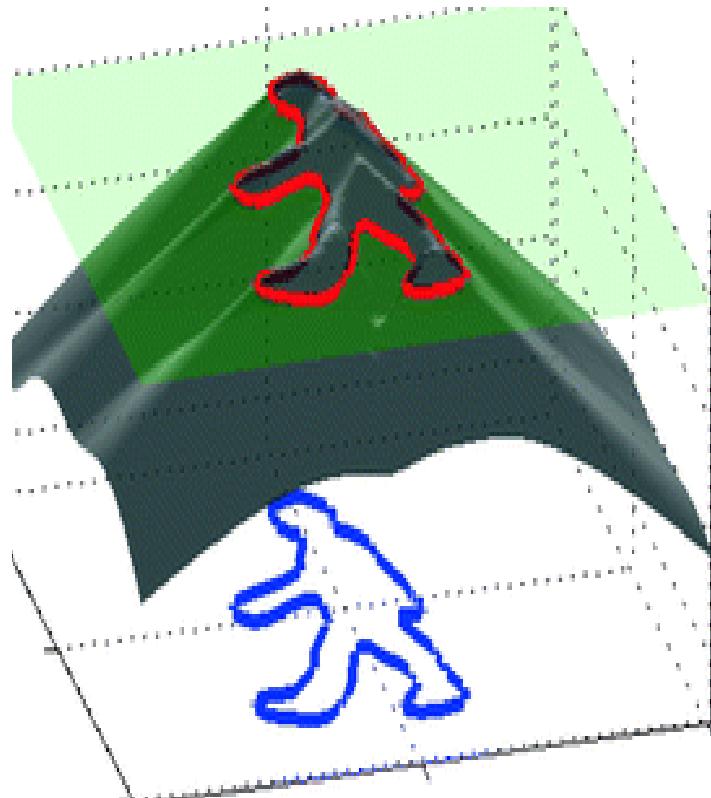
- Key ideas

- (1) Object boundary / surface representation

- (1) Represent object boundary / surface as zero-level set of ...

- (2) ... boundary's (signed) distance transform

- 1D Curve (2D surface) represented using a 2D (3D) image

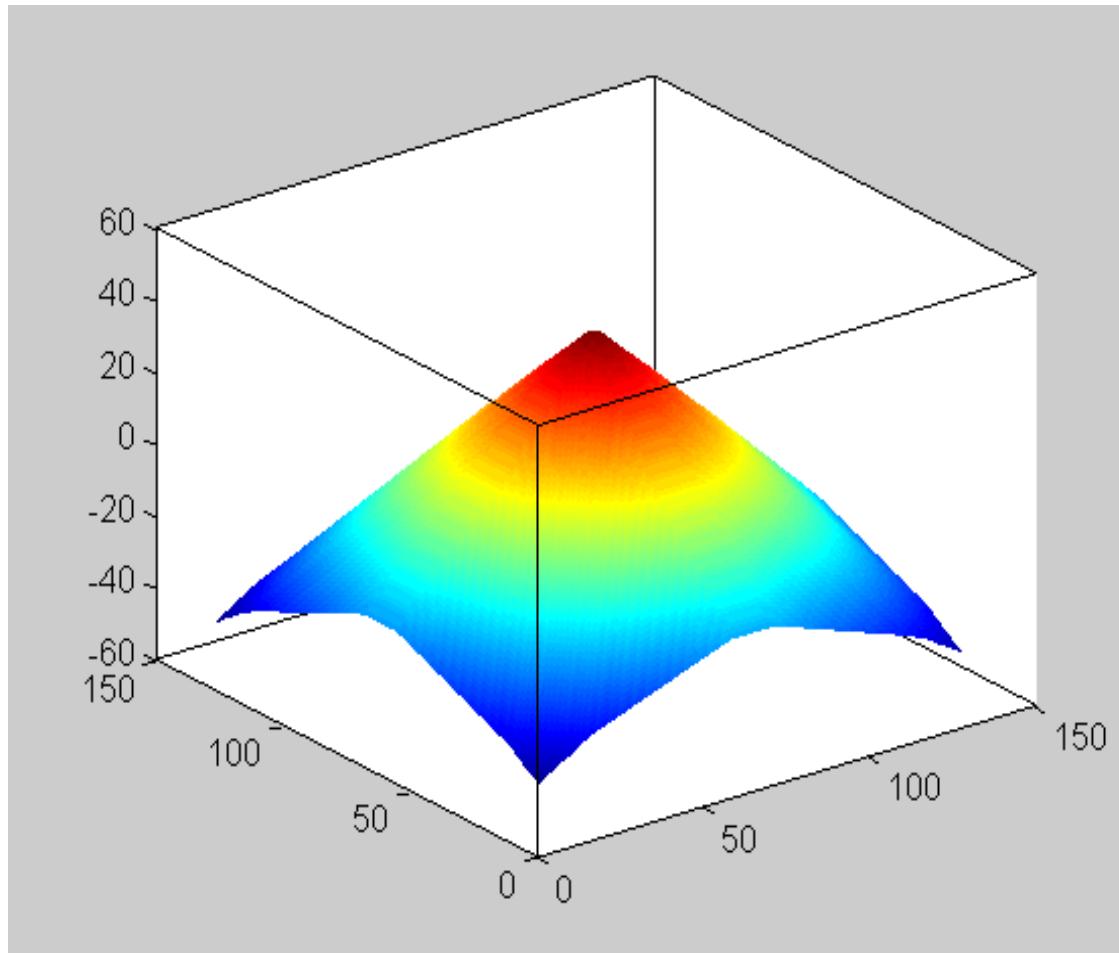


Level-Set Based Segmentation

- Key ideas

- (2) Surface evolution (to fit to the data)

- Modify distance transform to move surface
 - http://www.cs.swan.ac.uk/~csjason/rbf/gifs/1circle_levelset.gif



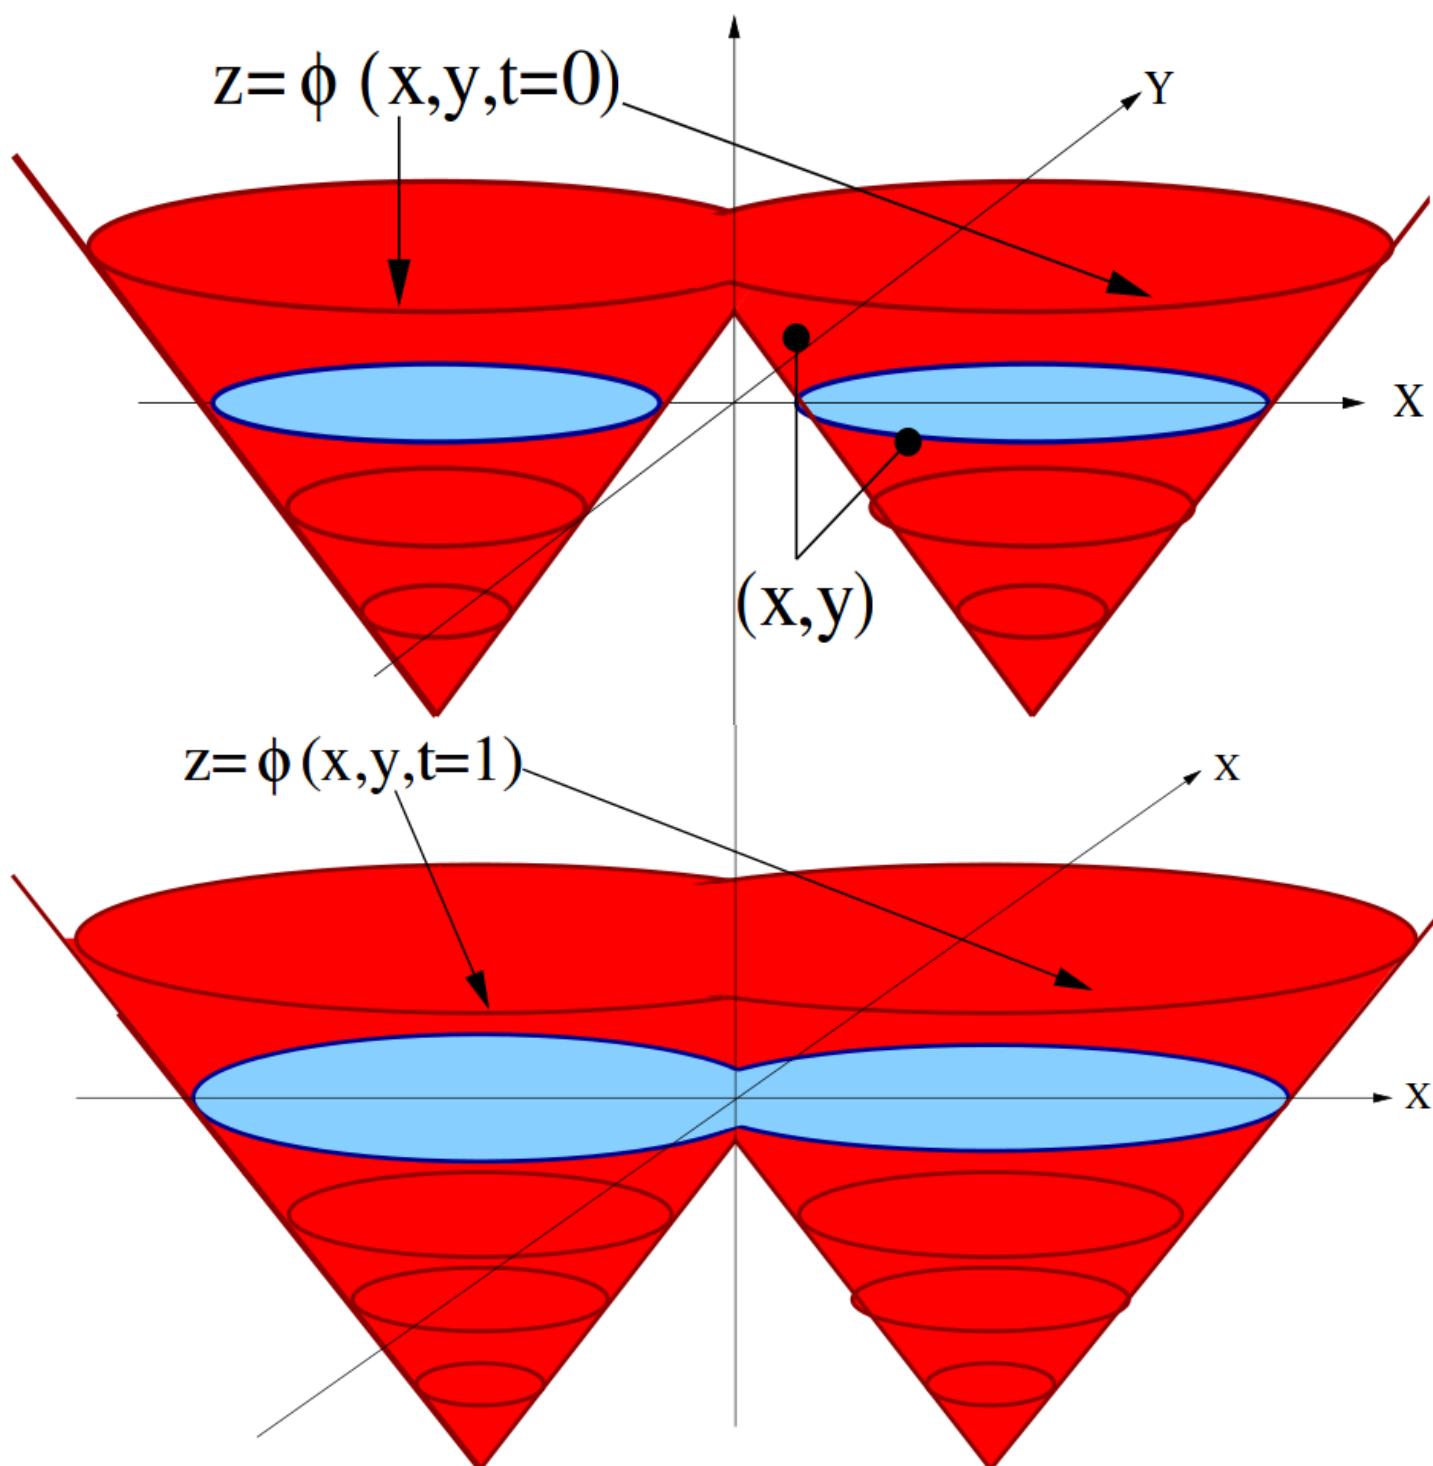
Level-Se

- Key ideas

- (2) Surface evc

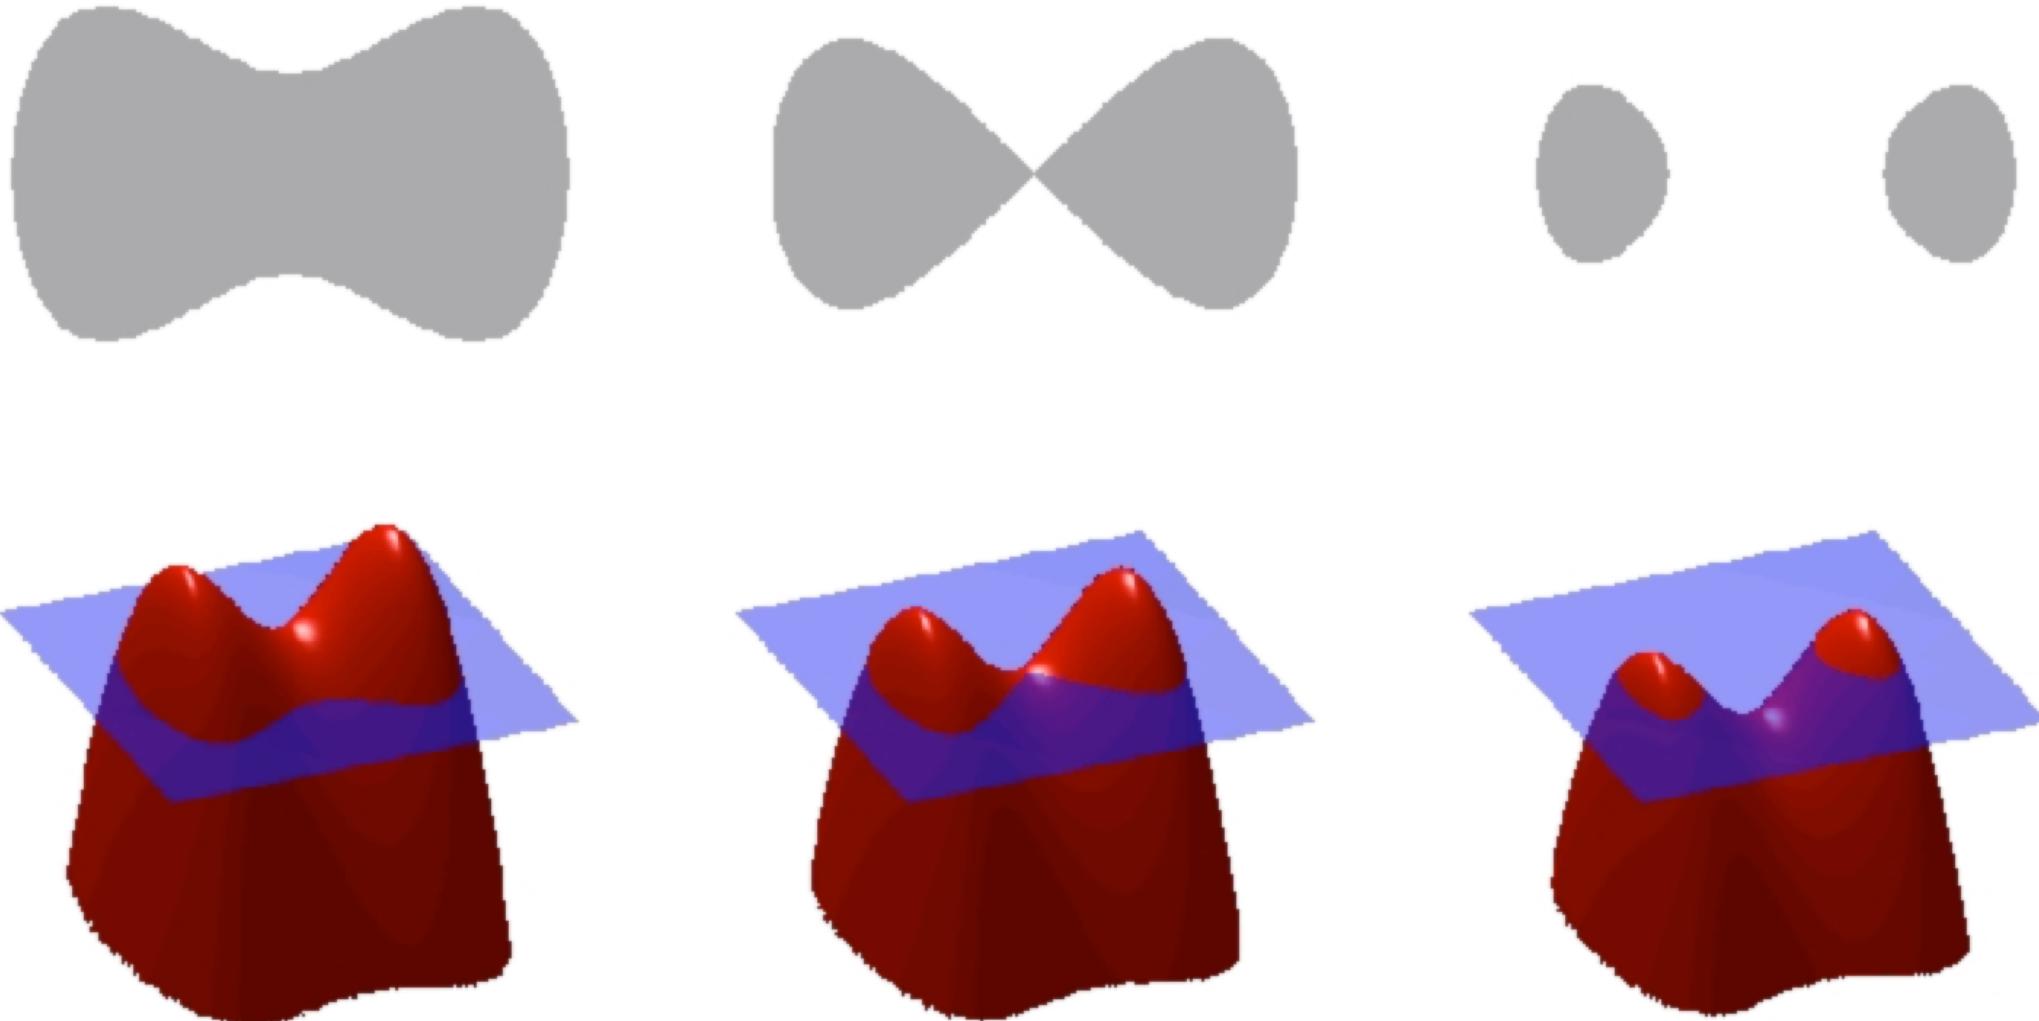
- Modify distance

Level Set Function z



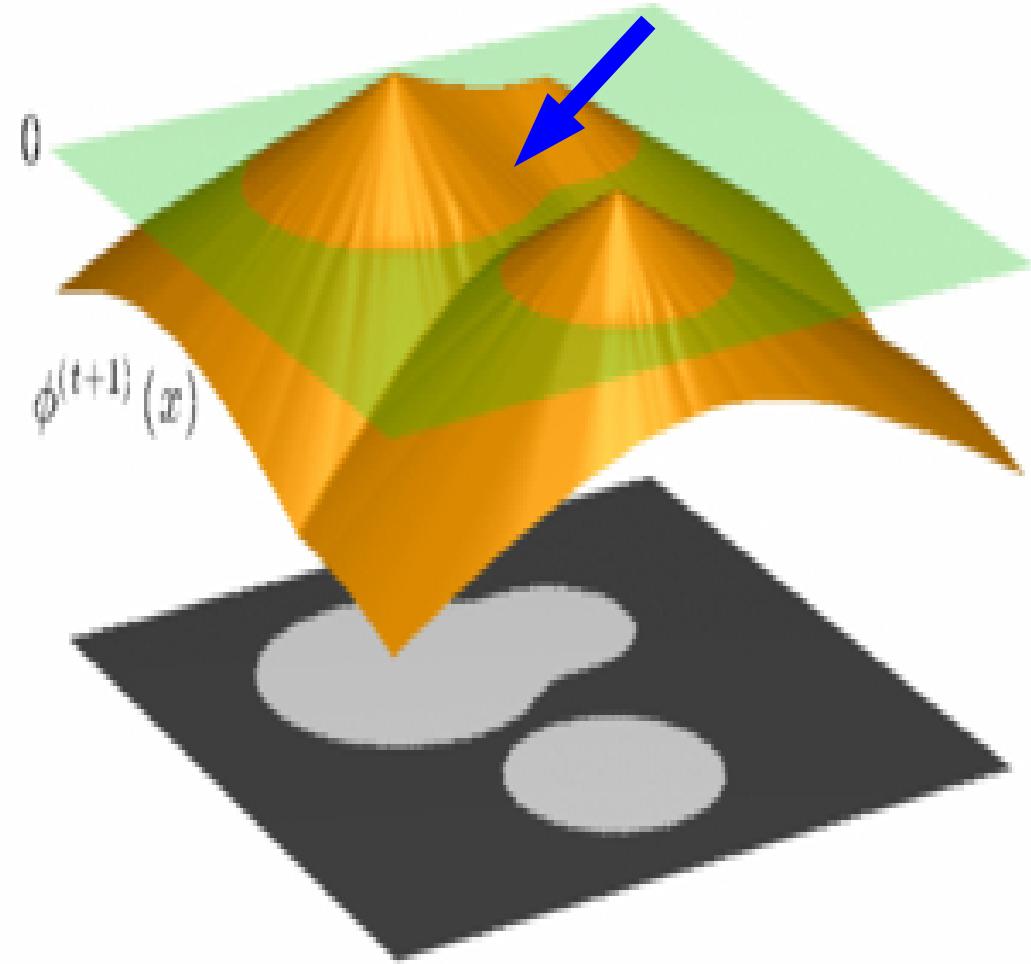
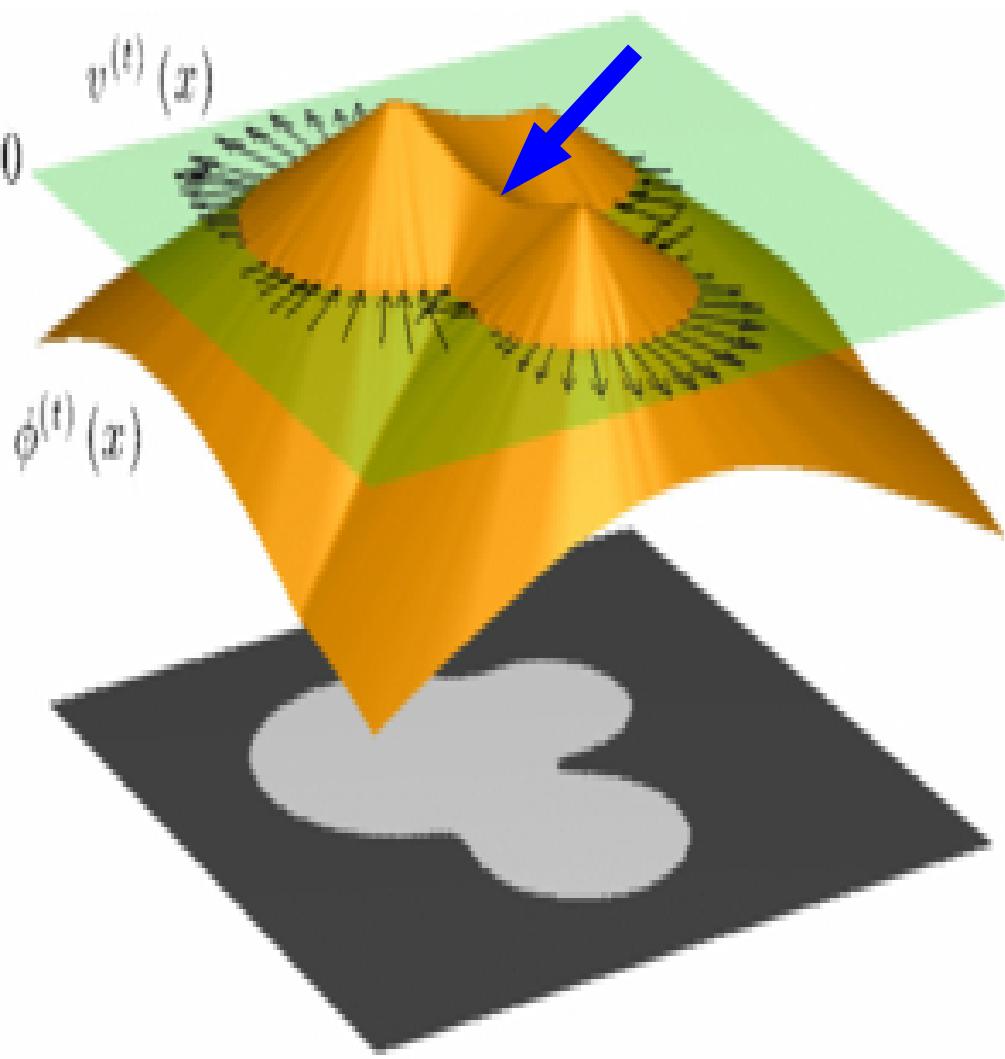
Level-Set Based Segmentation

- Key ideas
 - (2) Surface evolution
 - Modify distance transform to move surface



Level-Set Based Segmentation

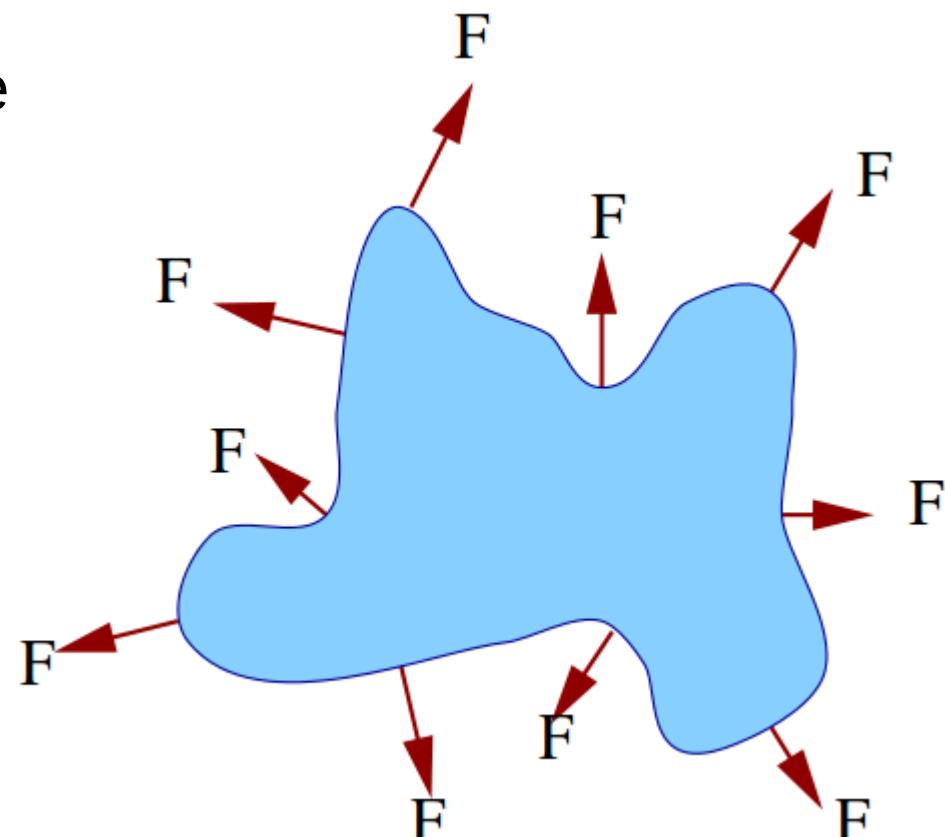
- Key ideas
 - (2) Surface evolution
 - Modify distance transform to move surface



Level-Set Based Segmentation

- Curve evolution / flow

- Each point moves according to a **force along its normal**
- Don't consider / define forces along tangent directions
 - Why ?
 - Such forces only **slide** points along curve
 - Such forces don't change curve

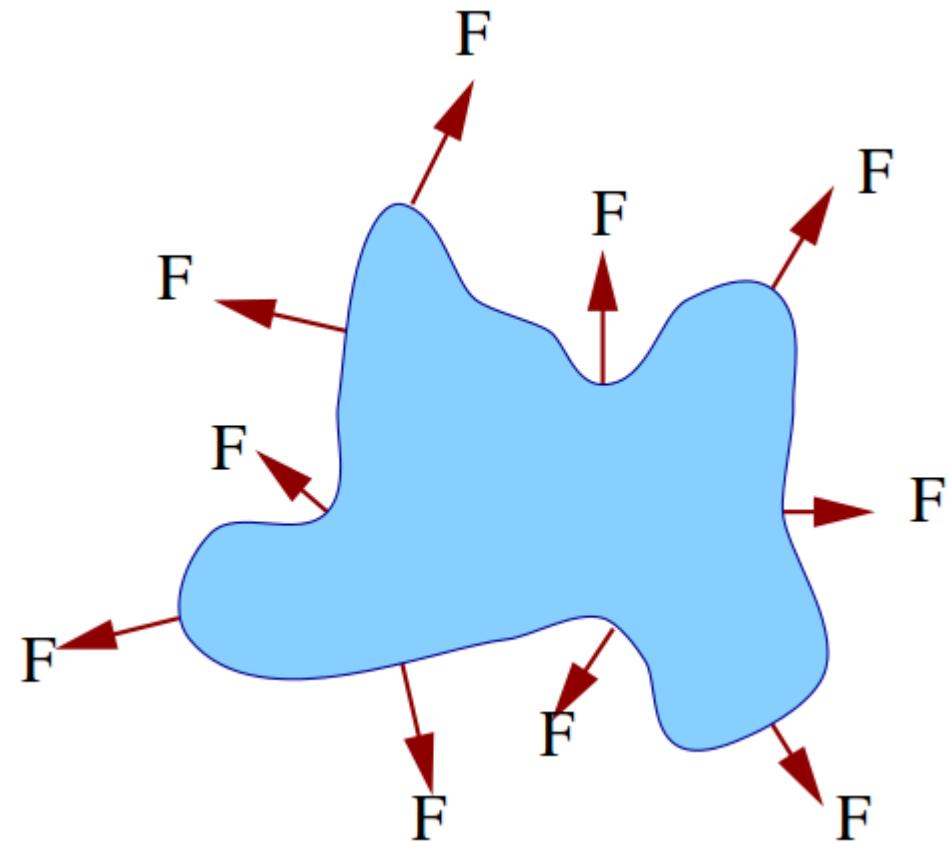


Level-Set Based Segmentation

- Advantages over mesh representations

- Avoid detecting intersections

- Self intersections
 - Intersections between components

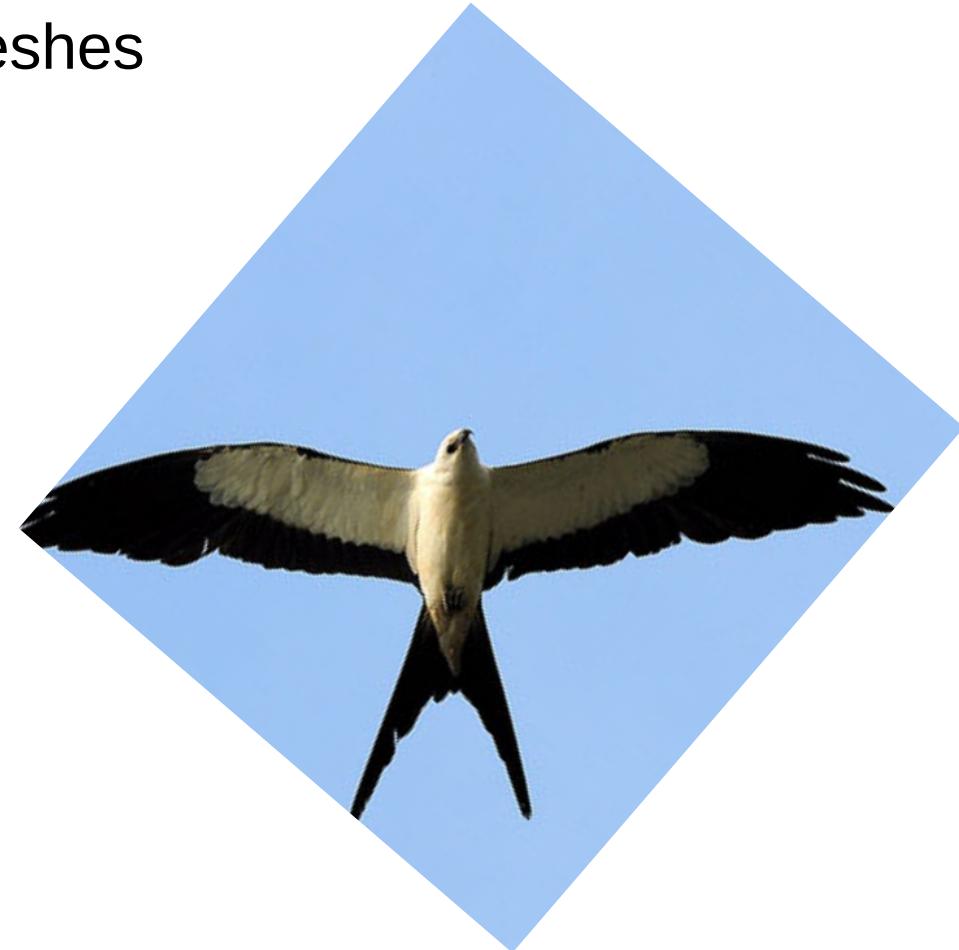
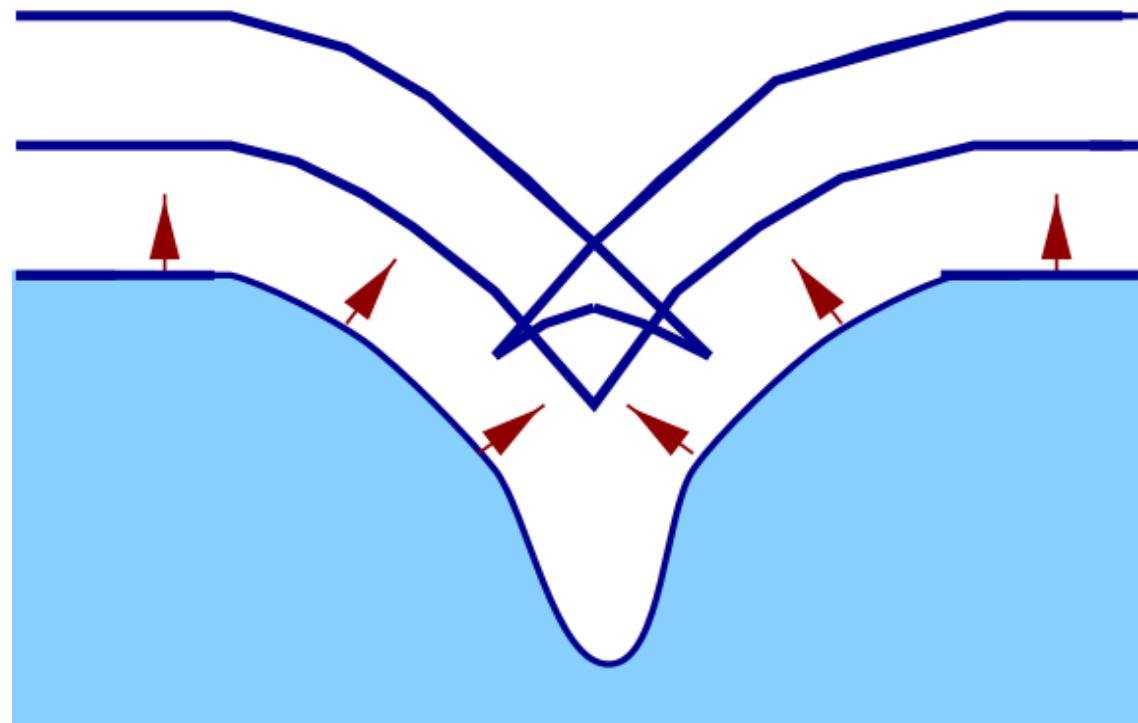


- Easy to handle splitting of 1 contour into 2 contours
 - Easy to handle merging of 2 contours into 1 contour

Level-Set Based Segmentation

- **Self intersection problem**

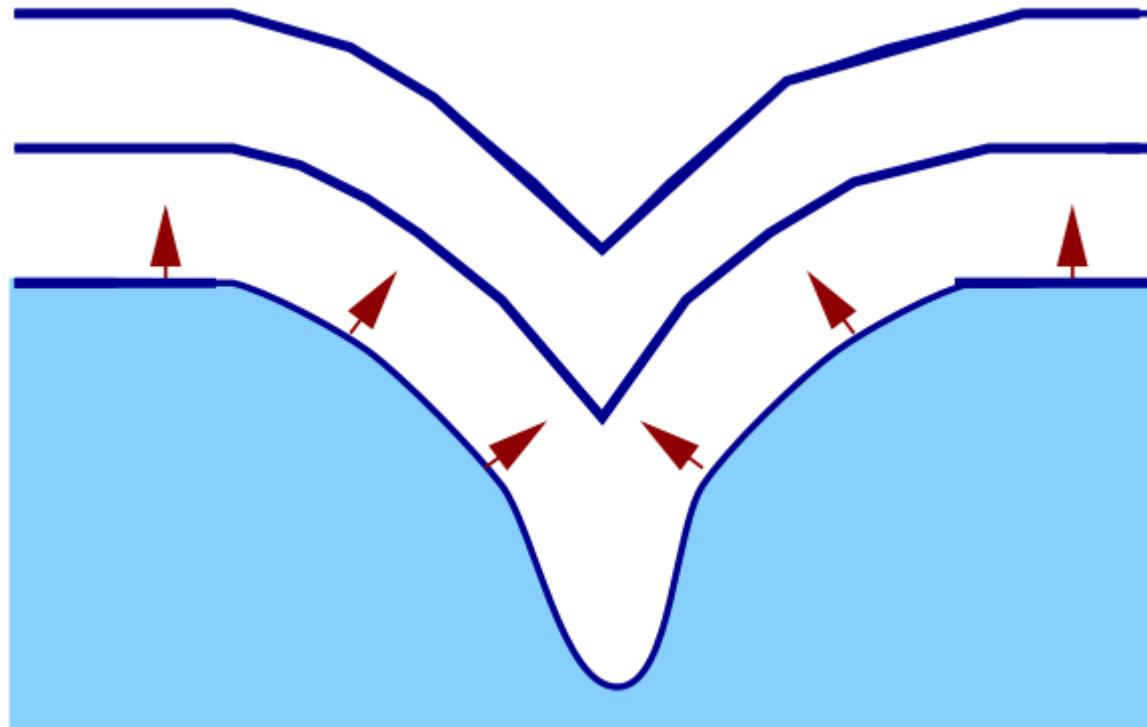
- Consider a sinusoid-like curve
- Each point moves along its normal with same force F
- “Swallowtail” problem
 - More difficult to handle with meshes



Level-Set Based Segmentation

- **Avoiding self intersections**

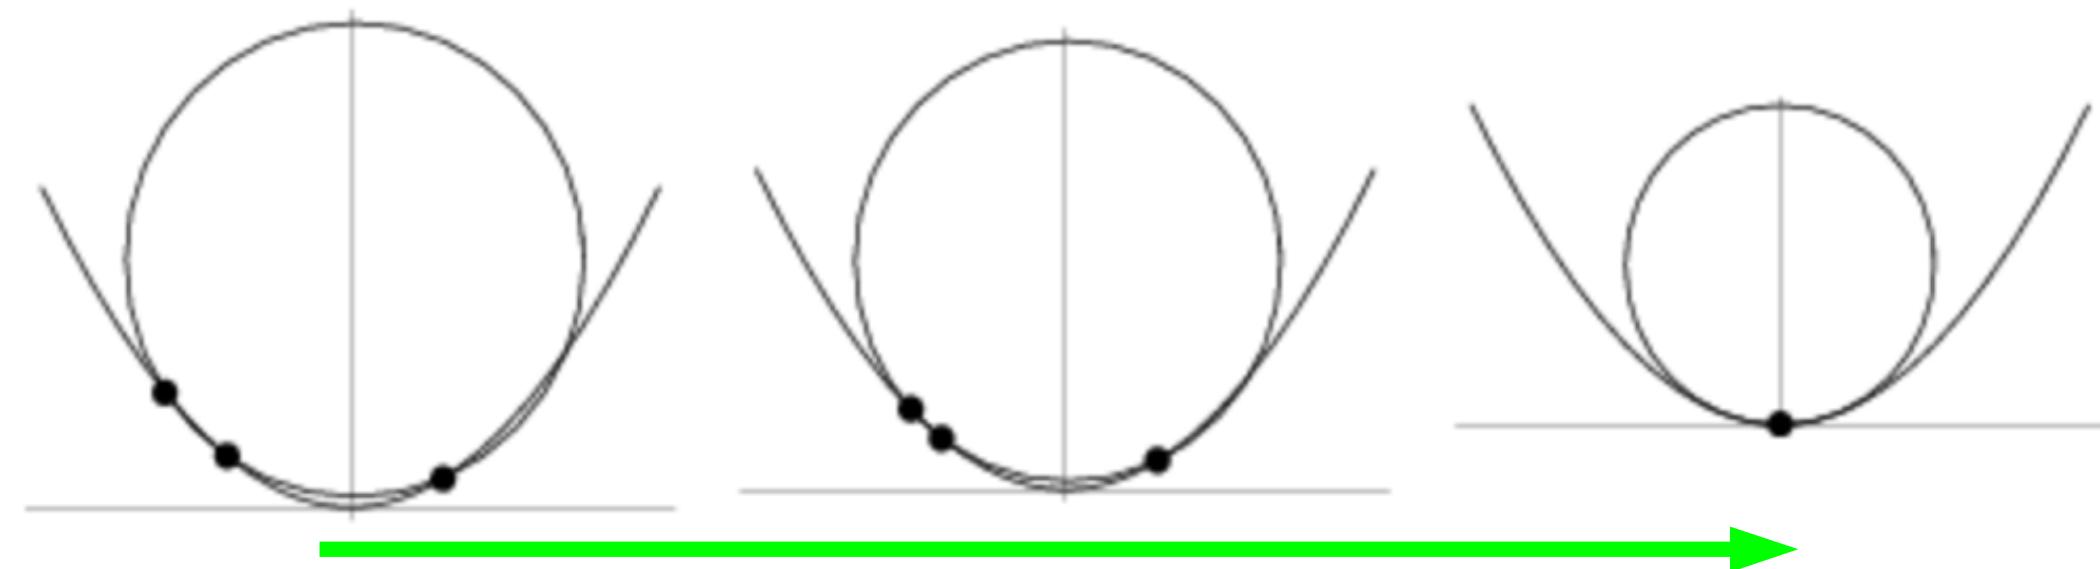
- This is more desirable:



- Must **avoid corners** too → undefined derivatives (e.g., normal direction)

Level-Set Based Segmentation

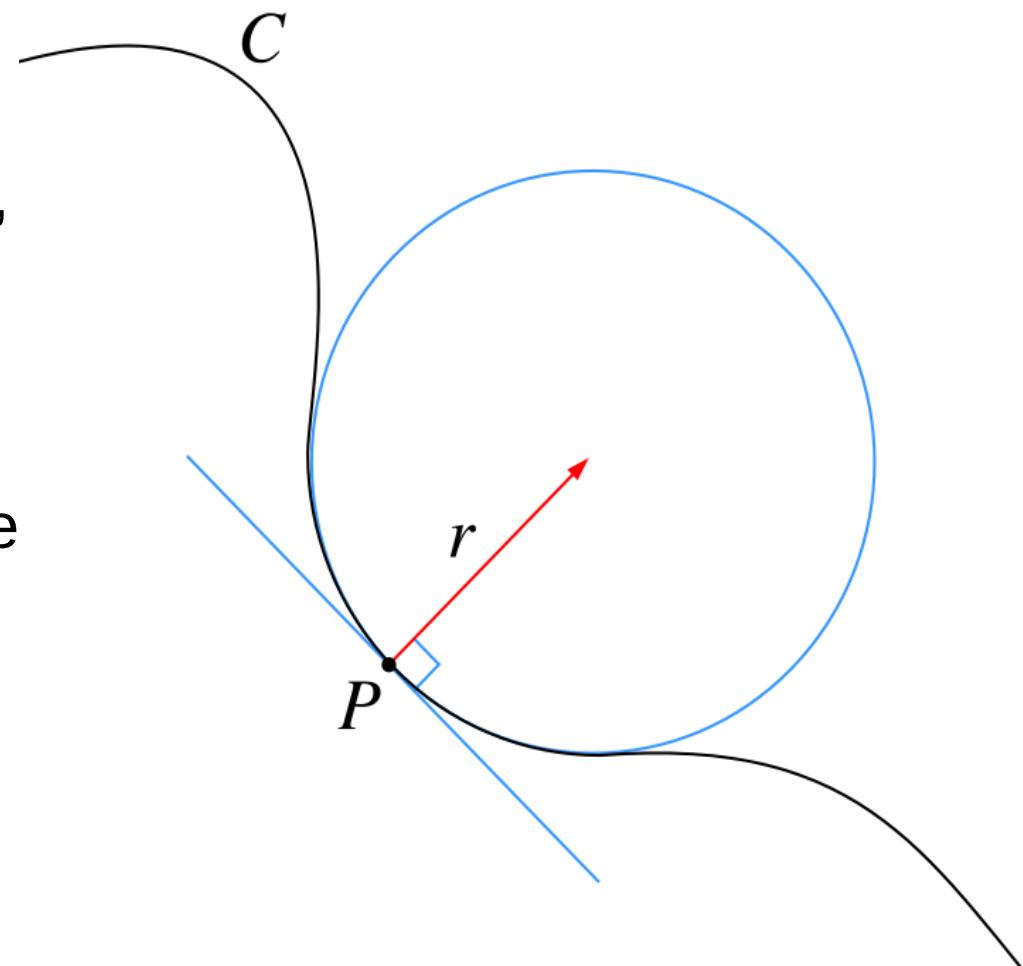
- Curvature of a (smooth) curve, at point P
= Curvature of its *osculating circle* at point P
 - **Osculating circle**, at P
 - Let point on 2D curve = $P = (f(t), g(t))$, where t = parameter
 - Limit of a sequence of circles passing through points corresponding to t_1, t, t_2 (s.t. $t_1 < t < t_2$) as $t_1 \rightarrow t$ and $t_2 \rightarrow t$



Level-Set Based Segmentation

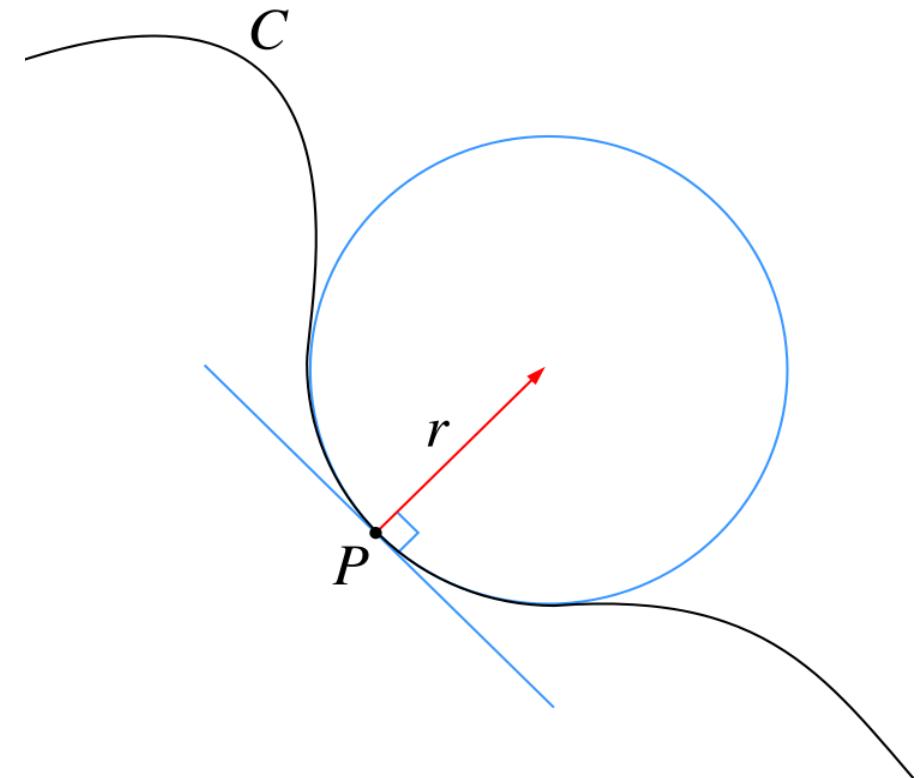
- Curvature of a (smooth) curve, at point P
= Curvature of its *osculating circle* at point P
 - Osculating circle, at P

- One among circles, all tangent to curve tangent at P, which **best approximates the curve near P**
- Center = center of curvature
- Radius r = radius of curvature
- **Curvature of curve = $1 / r$**



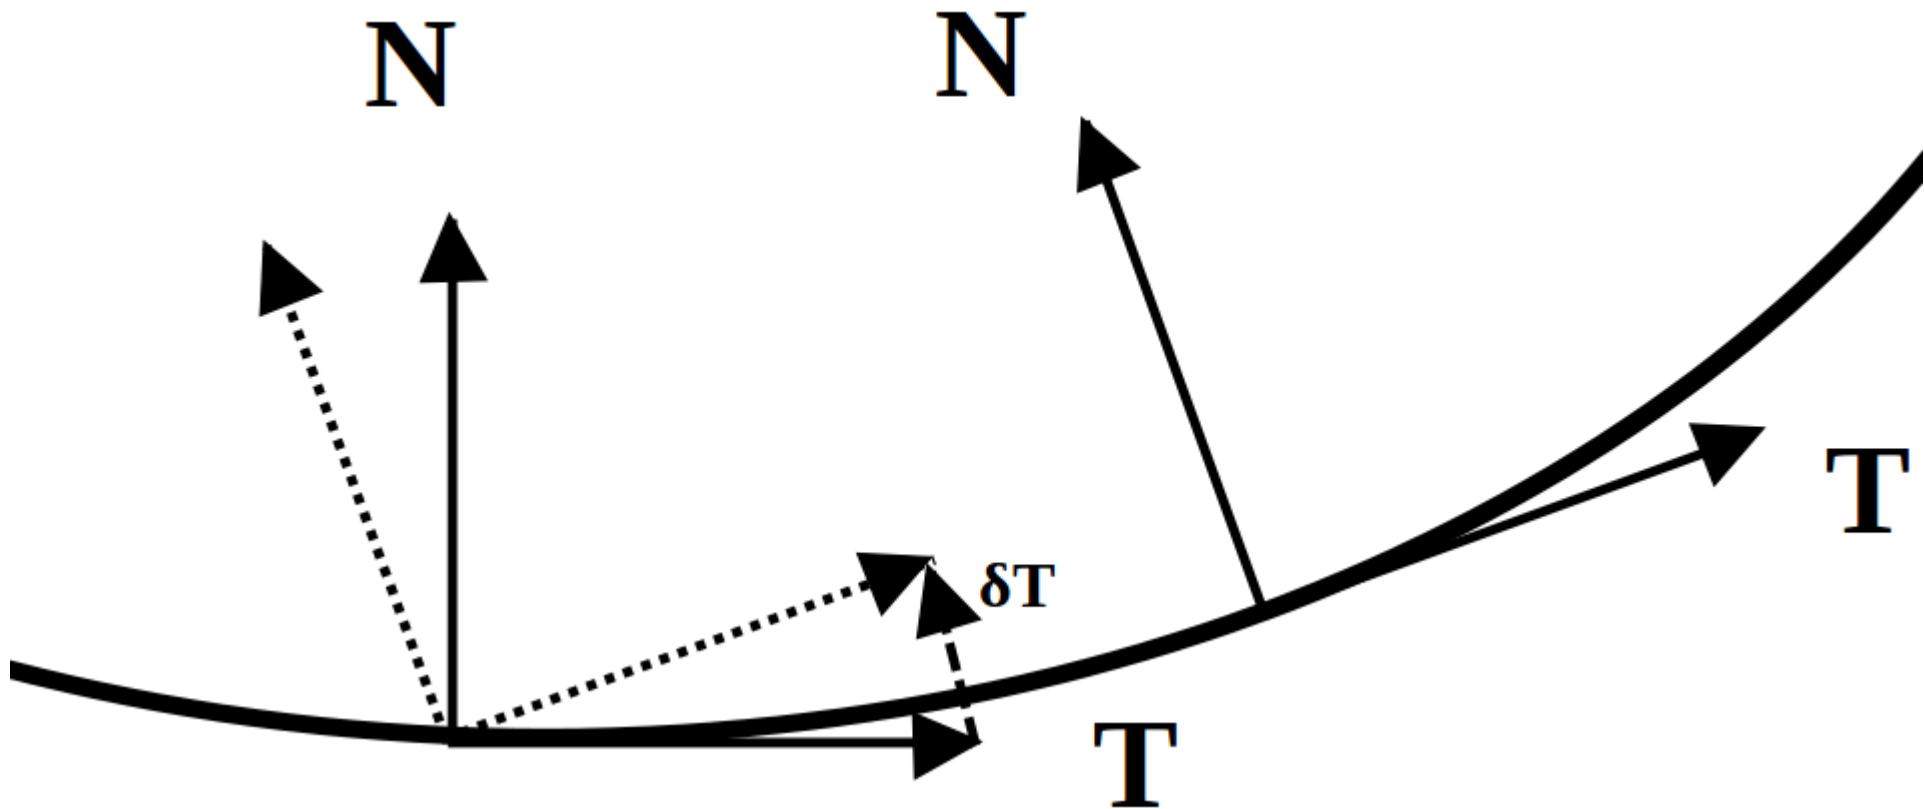
Level-Set Based Segmentation

- Curvature of a (smooth) curve
 - What is curvature of a **circle** at any point ?
 - Radius of osculating circle = ?
 - What is the curvature of a **line** at any point ?
 - Radius of osculating circle = ?



Level-Set Based Segmentation

- Curvature of a (smooth) curve
 - = magnitude of rate of change of tangent vector (or normal vector)
 - = $\| dT(s) / ds \|$, where s = arc length (NOT parameter)

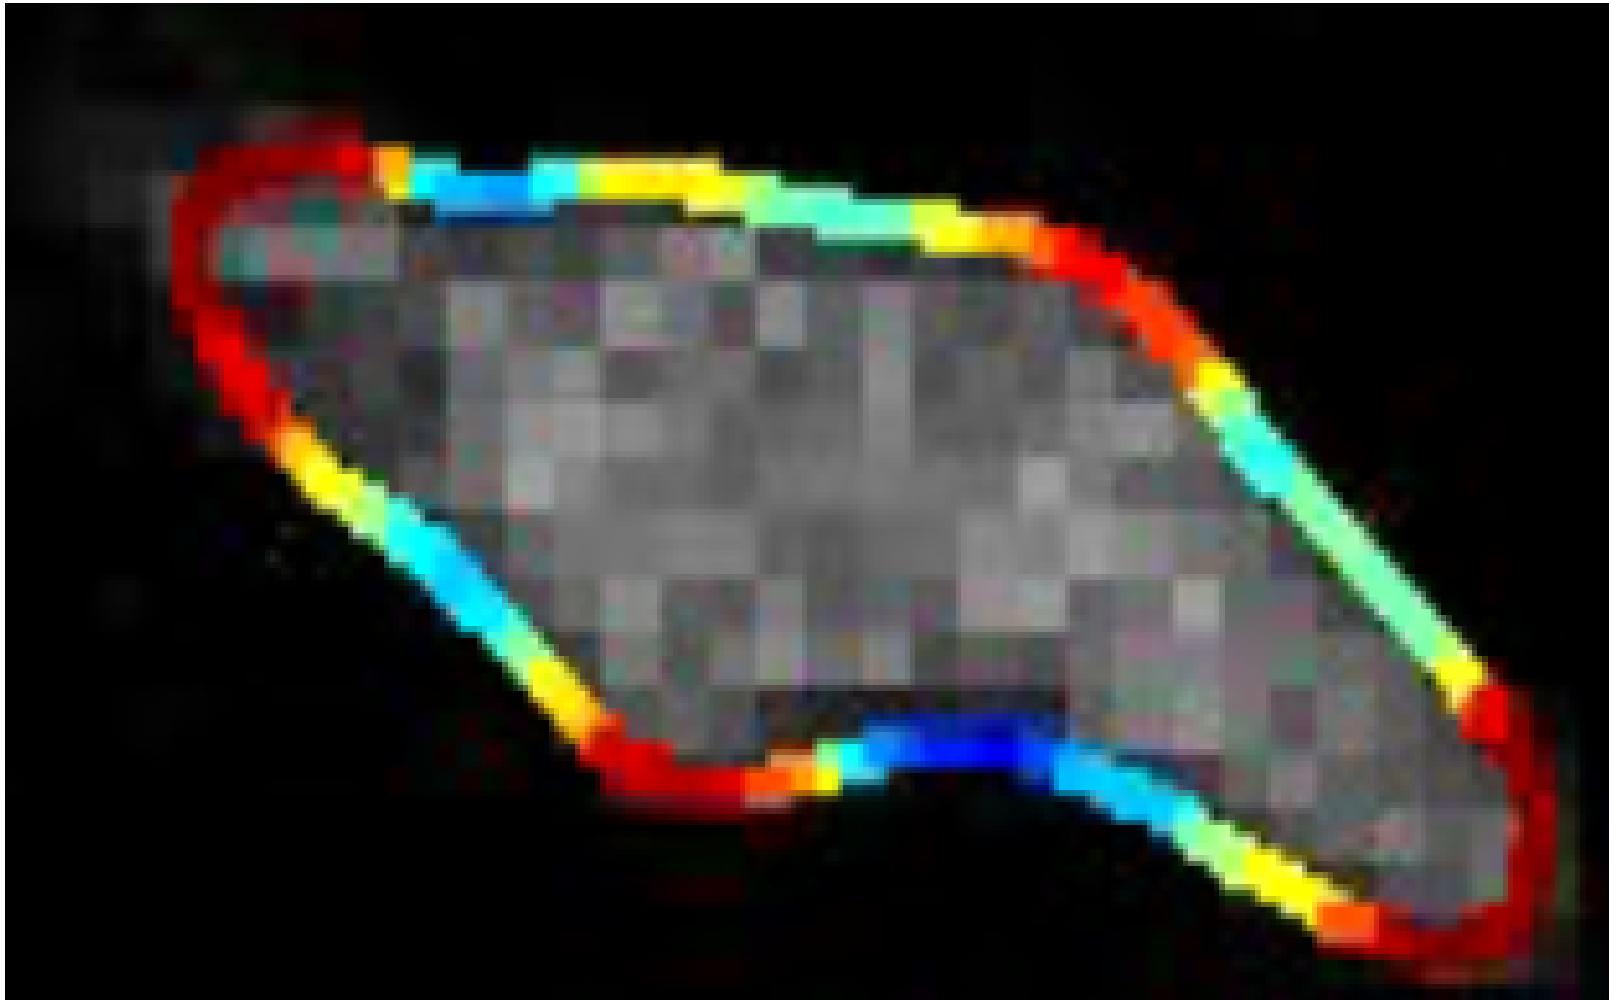


Level-Set Based Segmentation

- Curvature of a (smooth) curve
 - Signed curvature: depends on orientation of curve
 - e.g., whether curve goes from left → right or vice versa
 - Sinusoidal curve has points with positive and negative curvature

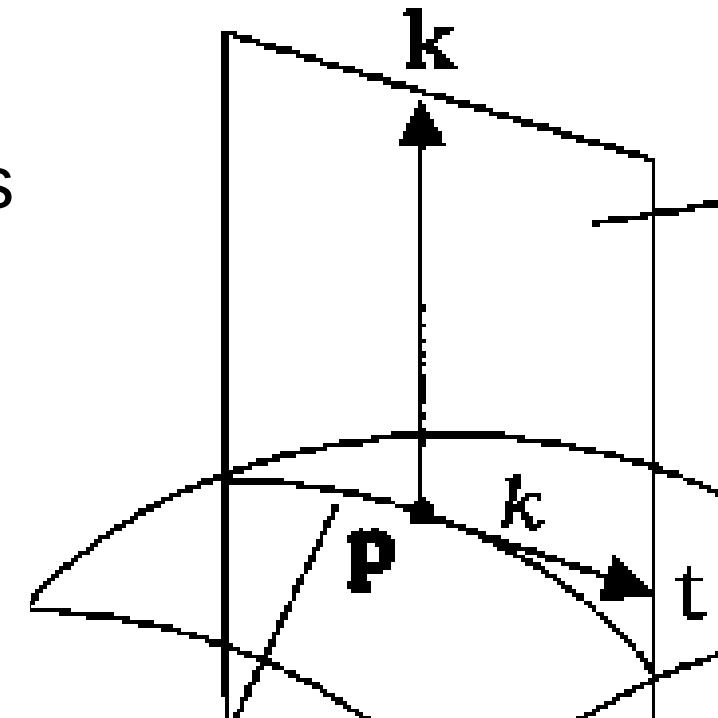
Level-Set Based Segmentation

- Curvature of a (smooth) curve
 - <http://upload.wikimedia.org/wikipedia/commons/3/3e/Cell-Shape-Dynamics-From-Waves-to-Migration-pcbi.1002392.s007.ogv>
 - Red = large positive
 - Blue = large negative



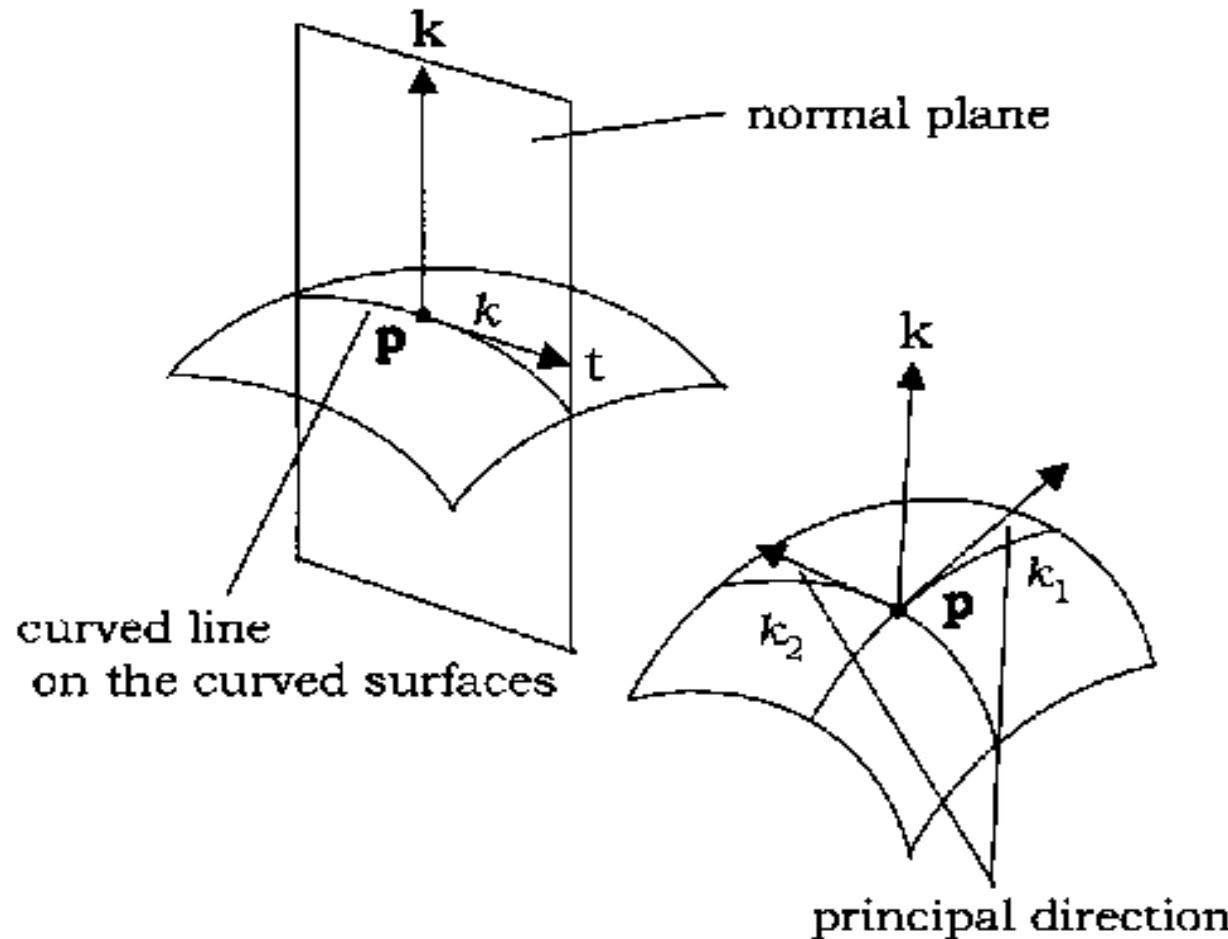
Level-Set Based Segmentation

- Curvature of a (smooth) surface
 - Consider a 1D curve on 2D surfaces passing through P
 - Consider a tangent vector 't' at P
 - Many tangent vectors
 - Consider a normal vector 'n' at P
 - Unique upto sign
 - Consider a plane passing through P, and containing a single 't' and 'n'
 - Consider curvature of curve on surface within this plane (at point P)
 - Normal curvature



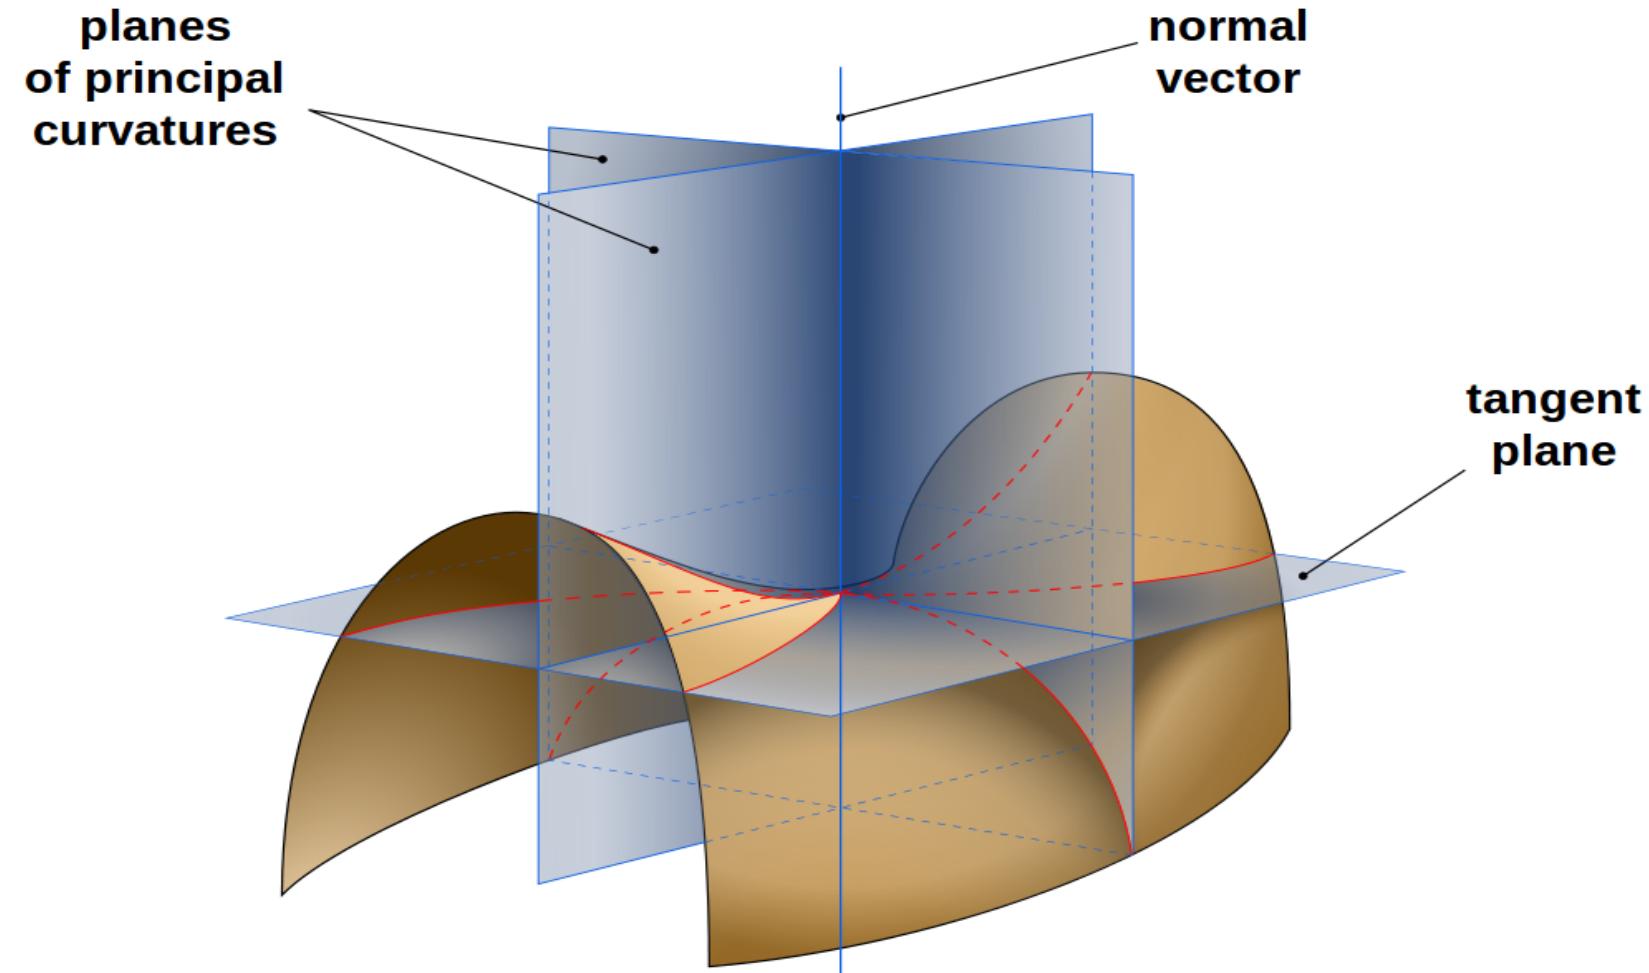
Level-Set Based Segmentation

- Curvature of a (smooth) surface
 - Over all possible tangent vectors:
 - Maximum curvature and minimum curvature are called **principal curvatures**



Level-Set Based Segmentation

- Curvature of a (smooth) surface
 - Principal curvatures
 - At saddle point in picture: $k_1 > 0, k_2 < 0, |k_1| > |k_2|$



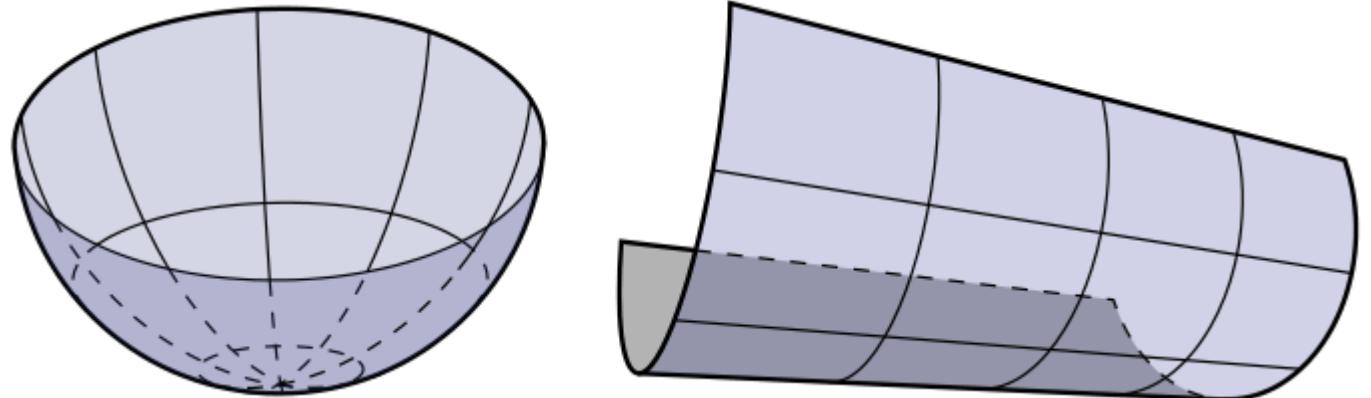
Level-Set Based Se

- Saddle



Level-Set Based Segmentation

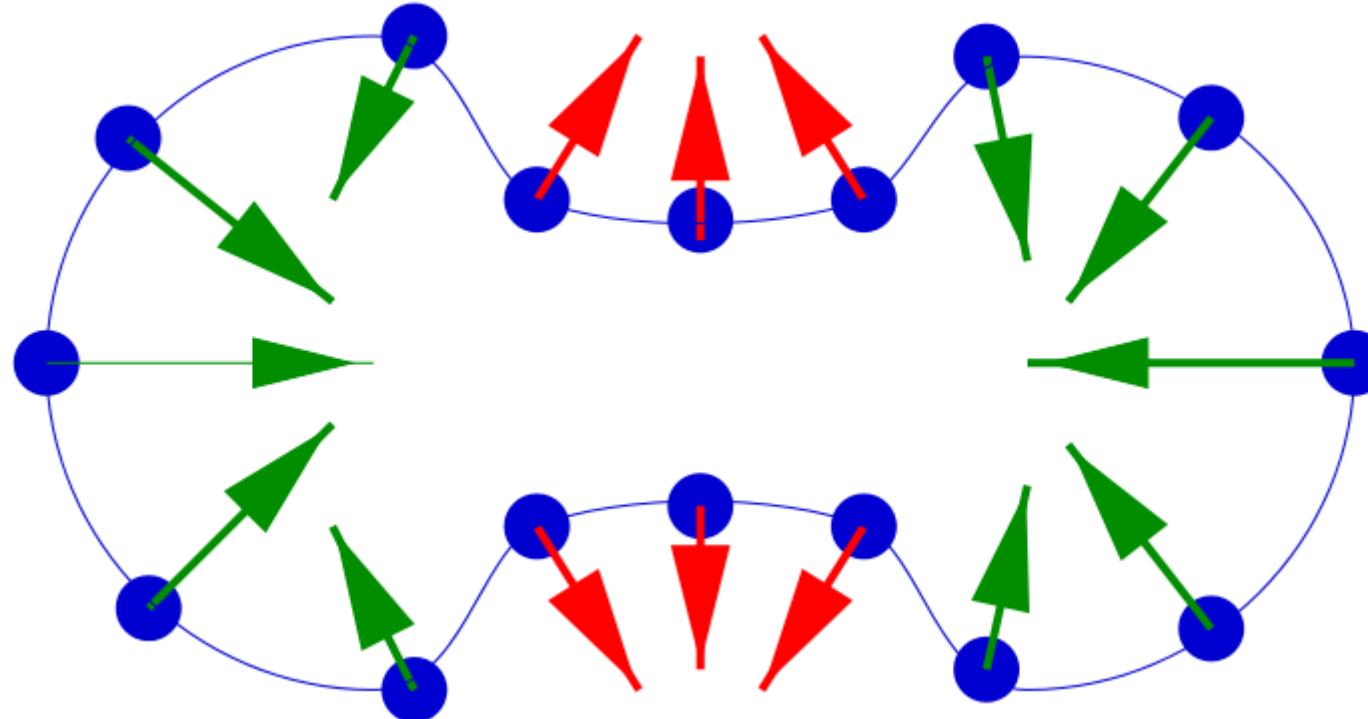
- Curvature of a (smooth) surface
 - Principal curvatures: k_1, k_2
 - **Gaussian curvature** = $k_1 * k_2$
 - **Mean curvature** = $(k_1 + k_2) / 2$
- What are the principal curvatures for any point on:
 - A sphere ?
 - A cylinder ?
 - A plane ?



Level-Set Based Segmentation

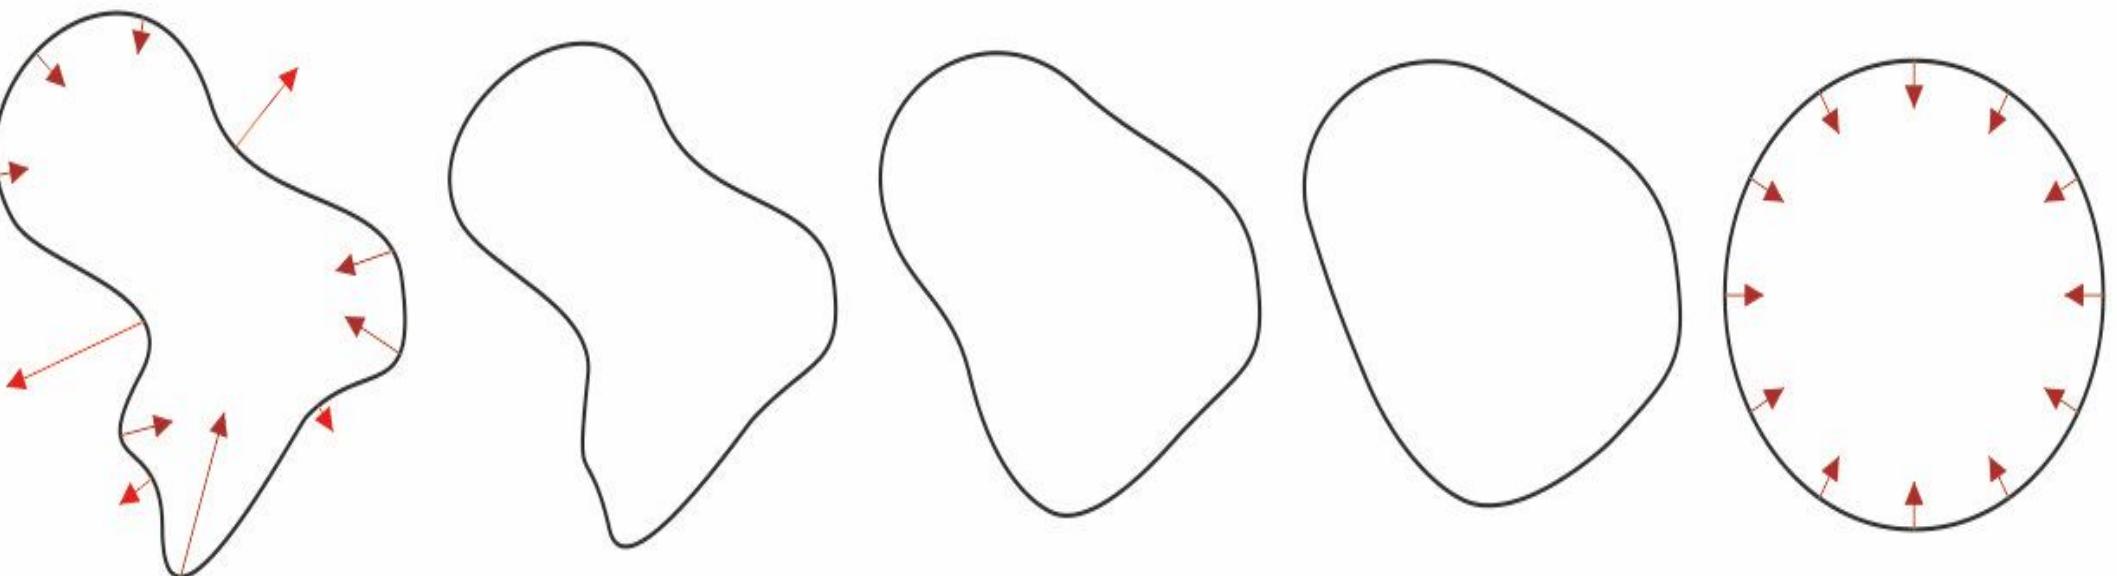
- Mean curvature flow

- Each point on curve moves along its normal driven by a force proportional to its **mean curvature**
 - Move along outward normal
 - Force = $(-1) * \text{signed mean curvature}$



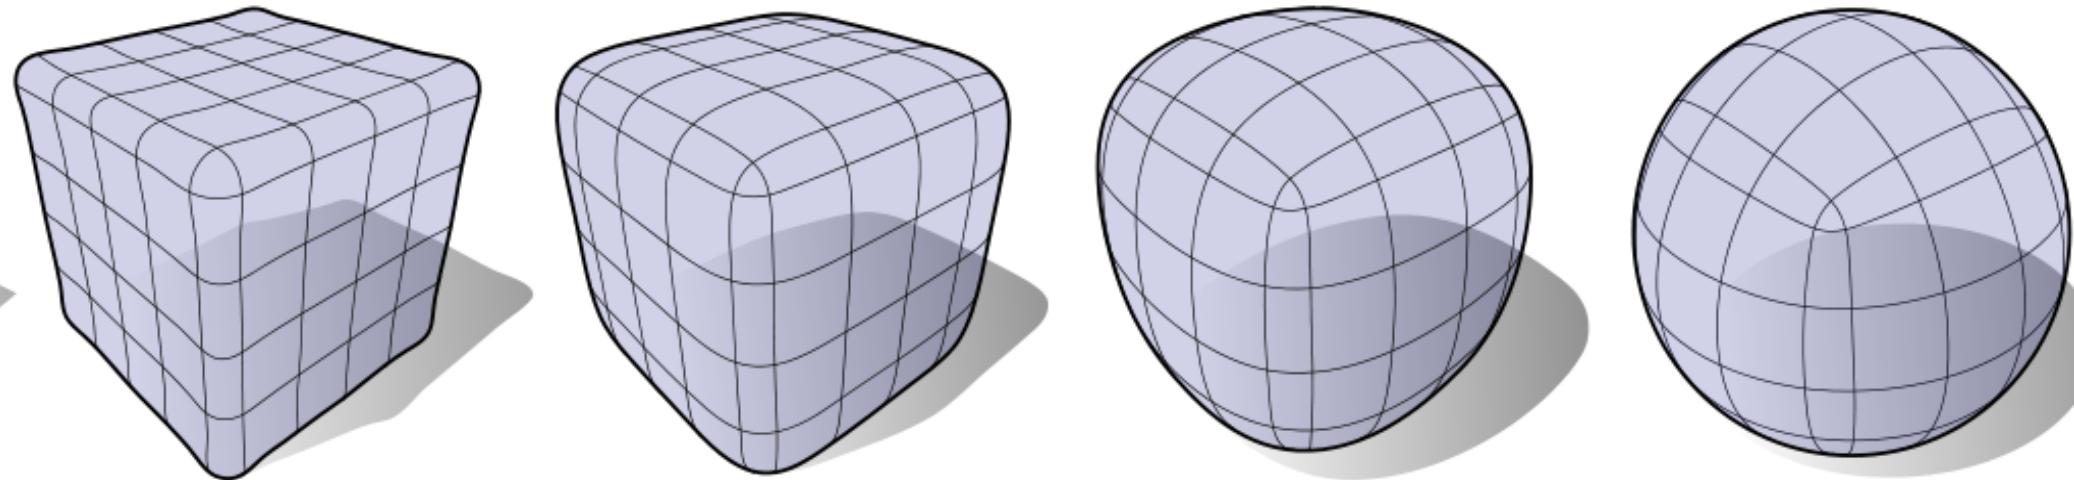
Level-Set Based Segmentation

- Mean curvature flow
 - Each point on curve moves (direction and magnitude) based on **mean curvature**



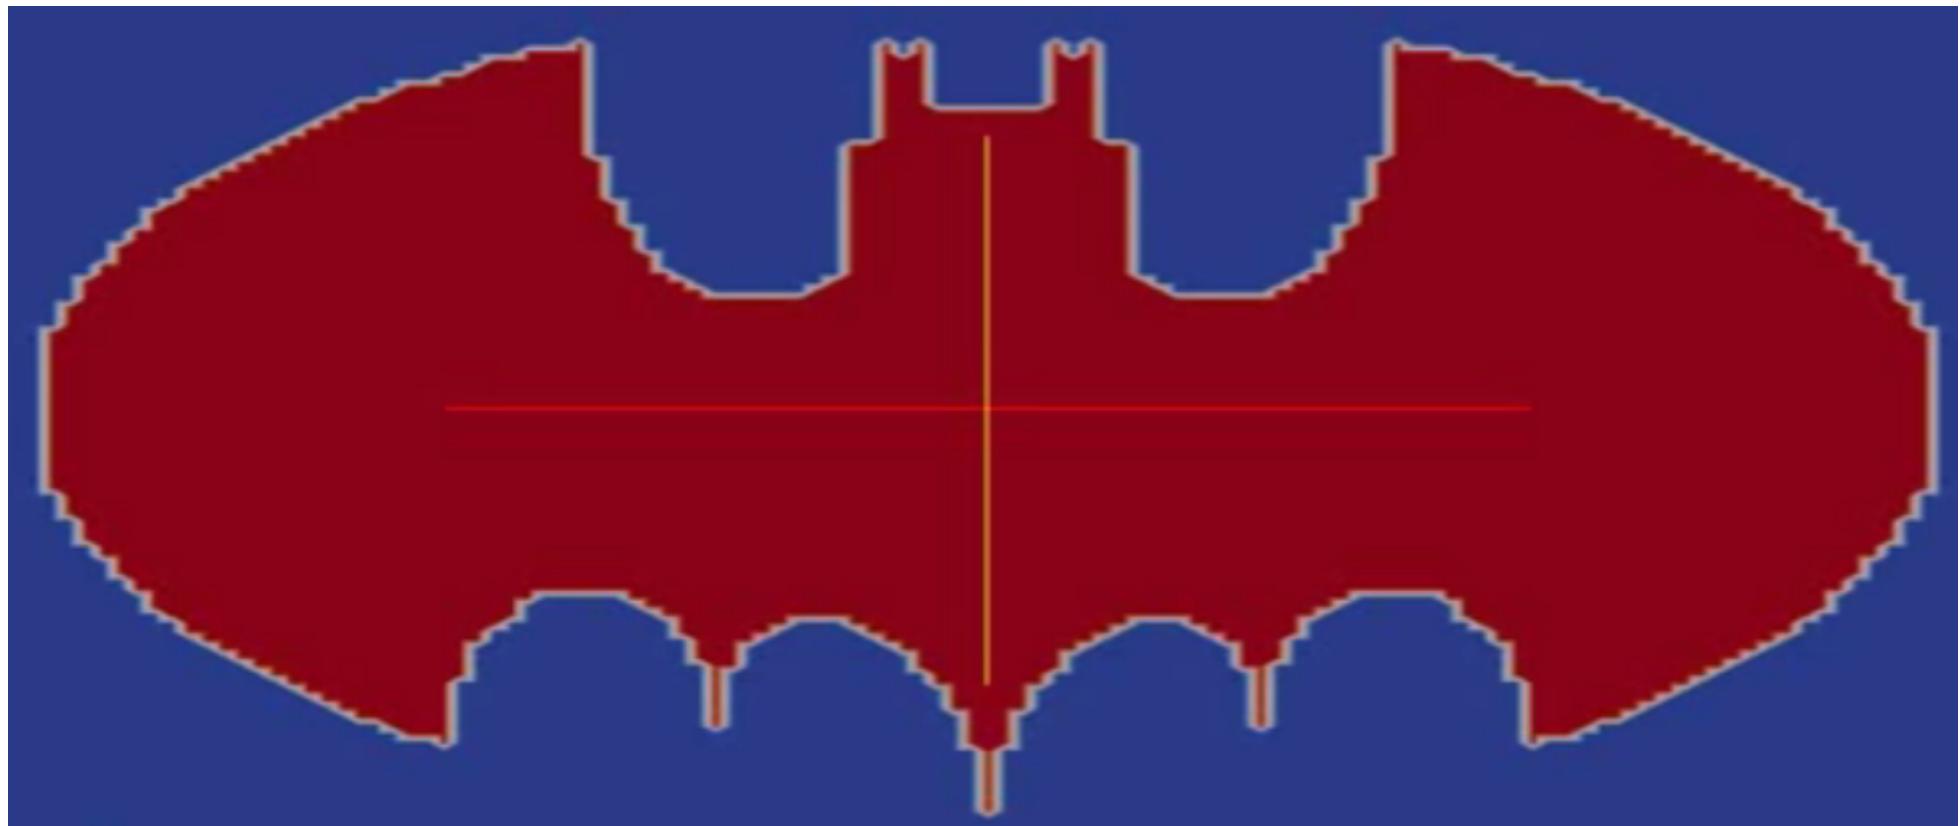
Level-Set Based Segmentation

- Mean curvature flow
 - Each point on curve moves (direction and magnitude) based on **mean curvature**



Level-Set Based Segmentation

- Mean curvature flow
 - Each point on curve moves (direction and magnitude) based on **mean curvature**
 - <https://www.youtube.com/watch?v=TeuLhhR7UxM>



Level-Set Based Segmentation

- Mean curvature flow
 - Theorem: **Every** simple closed curve (in 2D), evolving under mean curvature flow, collapses smoothly to a single point, without crossing over itself
 - Implications:
 - No matter how complicated / convoluted a curve might be ...
 - ... the curve will evolve into a circle-like curve

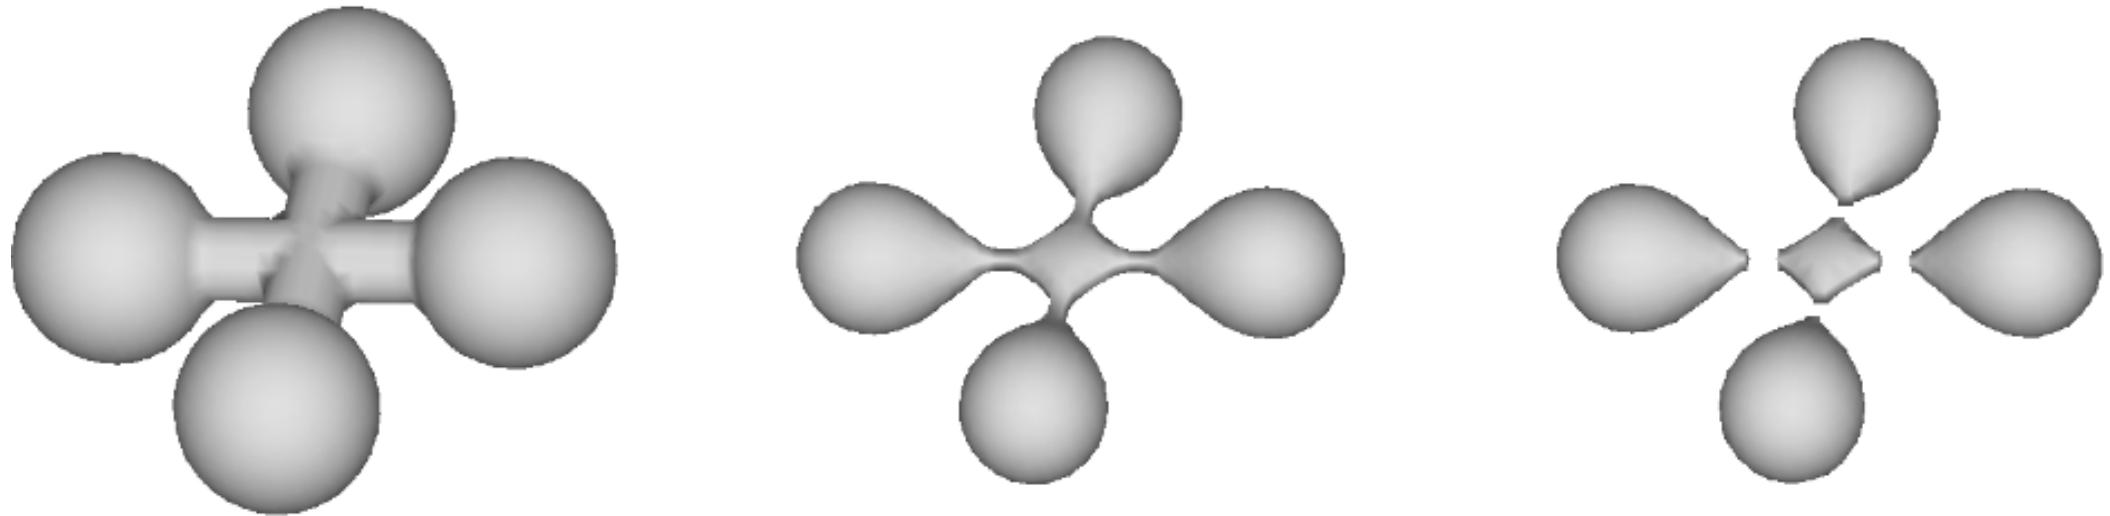
Level-Set Based Segmentation

- Mean curvature flow



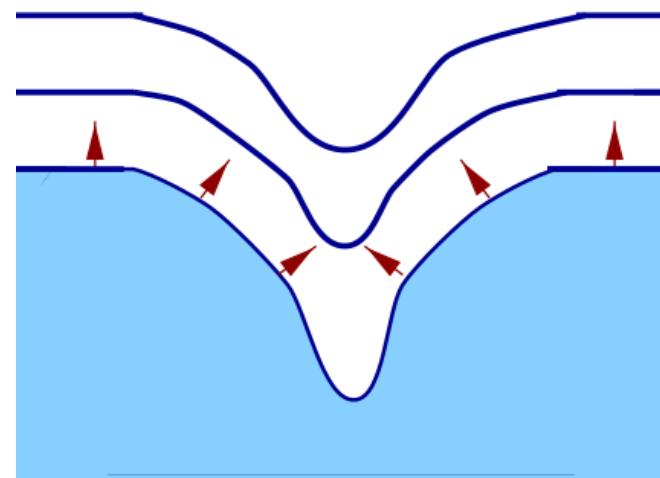
Level-Set Based Segmentation

- Mean curvature flow
 - Slightly different behavior in 3D

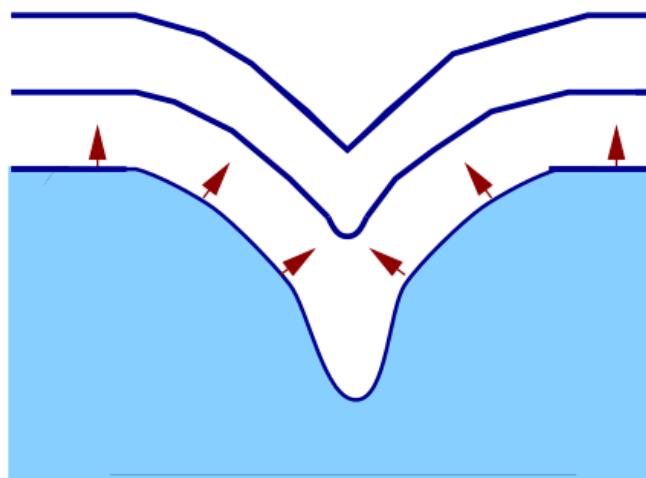


Level-Set Based Segmentation

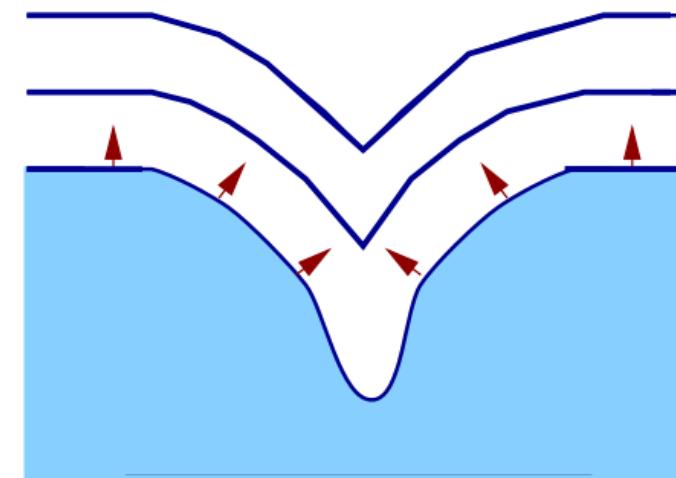
- Smooth out corners using mean curvature



Some Viscosity: $F = 1 - .1\kappa$



Less Viscosity: $F = 1 - .01\kappa$



No Viscosity: $F = 1$

Level-Set Based Segmentation

- Evolving curves
- Consider a subset Ω of \mathbb{R}^2
- Consider an oriented curve, evolving with time

$$\mathcal{C} \subset \Omega, \quad \mathcal{C} : \begin{cases} [0, 1] \times \mathbb{R}^+ \rightarrow \mathbb{R}^2 \\ (q, t) \mapsto \mathcal{C}(q, t) \end{cases}$$

- q = parametrization
- t = time
- Embed evolving curve C , at any time t ,
as zero level set of a time-varying function $\Phi(x, t)$
 - $\Phi : \Omega \times \mathbb{R}^+ \rightarrow \mathbb{R}$, Φ takes negative values inside curve
- Question: how to evolve Φ ?

Level-Set Based Segmentation

- Evolving curves
- Design principle: $\Phi(C(q,t), t) = 0$
 - At any time t , Φ should take value 0 when evaluated at points located on curve $C(q,t)$ at **same time t**
- Differentiate $\Phi(C(.,.), .)$ w.r.t. t
 - 1) $\Phi(C(q,t'), t'')$ changes because t'' changes, keeping location $C(q,t')$ fixed
 - 2) $\Phi(C(q,t'), t'')$ changes because location $C(q,t')$ changes, keeping time t'' fixed

$$\frac{\partial \Phi}{\partial t}(C(q, t), t) + \langle \nabla \Phi(C(q, t), t), \frac{\partial C}{\partial t}(q, t) \rangle = 0$$

Level-Set Based Segmentation

- Evolving curves

$$\frac{\partial \Phi}{\partial t}(\mathcal{C}(q, t), t) + \langle \nabla \Phi(\mathcal{C}(q, t), t), \frac{\partial \mathcal{C}}{\partial t}(q, t) \rangle = 0$$

- For a curve where each point moves in outward normal direction based on a force of magnitude F

$$\frac{\partial \mathcal{C}}{\partial t}(q, t) = F \vec{n}(q, t)$$

- where $n(q, t)$ is outward **unit normal** to curve at point parameterized by q at time t
- But, we know, $n(q, t) = \nabla \Phi(C(q, t)) / \| \nabla \Phi(C(q, t)) \|$

Level-Set Based Segmentation

- Evolving curves

$$\frac{\partial \Phi}{\partial t}(\mathcal{C}(q, t), t) + F|\nabla \Phi(\mathcal{C}(q, t), t)| = 0$$

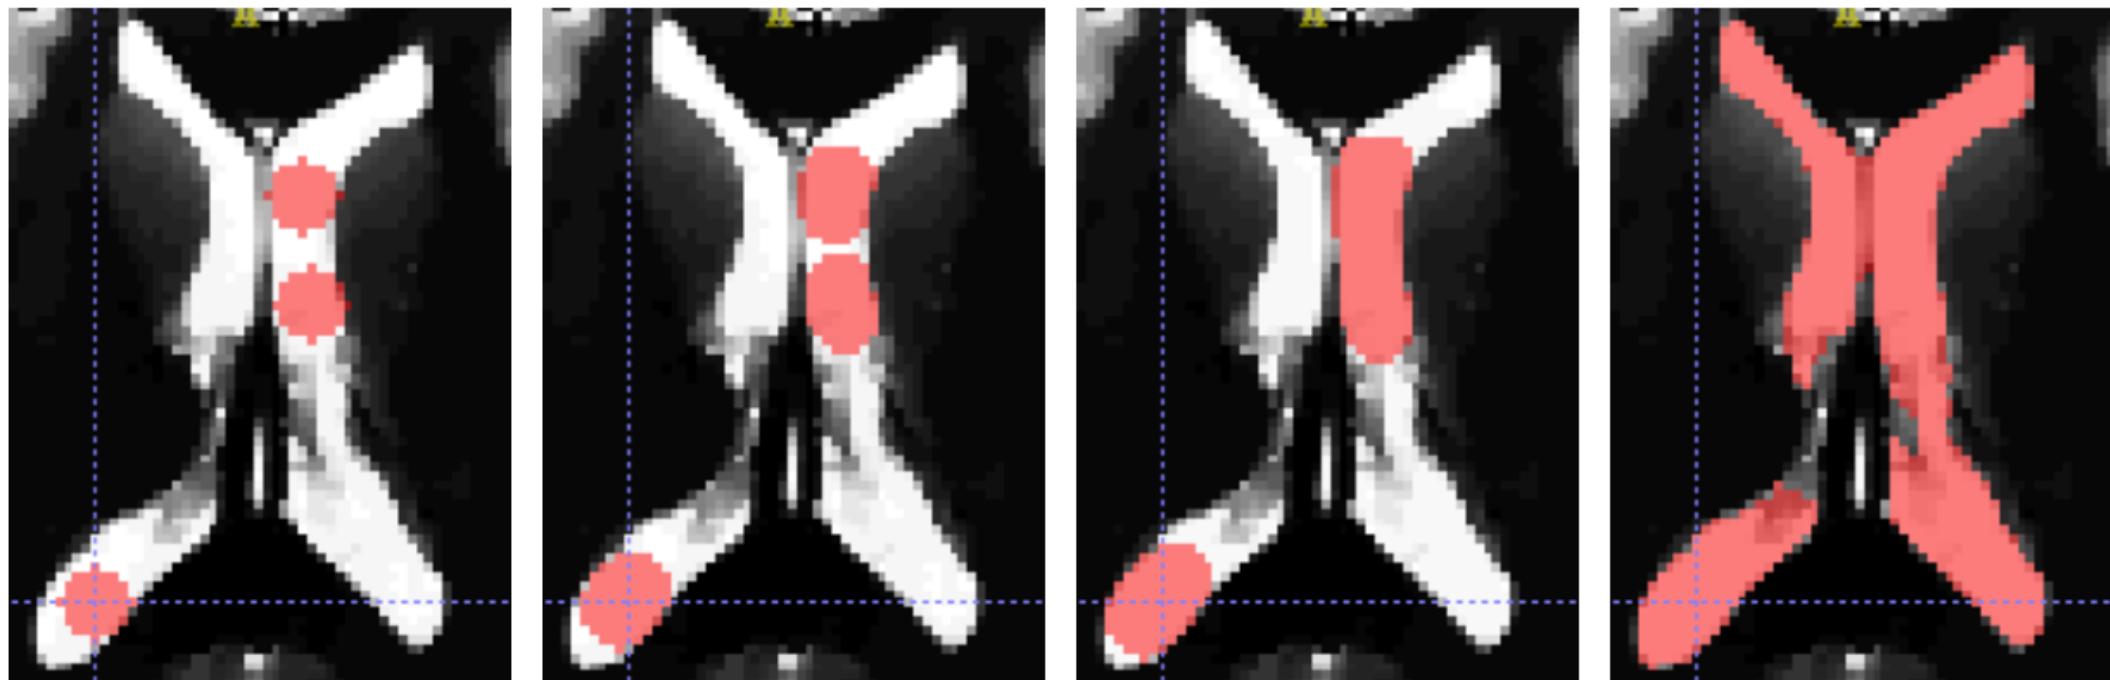
- This equation can be considered as defined on the whole domain $\Omega \times \mathbb{R}^+$
 - Requires specification of forces F on entire domain (instead of just on one level set)
 - And $\Phi(.,t)$ is maintained as a distance transform
- This gives the PDE:

$$\frac{\partial \Phi}{\partial t} + F|\nabla \Phi| = 0, \text{ on } \Omega \times \mathbb{R}^+$$

- with initial condition $\Phi(x, 0) = \Phi_0$

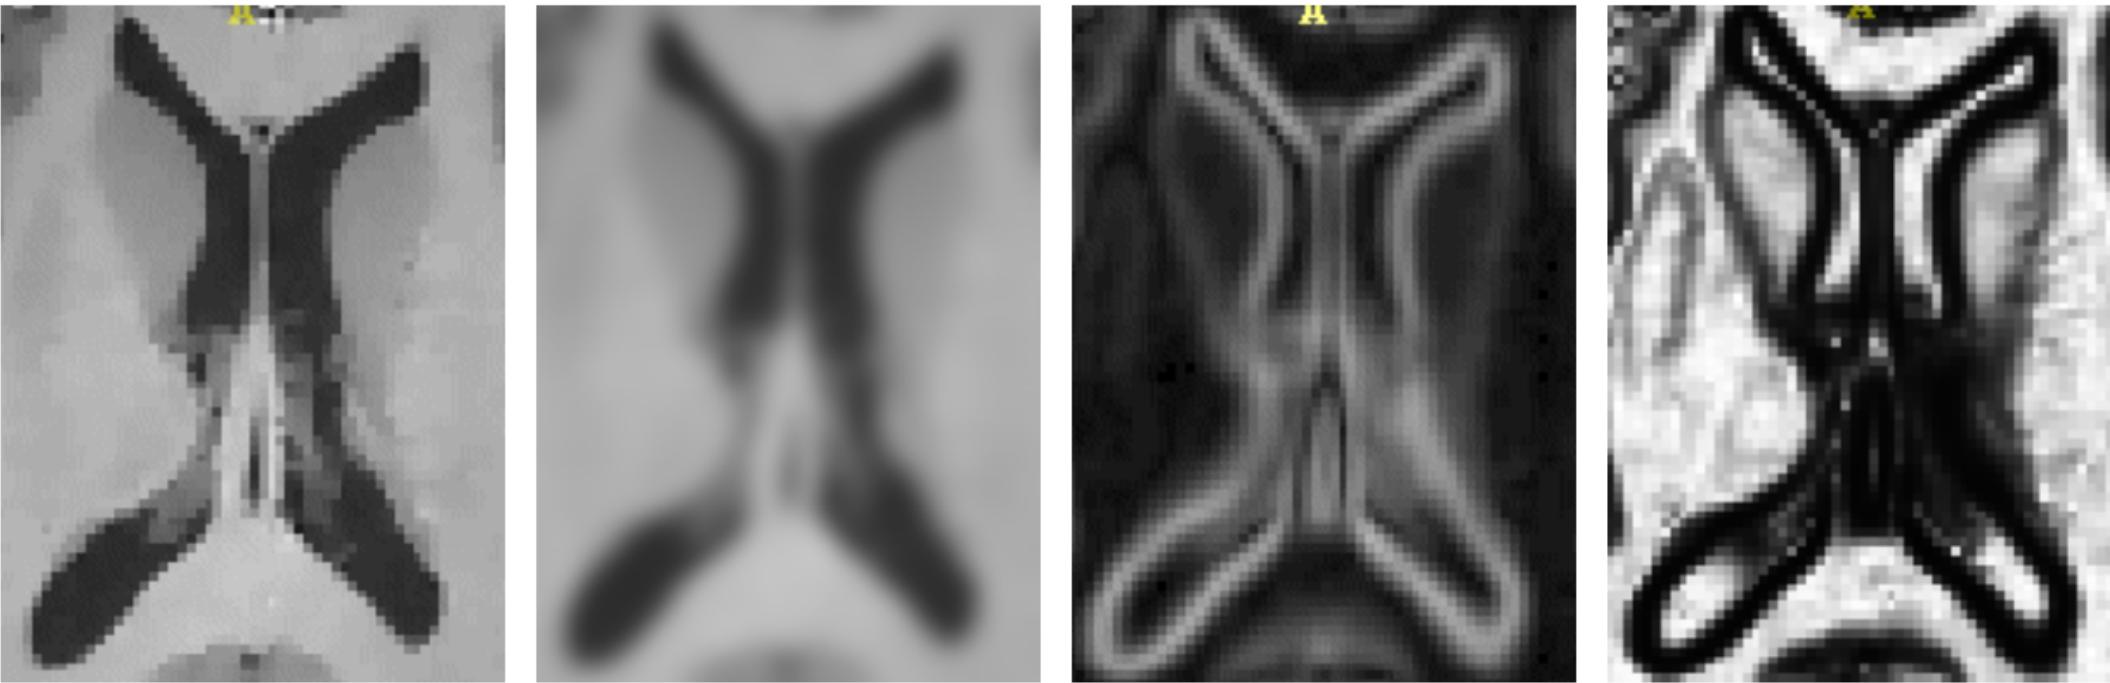
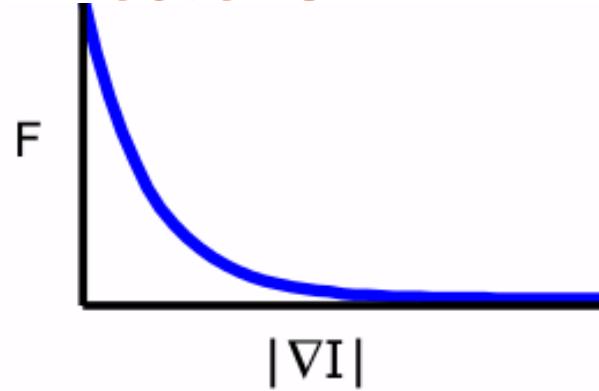
Level-Set Based Segmentation

- Anatomical structure segmentation
 - How to make the following happen ?
 - **How to design $F(x)$?**



Level-Set Based Segmentation

- Anatomical structure segmentation
 - Designing spatially-varying speed $F(x)$



- **What is a limitation of fast-marching based segmentation ?**

Level-Set Based Segmentation

- Limitations of fast-marching segmentation
 - Level-set curve always moves in 1 direction (outwards)
 - $F > 0$, always, for all pixel locations
 - What if we miss an edge ?
- We want the level-set curve to be able to move BOTH outwards and inwards
 - F can be positive and negative
 - Use intensity / texture statistics to design F , instead of using only edge information

L



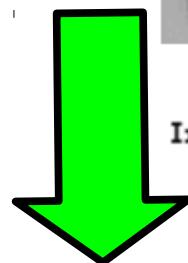
used S



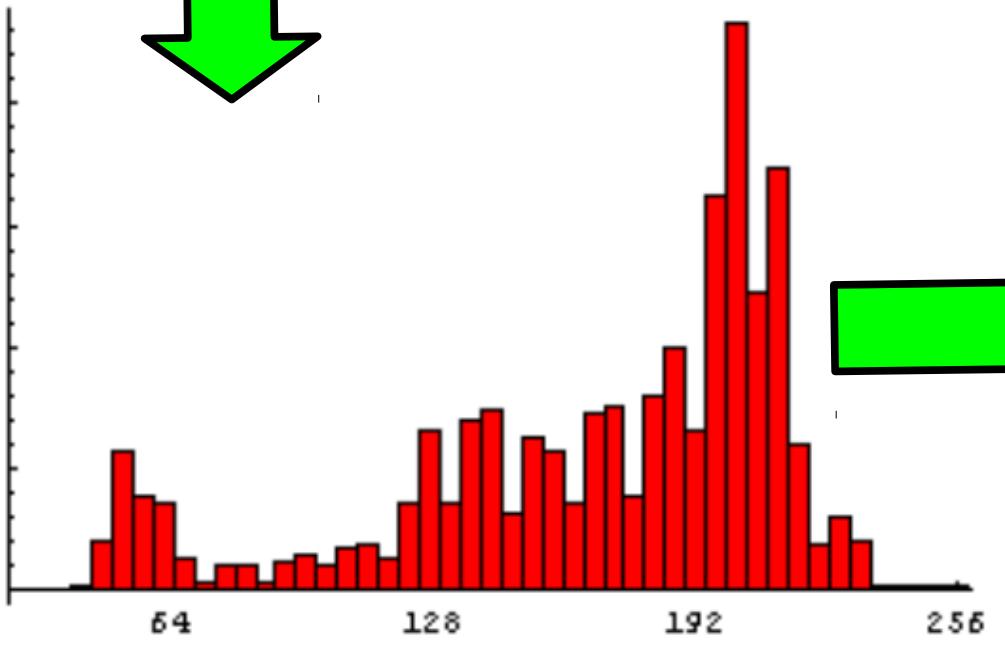
1.0

0.0

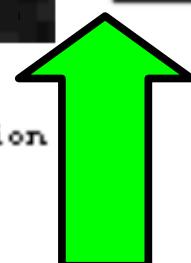
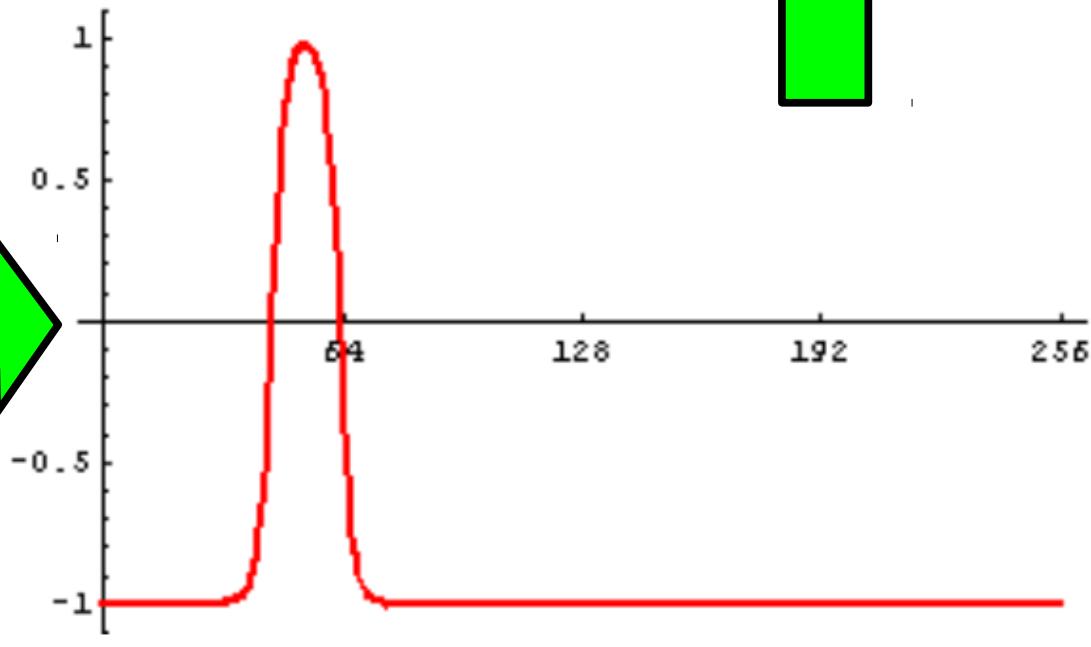
-1.0



Intensity Histogram

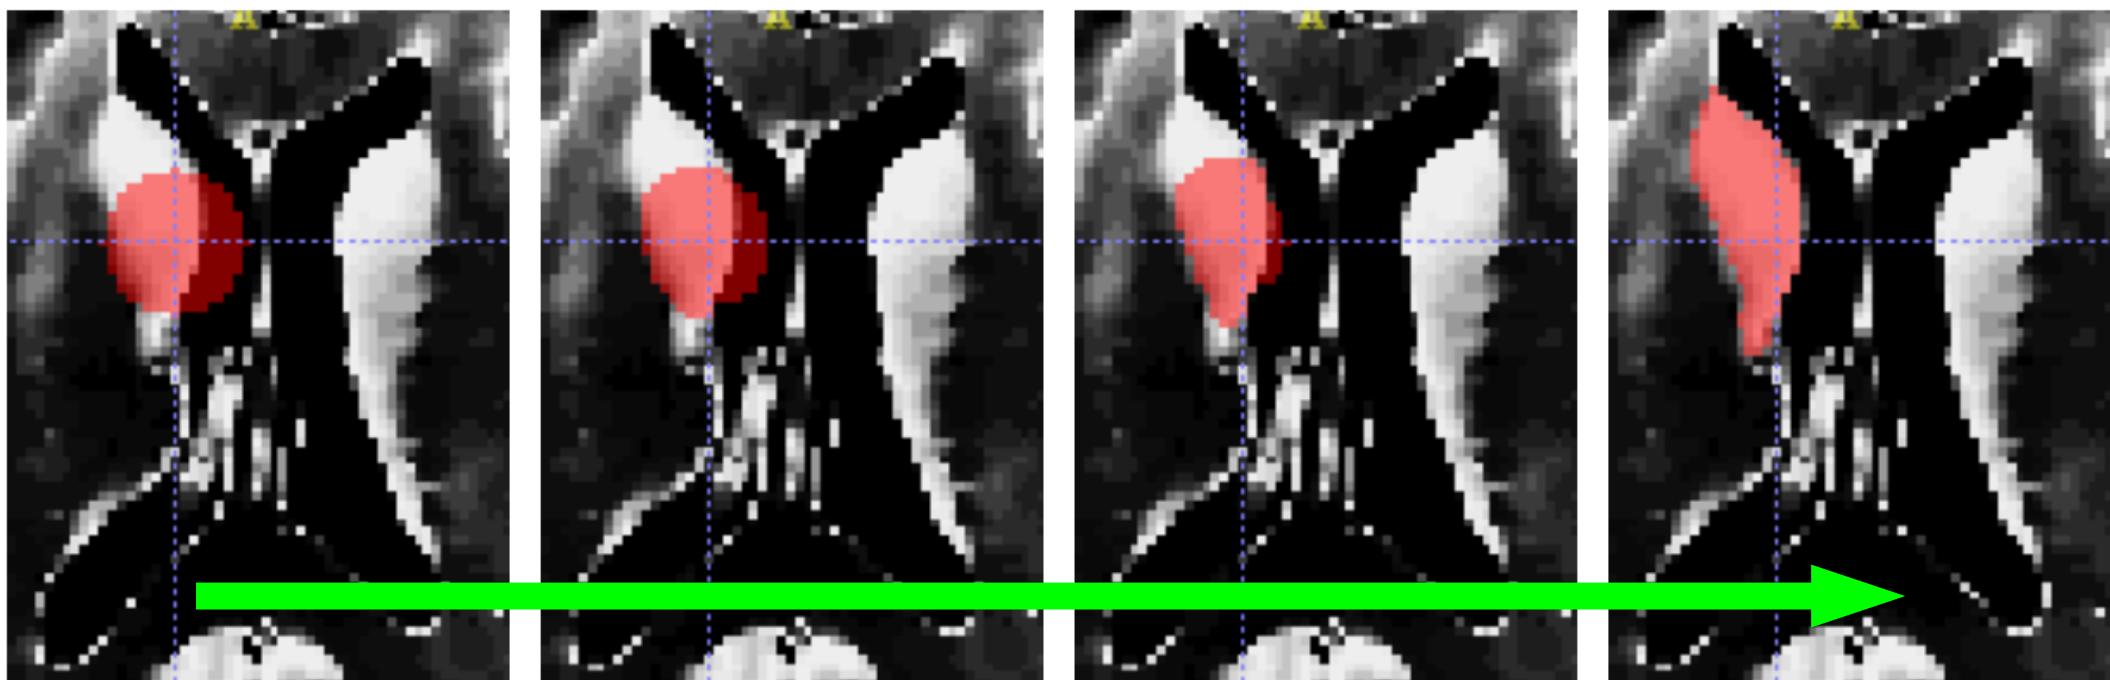
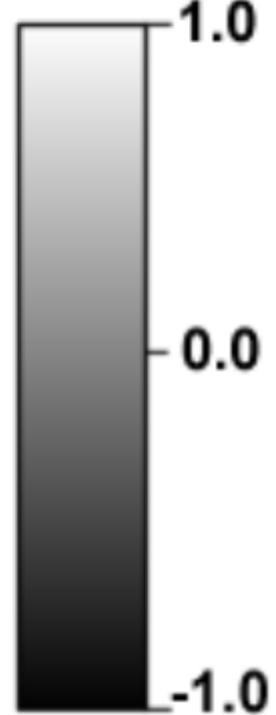


Remapping Function



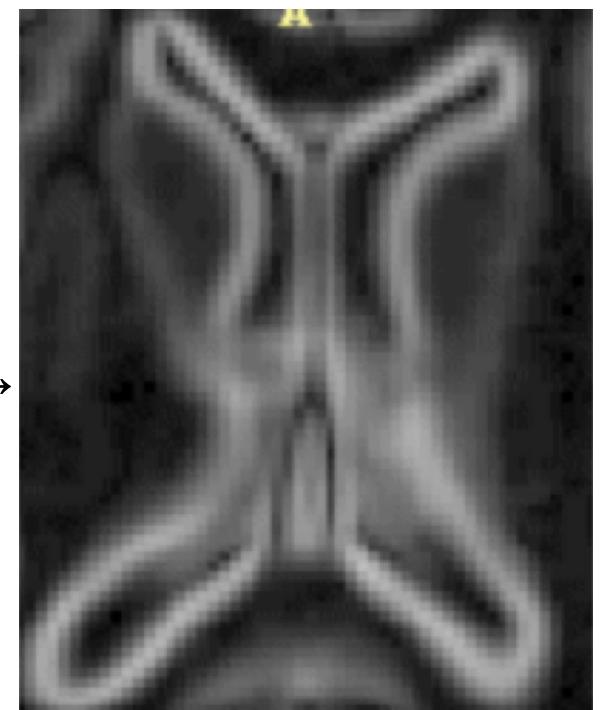
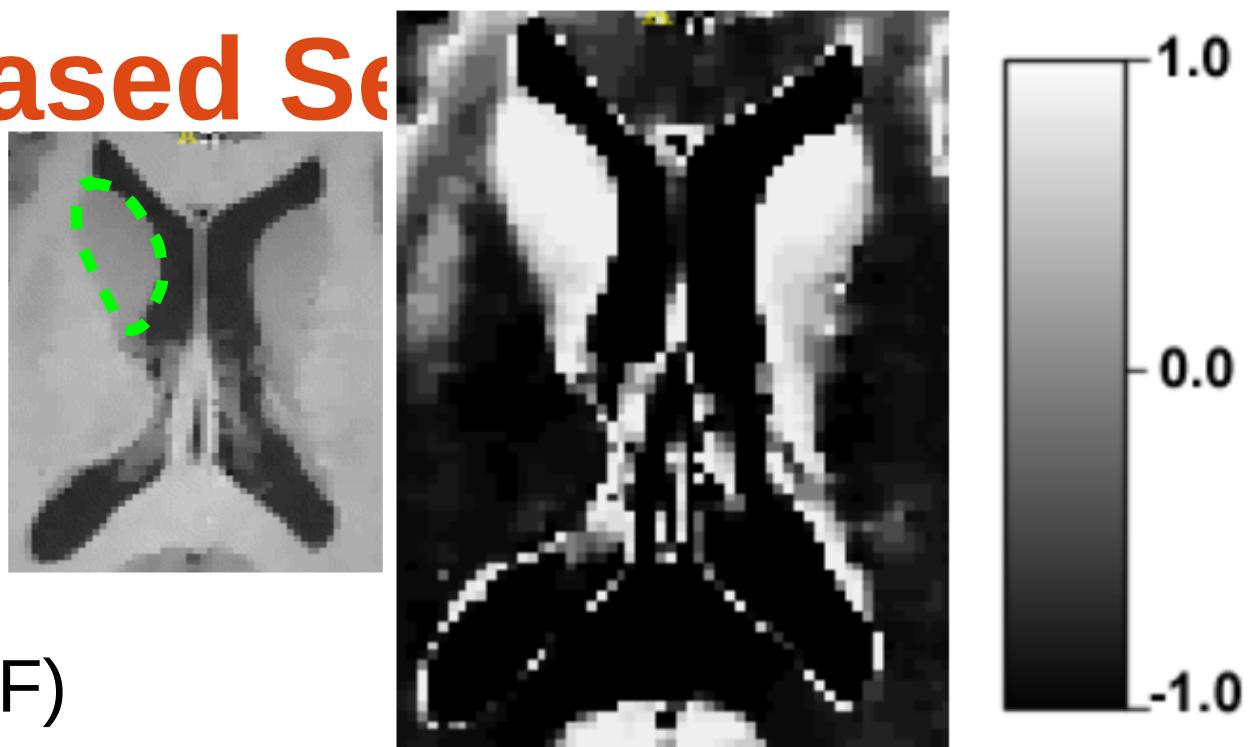
Level-Set Based S

- Level set evolution



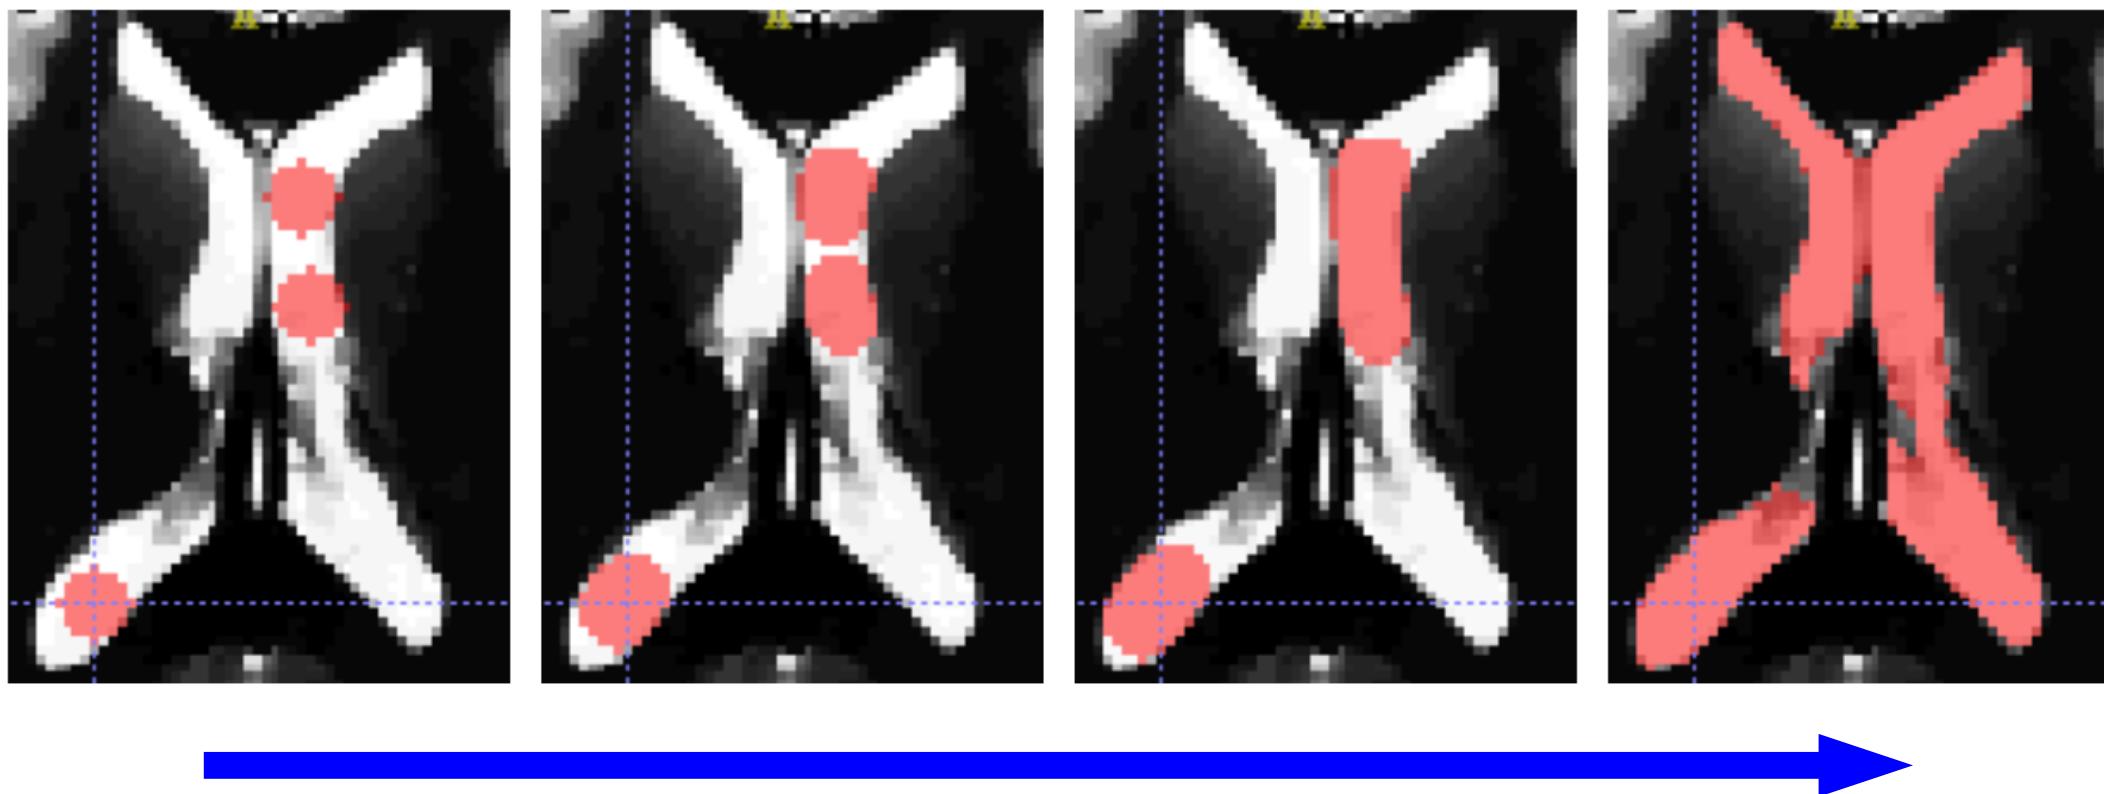
Level-Set Based Segmentation

- Types of forces
 - 'Propagation' force
 - Proportional to $F(x)$
 - Curvature force (MCF)
 - For maintaining smoothness of segmentation
 - 'Advection' force
 - Proportional to inner product of
 - (i) level-set normal and
 - (ii) gradient of edge (grad mag) image →
 - Can prevent leakage by pulling level set to edges



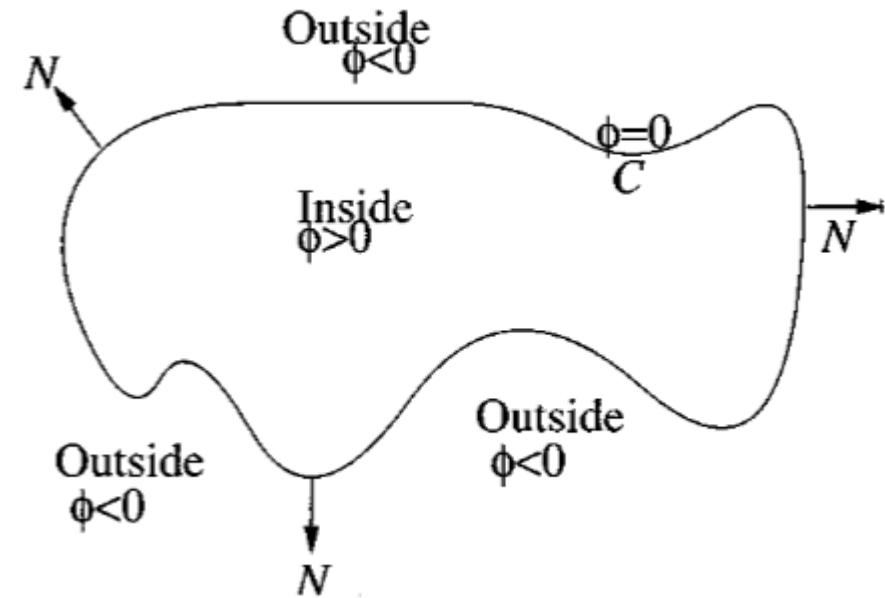
Level-Set Based Segmentation

- Initialization
 - Can be disconnected blobs
 - Merging handled easily



Level-Set Based Segmentation

- Optimizing segmentation + region-based intensity model
- Assume 2 regions: foreground (fg) + background (bg)
- Optimize segmentation to minimize:
 - Length of boundary
 - + (variance of fg * area of fg)
 - + (variance of bg * area of bg)
- If $\varphi(x)$ = distance function, then
 - boundary length =
$$\int_{\Omega} |\nabla H(\varphi(x))| dx = \int_{\Omega} \delta(\varphi(x)) |\nabla \varphi(x)| dx$$



Level-Set Based Segmentation

- Optimizing segmentation + intensity model
- Optimize
$$F(c_1, c_2, \mathcal{C}) = \mu \int_{\Omega} \delta(\Phi) |\nabla \Phi| dx dy + \lambda_1 \int_{\Omega} (I(x, y) - c_1)^2 H(\Phi(x, y)) dx dy + \lambda_2 \int_{\Omega} (I(x, y) - c_2)^2 (1 - H(\Phi(x, y))) dx dy$$
- Given a segmentation $\Phi(x)$,
optimal values for c_1, c_2 are = ?

Level-Set Based Segmentation

- Optimizing segmentation + intensity model
- Optimize
$$F(c_1, c_2, \mathcal{C}) = \mu \int_{\Omega} \delta(\Phi) |\nabla \Phi| dx dy + \lambda_1 \int_{\Omega} (I(x, y) - c_1)^2 H(\Phi(x, y)) dx dy + \lambda_2 \int_{\Omega} (I(x, y) - c_2)^2 (1 - H(\Phi(x, y))) dx dy$$
- Given c_1, c_2 , how to modify segmentation $\Phi(x)$?
- Gradient descent on functional $F(\cdot)$
(uses Euler-Lagrange equation)

$$\frac{\partial \Phi}{\partial t} = \delta_\epsilon(\Phi) \left(\mu \operatorname{div}\left(\frac{\nabla \Phi}{|\nabla \Phi|}\right) - \lambda_1 (I - c_1)^2 + \lambda_2 (I - c_2)^2 \right)$$

Level-Set Based Segmentation

- Optimizing segmentation + intensity model
- Solve a PDE

$$\frac{\partial \Phi}{\partial t} = \delta_\epsilon(\Phi) \left(\mu \operatorname{div}\left(\frac{\nabla \Phi}{|\nabla \Phi|}\right) - \lambda_1 (I - c_1)^2 + \lambda_2 (I - c_2)^2 \right)$$

- For surfaces in 3D,
Divergence of unit normal = mean curvature * (-1)

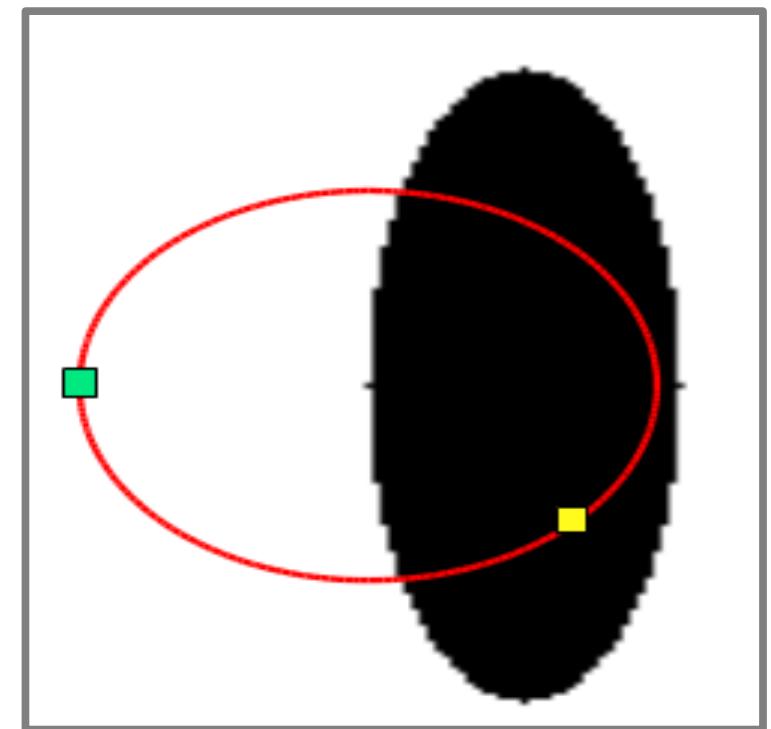
Level-Set Based Segmentation

- Optimizing segmentation + intensity model

$$\frac{\partial \Phi}{\partial t} = \delta_\epsilon(\Phi) \left(\mu \operatorname{div}\left(\frac{\nabla \Phi}{|\nabla \Phi|}\right) - \lambda_1(I - c_1)^2 + \lambda_2(I - c_2)^2 \right)$$

- Example

- Let c_1 : inside : closer to black
- Let c_2 : outside : closer to white
- **Green** point:
 $(I - c_2) = \text{small}$
 $(I - c_1) = \text{large}$
Moves inwards
- **Yellow** point:
Moves outwards



Oliver Heaviside

- Self-taught electrical engineer, mathematician, physicist
- Parents could not keep him at school after he was 16
- 1st winner of Faraday Medal (1922): top medal award by Institution of Engineering and Technology (previously called Institution of Electrical Engineers)
 - Donald Knuth won it in 2011

Oliver Heaviside



Mathematics is an experimental science, and definitions do not come first, but later on.

~ Oliver Heaviside

Oliver Heaviside

Oliver Heaviside

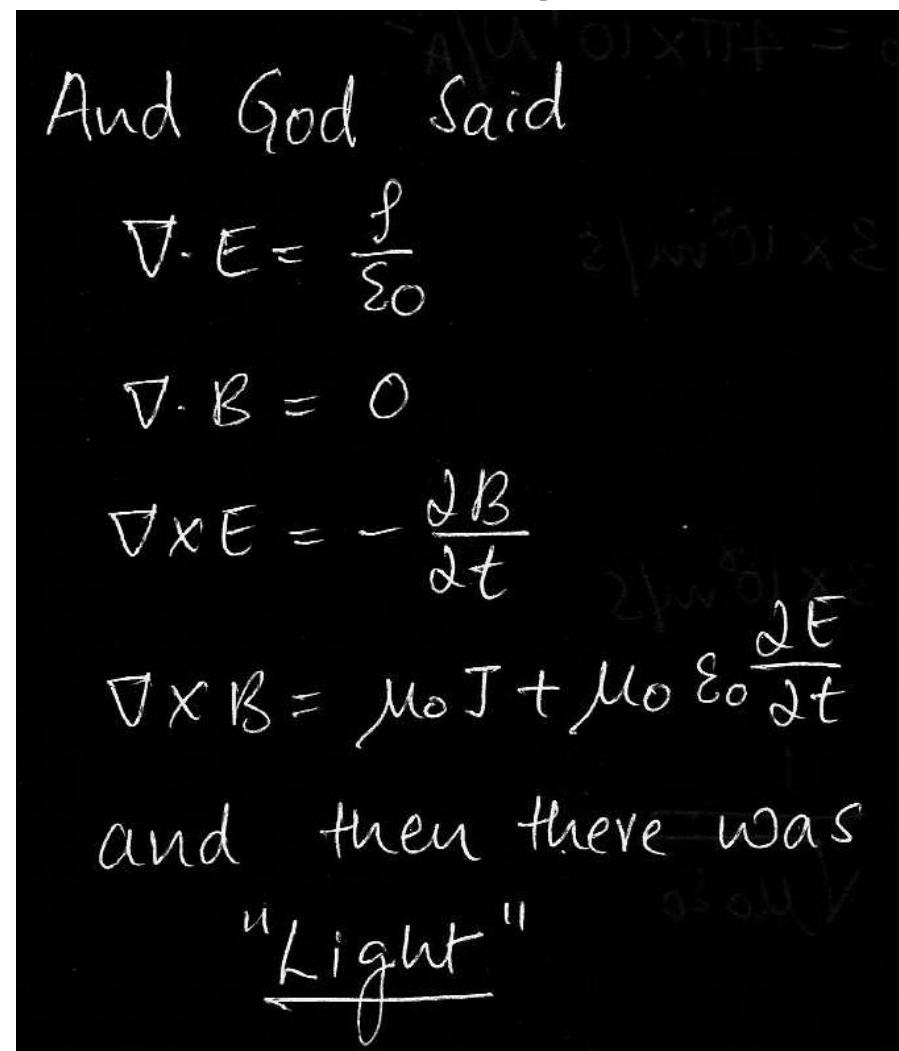


Why should I refuse a good dinner simply because I don't understand the digestive processes involved?

AZ QUOTES

Oliver Heaviside

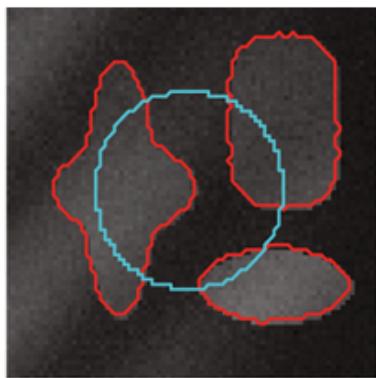
- Maxwell's formulation (1861-62) of electromagnetism consisted of 20 equations in 20 variables.
- Heaviside used curl and divergence operators of vector calculus to reformulate 12 of these 20 equations into 4 equations in 4 variables (B , E , J , ρ), the form by which they have been known ever since



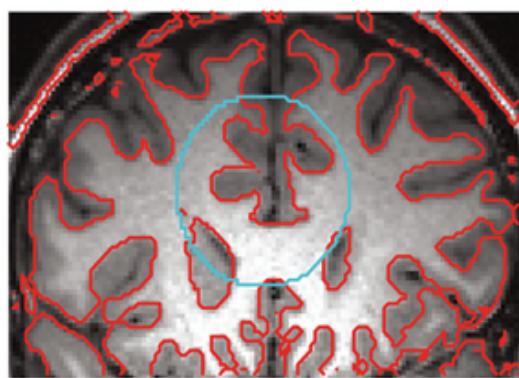
James Watt

- James Watt built steam engine in 1781
 - Lead to the industrial revolution
- When was theory of thermodynamics developed ?
 - Thermodynamics developed out of a desire to increase the efficiency of early steam engines
 - Carnot, 1824
 - Lord Kelvin, 1854
 - Gibbs, 1873-76, Statistical mechanics / thermodynamics
 - Explained classical thermodynamics as a natural result of statistics, classical mechanics, quantum theory at micro level

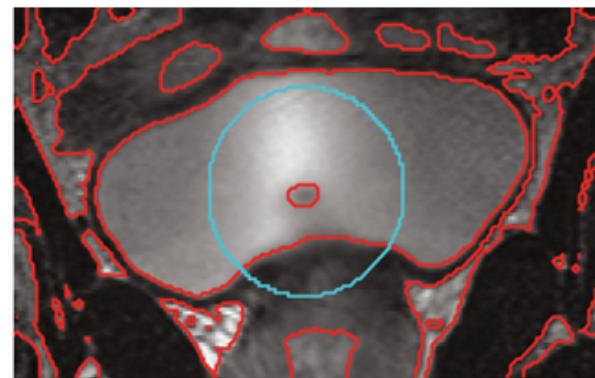
Level-Set Based Segmentation



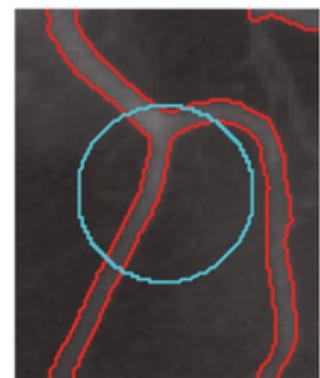
(a)



(b)



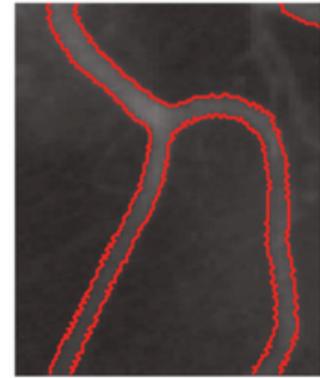
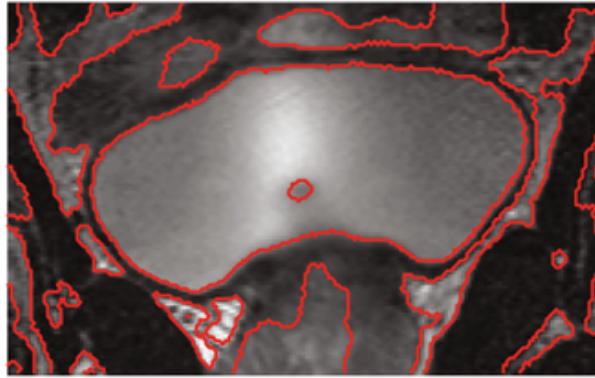
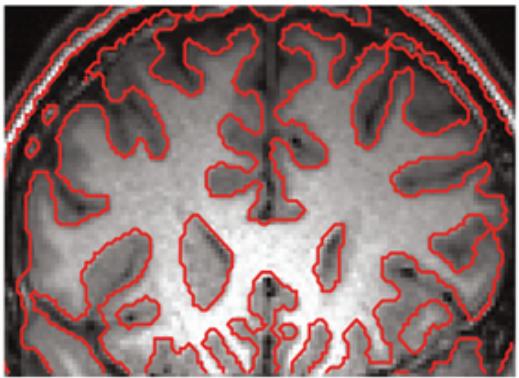
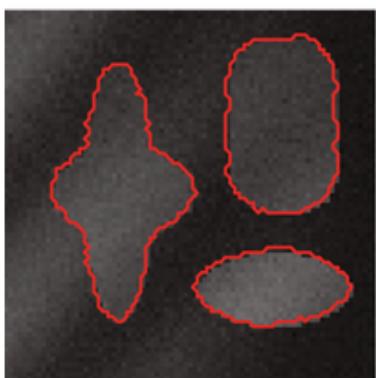
(c)



(d)



(e)



Level-Set Based Segmentation



(a)



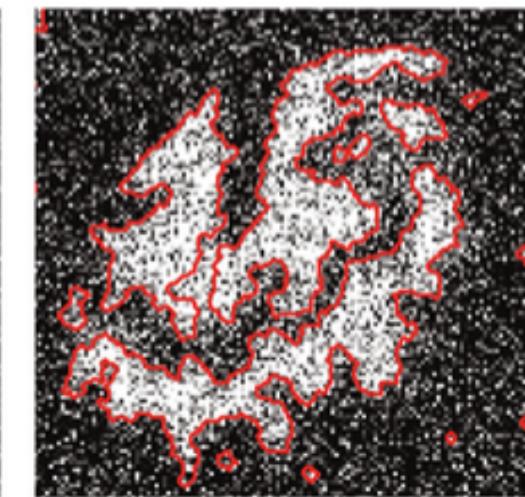
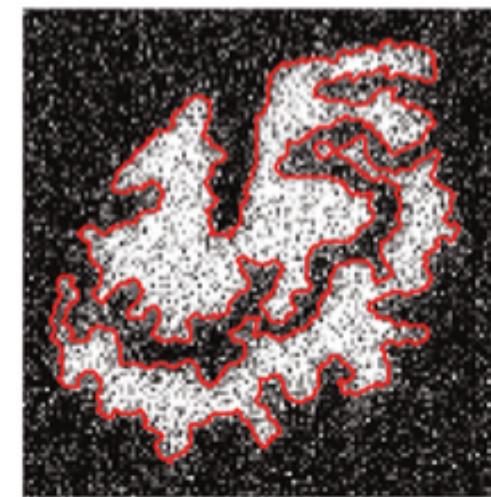
(b)



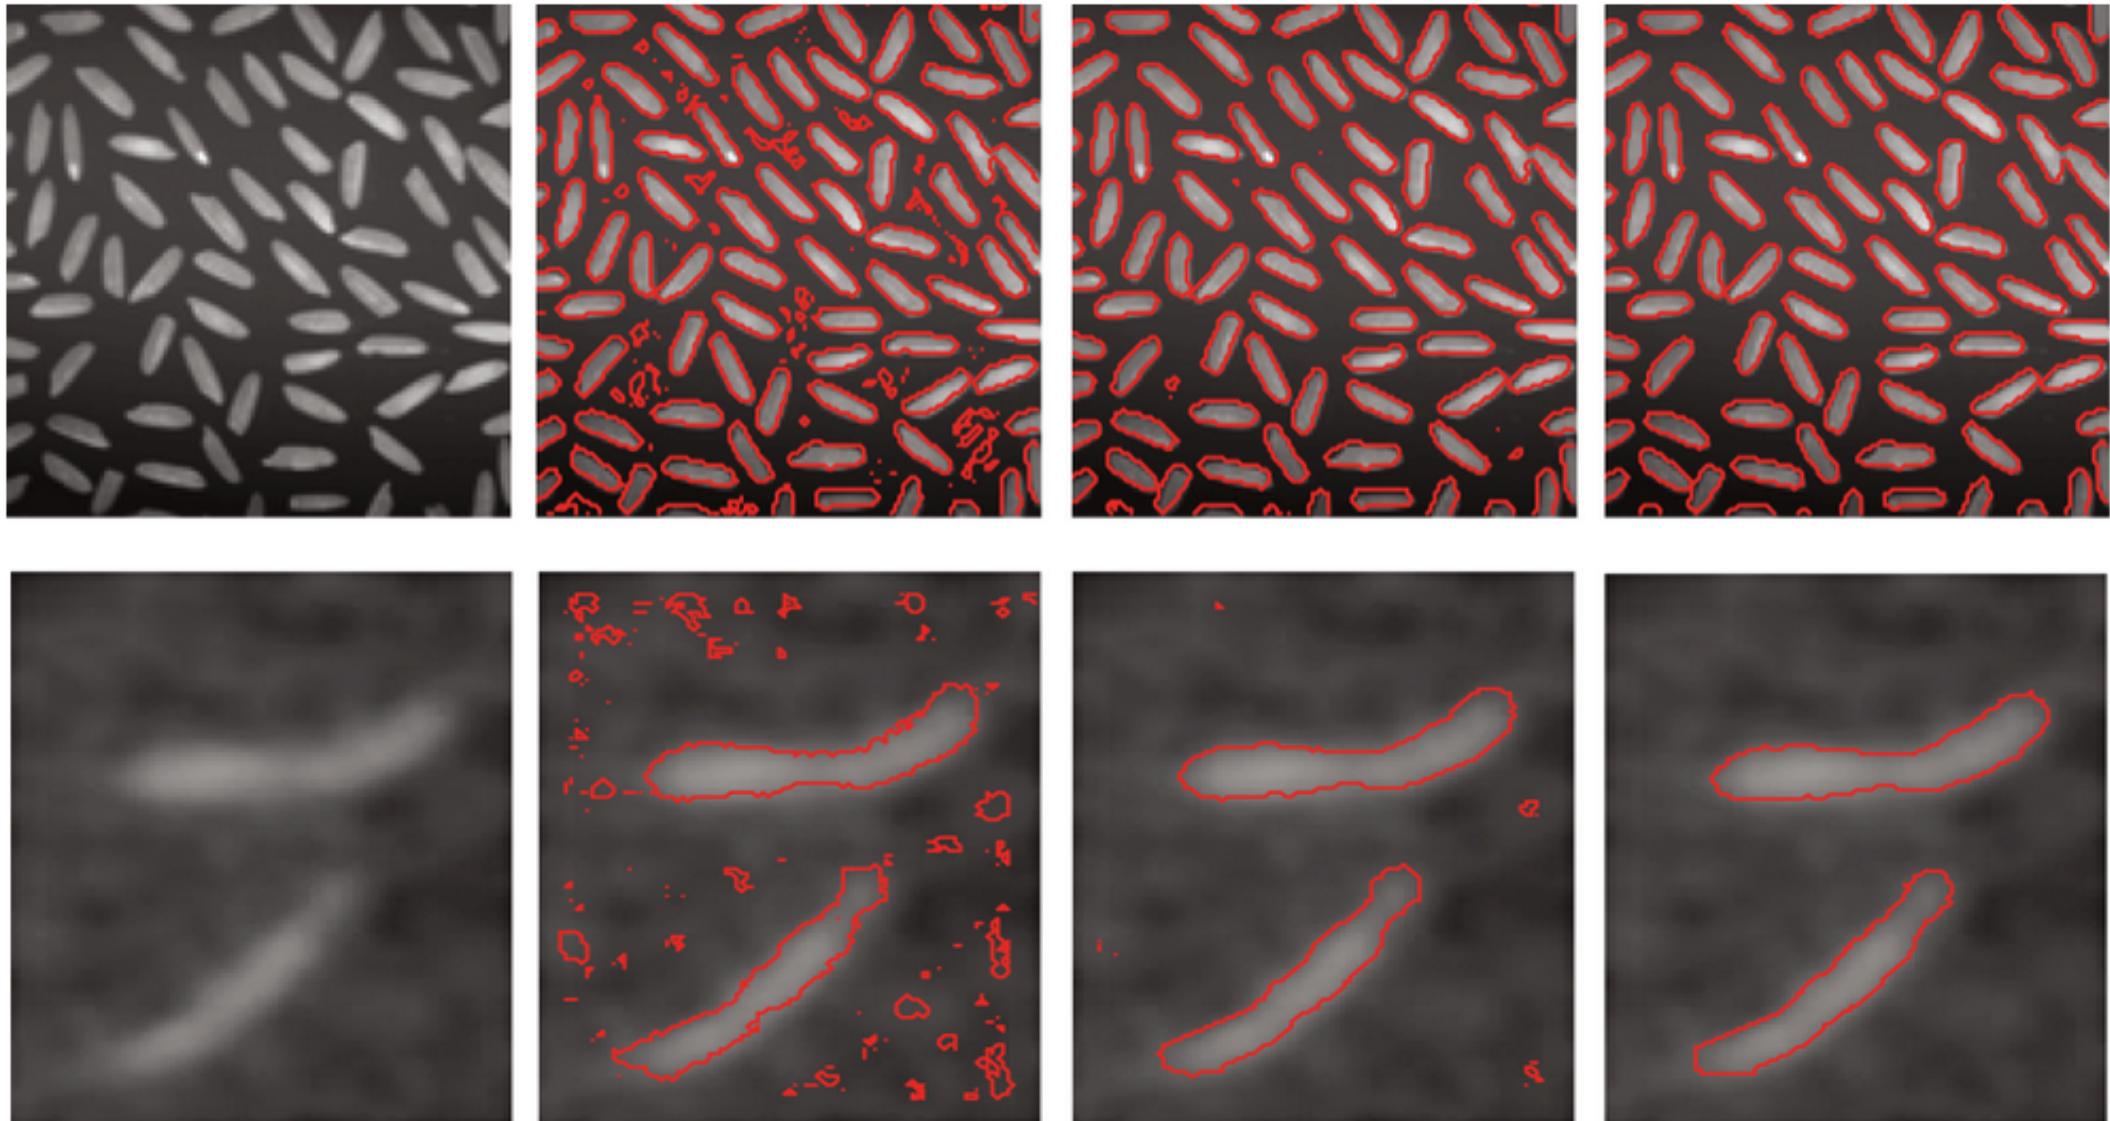
(c)



(d)



Level-Set Based Segmentation



Level-Set Based Segmentation

- Increasing regularization $\mu \rightarrow \dots$

